

# Machine Learning Project Report

## **Group Members:**

Name	Student No.	Section
Kunal Naresh Kumar Thapar	1107686	COMP-5413-FA
Rahul Niraj Singh	1099198	COMP-5413-FA

## **Task:**

Object-centric Image recognition task

## **Dataset:**

CIFAR10, Caltech101, Caltech256, CIFAR100

# Machine Learning Project Report

## CIFAR-10

### Introduction:

The CIFAR-10 dataset consists of 60000 32x32 colour images in 10 classes, with 6000 images per class. There are 50000 training images and 10000 test images.

The dataset is divided into five training batches and one test batch, each with 10000 images. The test batch contains exactly 1000 randomly-selected images from each class. The training batches contain the remaining images in random order, but some training batches may contain more images from one class than another. Between them, the training batches contain exactly 5000 images from each class.

Dataset Link: <https://www.cs.toronto.edu/~kriz/cifar-10-python.tar.gz>

### Model:

We have used Keras model based on VGG16 architecture for CIFAR-10 dataset.

### VGG16 Architecture:

This network is characterized by its simplicity, using only 3x3 convolutional layers stacked on top of each other in increasing depth. Reducing volume size is handled by max pooling. Two fully-connected layers, each with 4,096 nodes are then followed by a softmax classifier.

### Source Code:

```
from __future__ import print_function
import keras
from keras.datasets import cifar10
from keras.preprocessing.image import ImageDataGenerator
from keras.models import Sequential
from keras.layers import Dense, Dropout, Activation, Flatten
from keras.layers import Conv2D, MaxPooling2D, BatchNormalization
from keras import optimizers
import numpy as np
from keras.layers.core import Lambda
from keras import backend as K
from keras import regularizers
```

# Machine Learning Project Report

```
import time
from sklearn.metrics import accuracy_score

class cifar10vgg:
    def __init__(self,train=True):
        self.num_classes = 10
        self.weight_decay = 0.0005
        self.x_shape = [32,32,3]

        self.model = self.build_model()
        if train:
            self.model = self.train(self.model)
        else:
            self.model.load_weights('cifar10vgg.h5')

    def build_model(self):
        # Build the network of vgg for 10 classes with massive dropout and weight decay as described in the
        # paper.

        model = Sequential()
        weight_decay = self.weight_decay

        model.add(Conv2D(64, (3, 3), padding='same',
                        input_shape=self.x_shape,kernel_regularizer=regularizers.l2(weight_decay)))
        model.add(Activation('relu'))
        model.add(BatchNormalization())
        model.add(Dropout(0.3))

        model.add(Conv2D(64, (3, 3), padding='same',kernel_regularizer=regularizers.l2(weight_decay)))
        model.add(Activation('relu'))
        model.add(BatchNormalization())

        model.add(MaxPooling2D(pool_size=(2, 2)))

        model.add(Conv2D(128, (3, 3), padding='same',kernel_regularizer=regularizers.l2(weight_decay)))
        model.add(Activation('relu'))
        model.add(BatchNormalization())
        model.add(Dropout(0.4))

        model.add(Conv2D(128, (3, 3), padding='same',kernel_regularizer=regularizers.l2(weight_decay)))
        model.add(Activation('relu'))
        model.add(BatchNormalization())

        model.add(MaxPooling2D(pool_size=(2, 2)))

        model.add(Conv2D(256, (3, 3), padding='same',kernel_regularizer=regularizers.l2(weight_decay)))
```

# Machine Learning Project Report

```
model.add(Activation('relu'))
model.add(BatchNormalization())
model.add(Dropout(0.4))
```

```
model.add(Conv2D(256, (3, 3), padding='same',kernel_regularizer=regularizers.l2(weight_decay)))
model.add(Activation('relu'))
model.add(BatchNormalization())
model.add(Dropout(0.4))
```

```
model.add(Conv2D(256, (3, 3), padding='same',kernel_regularizer=regularizers.l2(weight_decay)))
model.add(Activation('relu'))
model.add(BatchNormalization())
```

```
model.add(MaxPooling2D(pool_size=(2, 2)))
```

```
model.add(Conv2D(512, (3, 3), padding='same',kernel_regularizer=regularizers.l2(weight_decay)))
model.add(Activation('relu'))
model.add(BatchNormalization())
model.add(Dropout(0.4))
```

```
model.add(Conv2D(512, (3, 3), padding='same',kernel_regularizer=regularizers.l2(weight_decay)))
model.add(Activation('relu'))
model.add(BatchNormalization())
model.add(Dropout(0.4))
```

```
model.add(Conv2D(512, (3, 3), padding='same',kernel_regularizer=regularizers.l2(weight_decay)))
model.add(Activation('relu'))
model.add(BatchNormalization())
```

```
model.add(MaxPooling2D(pool_size=(2, 2)))
```

```
model.add(Conv2D(512, (3, 3), padding='same',kernel_regularizer=regularizers.l2(weight_decay)))
model.add(Activation('relu'))
model.add(BatchNormalization())
model.add(Dropout(0.4))
```

```
model.add(Conv2D(512, (3, 3), padding='same',kernel_regularizer=regularizers.l2(weight_decay)))
model.add(Activation('relu'))
model.add(BatchNormalization())
model.add(Dropout(0.4))
```

```
model.add(Conv2D(512, (3, 3), padding='same',kernel_regularizer=regularizers.l2(weight_decay)))
model.add(Activation('relu'))
model.add(BatchNormalization())
```

```
model.add(MaxPooling2D(pool_size=(2, 2)))
```

# Machine Learning Project Report

```
model.add(Dropout(0.5))
```

```
model.add(Flatten())
model.add(Dense(512,kernel_regularizer=regularizers.l2(weight_decay)))
model.add(Activation('relu'))
model.add(BatchNormalization())
```

```
model.add(Dropout(0.5))
model.add(Dense(self.num_classes))
model.add(Activation('softmax'))
return model
```

```
def normalize(self,X_train,X_test):
    #this function normalize inputs for zero mean and unit variance
    # it is used when training a model.
    # Input: training set and test set
    # Output: normalized training set and test set according to the training set statistics.
    mean = np.mean(X_train,axis=(0,1,2,3))
    std = np.std(X_train, axis=(0, 1, 2, 3))
    X_train = (X_train-mean)/(std+1e-7)
    X_test = (X_test-mean)/(std+1e-7)
    return X_train, X_test
```

```
def normalize_production(self,x):
    #this function is used to normalize instances in production according to saved training set statistics
    # Input: X - a training set
    # Output X - a normalized training set according to normalization constants.
```

```
    #these values produced during first training and are general for the standard cifar10 training set normalization
    mean = 120.707
    std = 64.15
    return (x-mean)/(std+1e-7)
```

```
def predict(self,x,normalize=True,batch_size=50):
    if normalize:
        x = self.normalize_production(x)
    return self.model.predict(x,batch_size)
```

```
def train(self,model):

    #training parameters
    batch_size = 128
    maxepoches = 250
    learning_rate = 0.1
    lr_decay = 1e-6
```

# Machine Learning Project Report

```
lr_drop = 20
# The data, shuffled and split between train and test sets:
(x_train, y_train), (x_test, y_test) = cifar10.load_data()
x_train = x_train.astype('float32')
x_test = x_test.astype('float32')
x_train, x_test = self.normalize(x_train, x_test)

y_train = keras.utils.to_categorical(y_train, self.num_classes)
y_test = keras.utils.to_categorical(y_test, self.num_classes)

def lr_scheduler(epoch):
    return learning_rate * (0.5 ** (epoch // lr_drop))
reduce_lr = keras.callbacks.LearningRateScheduler(lr_scheduler)

#data augmentation
datagen = ImageDataGenerator(
    featurewise_center=False, # set input mean to 0 over the dataset
    samplewise_center=False, # set each sample mean to 0
    featurewise_std_normalization=False, # divide inputs by std of the dataset
    samplewise_std_normalization=False, # divide each input by its std
    zca_whitening=False, # apply ZCA whitening
    rotation_range=15, # randomly rotate images in the range (degrees, 0 to 180)
    width_shift_range=0.1, # randomly shift images horizontally (fraction of total width)
    height_shift_range=0.1, # randomly shift images vertically (fraction of total height)
    horizontal_flip=True, # randomly flip images
    vertical_flip=False) # randomly flip images
# (std, mean, and principal components if ZCA whitening is applied).
start = time.time()
datagen.fit(x_train)

#optimization details
sgd = optimizers.SGD(lr=learning_rate, decay=lr_decay, momentum=0.9, nesterov=True)
model.compile(loss='categorical_crossentropy', optimizer=sgd, metrics=['accuracy'])

# training process in a for loop with learning rate drop every 25 epoches.

historytemp = model.fit_generator(datagen.flow(x_train, y_train,
                                              batch_size=batch_size),
                                steps_per_epoch=x_train.shape[0] // batch_size,
                                epochs=maxepoches,
                                validation_data=(x_test, y_test),callbacks=[reduce_lr],verbose=2)
model.save_weights('cifar10vgg.h5')
end = time.time()
print("Training Time: ",(end-start))
return model
```

# Machine Learning Project Report

```
if __name__ == '__main__':

    (x_train, y_train), (x_test, y_test) = cifar10.load_data()
    x_train = x_train.astype('float32')
    x_test = x_test.astype('float32')

    y_train = keras.utils.to_categorical(y_train, 10)
    y_test = keras.utils.to_categorical(y_test, 10)

    model = cifar10vgg()

    predicted_x = model.predict(x_test)
    accuracy = accuracy_score(np.argmax(predicted_x,1) , np.argmax(y_test,1))
    print("Accuracy: ",accuracy*100,"%" )
```

## Output:

### First Run:

```
Epoch 1/250
- 32s - loss: 19.2023 - acc: 0.1889 - val_loss: 17.2594 - val_acc: 0.1217
Epoch 2/250
- 27s - loss: 13.6068 - acc: 0.2380 - val_loss: 10.1019 - val_acc: 0.1634
Epoch 3/250
- 28s - loss: 7.5740 - acc: 0.2756 - val_loss: 6.2211 - val_acc: 0.1888
Epoch 4/250
- 28s - loss: 5.1312 - acc: 0.3028 - val_loss: 5.4339 - val_acc: 0.1549
Epoch 5/250
- 28s - loss: 5.0240 - acc: 0.2932 - val_loss: 4.6637 - val_acc: 0.2218
Epoch 6/250
- 27s - loss: 3.5551 - acc: 0.3966 - val_loss: 3.6040 - val_acc: 0.1930
Epoch 7/250
- 28s - loss: 2.4914 - acc: 0.4785 - val_loss: 2.3857 - val_acc: 0.4381
Epoch 8/250
- 27s - loss: 1.9851 - acc: 0.5427 - val_loss: 1.9532 - val_acc: 0.5015
Epoch 9/250
- 28s - loss: 1.7069 - acc: 0.5986 - val_loss: 1.8182 - val_acc: 0.5782
Epoch 10/250
- 28s - loss: 1.5898 - acc: 0.6369 - val_loss: 1.5059 - val_acc: 0.6916
Epoch 11/250
- 28s - loss: 1.5718 - acc: 0.6557 - val_loss: 1.4599 - val_acc: 0.6838
```

# Machine Learning Project Report

Epoch 12/250  
- 27s - loss: 1.4975 - acc: 0.6826 - val\_loss: 1.4220 - val\_acc: 0.7159  
Epoch 13/250  
- 27s - loss: 1.4755 - acc: 0.6961 - val\_loss: 1.5002 - val\_acc: 0.6911  
Epoch 14/250  
- 27s - loss: 1.4805 - acc: 0.7036 - val\_loss: 1.3736 - val\_acc: 0.7338  
Epoch 15/250  
- 27s - loss: 1.4720 - acc: 0.7109 - val\_loss: 1.5302 - val\_acc: 0.6971  
Epoch 16/250  
- 27s - loss: 1.4873 - acc: 0.7099 - val\_loss: 1.5039 - val\_acc: 0.7041  
Epoch 17/250  
- 27s - loss: 1.4616 - acc: 0.7217 - val\_loss: 1.3945 - val\_acc: 0.7454  
Epoch 18/250  
- 27s - loss: 1.4736 - acc: 0.7263 - val\_loss: 1.6964 - val\_acc: 0.6644  
Epoch 19/250  
- 27s - loss: 1.4721 - acc: 0.7323 - val\_loss: 1.3471 - val\_acc: 0.7733  
Epoch 20/250  
- 27s - loss: 1.4864 - acc: 0.7302 - val\_loss: 1.5061 - val\_acc: 0.7329  
Epoch 21/250  
- 27s - loss: 1.3163 - acc: 0.7730 - val\_loss: 1.2446 - val\_acc: 0.7805  
Epoch 22/250  
- 27s - loss: 1.2260 - acc: 0.7799 - val\_loss: 1.1218 - val\_acc: 0.8046  
Epoch 23/250  
- 27s - loss: 1.1954 - acc: 0.7804 - val\_loss: 1.1525 - val\_acc: 0.7937  
Epoch 24/250  
- 26s - loss: 1.1939 - acc: 0.7833 - val\_loss: 1.1571 - val\_acc: 0.7919  
Epoch 25/250  
- 26s - loss: 1.1867 - acc: 0.7843 - val\_loss: 1.1062 - val\_acc: 0.8129  
Epoch 26/250  
- 26s - loss: 1.1968 - acc: 0.7862 - val\_loss: 1.1712 - val\_acc: 0.7984  
Epoch 27/250  
- 26s - loss: 1.1880 - acc: 0.7910 - val\_loss: 1.1188 - val\_acc: 0.8132  
Epoch 28/250  
- 26s - loss: 1.1982 - acc: 0.7912 - val\_loss: 1.1058 - val\_acc: 0.8210  
Epoch 29/250  
- 26s - loss: 1.2004 - acc: 0.7920 - val\_loss: 1.1170 - val\_acc: 0.8162  
Epoch 30/250  
- 26s - loss: 1.2014 - acc: 0.7956 - val\_loss: 1.1423 - val\_acc: 0.8137  
Epoch 31/250  
- 26s - loss: 1.2059 - acc: 0.7945 - val\_loss: 1.1261 - val\_acc: 0.8167  
Epoch 32/250  
- 26s - loss: 1.2034 - acc: 0.7964 - val\_loss: 1.1947 - val\_acc: 0.8004  
Epoch 33/250  
- 26s - loss: 1.2102 - acc: 0.7975 - val\_loss: 1.1822 - val\_acc: 0.8084  
Epoch 34/250  
- 26s - loss: 1.2160 - acc: 0.7965 - val\_loss: 1.1477 - val\_acc: 0.8262  
Epoch 35/250  
- 26s - loss: 1.2164 - acc: 0.8001 - val\_loss: 1.2160 - val\_acc: 0.7997  
Epoch 36/250  
- 26s - loss: 1.2203 - acc: 0.7993 - val\_loss: 1.1232 - val\_acc: 0.8275  
Epoch 37/250  
- 26s - loss: 1.2149 - acc: 0.8025 - val\_loss: 1.1194 - val\_acc: 0.8327



# Machine Learning Project Report

Epoch 38/250  
- 26s - loss: 1.2166 - acc: 0.8030 - val\_loss: 1.2035 - val\_acc: 0.8083  
Epoch 39/250  
- 26s - loss: 1.2274 - acc: 0.8020 - val\_loss: 1.1566 - val\_acc: 0.8187  
Epoch 40/250  
- 26s - loss: 1.2154 - acc: 0.8044 - val\_loss: 1.1749 - val\_acc: 0.8182  
Epoch 41/250  
- 26s - loss: 1.0978 - acc: 0.8364 - val\_loss: 0.9581 - val\_acc: 0.8696  
Epoch 42/250  
- 26s - loss: 1.0231 - acc: 0.8449 - val\_loss: 0.9451 - val\_acc: 0.8592  
Epoch 43/250  
- 26s - loss: 0.9858 - acc: 0.8440 - val\_loss: 0.9199 - val\_acc: 0.8638  
Epoch 44/250  
- 26s - loss: 0.9756 - acc: 0.8424 - val\_loss: 0.9457 - val\_acc: 0.8518  
Epoch 45/250  
- 26s - loss: 0.9623 - acc: 0.8455 - val\_loss: 0.8753 - val\_acc: 0.8745  
Epoch 46/250  
- 26s - loss: 0.9580 - acc: 0.8433 - val\_loss: 1.0023 - val\_acc: 0.8285  
Epoch 47/250  
- 26s - loss: 0.9539 - acc: 0.8432 - val\_loss: 0.8915 - val\_acc: 0.8608  
Epoch 48/250  
- 26s - loss: 0.9466 - acc: 0.8455 - val\_loss: 1.0366 - val\_acc: 0.8156  
Epoch 49/250  
- 26s - loss: 0.9617 - acc: 0.8418 - val\_loss: 0.8748 - val\_acc: 0.8720  
Epoch 50/250  
- 26s - loss: 0.9519 - acc: 0.8452 - val\_loss: 0.9278 - val\_acc: 0.8546  
Epoch 51/250  
- 26s - loss: 0.9555 - acc: 0.8442 - val\_loss: 0.9278 - val\_acc: 0.8555  
Epoch 52/250  
- 26s - loss: 0.9481 - acc: 0.8488 - val\_loss: 0.9034 - val\_acc: 0.8629  
Epoch 53/250  
- 26s - loss: 0.9440 - acc: 0.8503 - val\_loss: 0.9807 - val\_acc: 0.8423  
Epoch 54/250  
- 25s - loss: 0.9560 - acc: 0.8468 - val\_loss: 0.9160 - val\_acc: 0.8604  
Epoch 55/250  
- 26s - loss: 0.9564 - acc: 0.8485 - val\_loss: 0.9399 - val\_acc: 0.8535  
Epoch 56/250  
- 26s - loss: 0.9605 - acc: 0.8493 - val\_loss: 0.9484 - val\_acc: 0.8505  
Epoch 57/250  
- 26s - loss: 0.9576 - acc: 0.8502 - val\_loss: 0.9273 - val\_acc: 0.8602  
Epoch 58/250  
- 26s - loss: 0.9614 - acc: 0.8502 - val\_loss: 0.9504 - val\_acc: 0.8533  
Epoch 59/250  
- 26s - loss: 0.9616 - acc: 0.8496 - val\_loss: 0.9192 - val\_acc: 0.8653  
Epoch 60/250  
- 25s - loss: 0.9689 - acc: 0.8496 - val\_loss: 0.9281 - val\_acc: 0.8630  
Epoch 61/250  
- 26s - loss: 0.8729 - acc: 0.8768 - val\_loss: 0.8090 - val\_acc: 0.8937  
Epoch 62/250  
- 26s - loss: 0.8249 - acc: 0.8862 - val\_loss: 0.8332 - val\_acc: 0.8769  
Epoch 63/250  
- 26s - loss: 0.8053 - acc: 0.8854 - val\_loss: 0.8140 - val\_acc: 0.8796

# Machine Learning Project Report

Epoch 64/250  
- 26s - loss: 0.7862 - acc: 0.8862 - val\_loss: 0.7936 - val\_acc: 0.8856  
Epoch 65/250  
- 26s - loss: 0.7682 - acc: 0.8890 - val\_loss: 0.7498 - val\_acc: 0.8929  
Epoch 66/250  
- 26s - loss: 0.7573 - acc: 0.8898 - val\_loss: 0.7741 - val\_acc: 0.8825  
Epoch 67/250  
- 25s - loss: 0.7554 - acc: 0.8880 - val\_loss: 0.7396 - val\_acc: 0.8936  
Epoch 68/250  
- 26s - loss: 0.7462 - acc: 0.8880 - val\_loss: 0.7291 - val\_acc: 0.8959  
Epoch 69/250  
- 26s - loss: 0.7428 - acc: 0.8888 - val\_loss: 0.7072 - val\_acc: 0.8981  
Epoch 70/250  
- 25s - loss: 0.7430 - acc: 0.8877 - val\_loss: 0.7468 - val\_acc: 0.8890  
Epoch 71/250  
- 25s - loss: 0.7405 - acc: 0.8870 - val\_loss: 0.7448 - val\_acc: 0.8899  
Epoch 72/250  
- 26s - loss: 0.7409 - acc: 0.8858 - val\_loss: 0.7575 - val\_acc: 0.8858  
Epoch 73/250  
- 25s - loss: 0.7391 - acc: 0.8866 - val\_loss: 0.7401 - val\_acc: 0.8905  
Epoch 74/250  
- 26s - loss: 0.7409 - acc: 0.8868 - val\_loss: 0.7210 - val\_acc: 0.8901  
Epoch 75/250  
- 26s - loss: 0.7388 - acc: 0.8884 - val\_loss: 0.8357 - val\_acc: 0.8617  
Epoch 76/250  
- 26s - loss: 0.7382 - acc: 0.8874 - val\_loss: 0.7601 - val\_acc: 0.8826  
Epoch 77/250  
- 26s - loss: 0.7370 - acc: 0.8889 - val\_loss: 0.7198 - val\_acc: 0.8960  
Epoch 78/250  
- 26s - loss: 0.7356 - acc: 0.8878 - val\_loss: 0.7542 - val\_acc: 0.8824  
Epoch 79/250  
- 26s - loss: 0.7389 - acc: 0.8885 - val\_loss: 0.7284 - val\_acc: 0.8907  
Epoch 80/250  
- 26s - loss: 0.7369 - acc: 0.8888 - val\_loss: 0.7737 - val\_acc: 0.8822  
Epoch 81/250  
- 26s - loss: 0.6750 - acc: 0.9077 - val\_loss: 0.6911 - val\_acc: 0.9031  
Epoch 82/250  
- 26s - loss: 0.6451 - acc: 0.9129 - val\_loss: 0.6512 - val\_acc: 0.9099  
Epoch 83/250  
- 26s - loss: 0.6357 - acc: 0.9126 - val\_loss: 0.6334 - val\_acc: 0.9164  
Epoch 84/250  
- 26s - loss: 0.6195 - acc: 0.9181 - val\_loss: 0.6839 - val\_acc: 0.8996  
Epoch 85/250  
- 25s - loss: 0.6071 - acc: 0.9184 - val\_loss: 0.6758 - val\_acc: 0.9030  
Epoch 86/250  
- 26s - loss: 0.5996 - acc: 0.9196 - val\_loss: 0.6563 - val\_acc: 0.9052  
Epoch 87/250  
- 26s - loss: 0.5987 - acc: 0.9169 - val\_loss: 0.6187 - val\_acc: 0.9150  
Epoch 88/250  
- 26s - loss: 0.5895 - acc: 0.9195 - val\_loss: 0.6113 - val\_acc: 0.9150  
Epoch 89/250  
- 26s - loss: 0.5860 - acc: 0.9189 - val\_loss: 0.6247 - val\_acc: 0.9074

# Machine Learning Project Report

Epoch 90/250  
- 25s - loss: 0.5806 - acc: 0.9195 - val\_loss: 0.6087 - val\_acc: 0.9138  
Epoch 91/250  
- 26s - loss: 0.5747 - acc: 0.9202 - val\_loss: 0.6242 - val\_acc: 0.9068  
Epoch 92/250  
- 26s - loss: 0.5723 - acc: 0.9191 - val\_loss: 0.6404 - val\_acc: 0.9028  
Epoch 93/250  
- 26s - loss: 0.5714 - acc: 0.9194 - val\_loss: 0.6095 - val\_acc: 0.9104  
Epoch 94/250  
- 26s - loss: 0.5630 - acc: 0.9207 - val\_loss: 0.6242 - val\_acc: 0.9065  
Epoch 95/250  
- 26s - loss: 0.5643 - acc: 0.9197 - val\_loss: 0.6291 - val\_acc: 0.9071  
Epoch 96/250  
- 26s - loss: 0.5623 - acc: 0.9213 - val\_loss: 0.6162 - val\_acc: 0.9052  
Epoch 97/250  
- 26s - loss: 0.5656 - acc: 0.9184 - val\_loss: 0.6174 - val\_acc: 0.9077  
Epoch 98/250  
- 25s - loss: 0.5620 - acc: 0.9208 - val\_loss: 0.6310 - val\_acc: 0.9011  
Epoch 99/250  
- 25s - loss: 0.5592 - acc: 0.9206 - val\_loss: 0.6732 - val\_acc: 0.8890  
Epoch 100/250  
- 26s - loss: 0.5554 - acc: 0.9216 - val\_loss: 0.6391 - val\_acc: 0.9000  
Epoch 101/250  
- 26s - loss: 0.5239 - acc: 0.9307 - val\_loss: 0.5800 - val\_acc: 0.9174  
Epoch 102/250  
- 25s - loss: 0.5028 - acc: 0.9360 - val\_loss: 0.5999 - val\_acc: 0.9133  
Epoch 103/250  
- 25s - loss: 0.4956 - acc: 0.9377 - val\_loss: 0.5657 - val\_acc: 0.9188  
Epoch 104/250  
- 26s - loss: 0.4831 - acc: 0.9401 - val\_loss: 0.5611 - val\_acc: 0.9190  
Epoch 105/250  
- 25s - loss: 0.4764 - acc: 0.9411 - val\_loss: 0.5510 - val\_acc: 0.9227  
Epoch 106/250  
- 26s - loss: 0.4737 - acc: 0.9414 - val\_loss: 0.5818 - val\_acc: 0.9127  
Epoch 107/250  
- 26s - loss: 0.4633 - acc: 0.9433 - val\_loss: 0.5500 - val\_acc: 0.9225  
Epoch 108/250  
- 25s - loss: 0.4668 - acc: 0.9413 - val\_loss: 0.5684 - val\_acc: 0.9165  
Epoch 109/250  
- 26s - loss: 0.4597 - acc: 0.9424 - val\_loss: 0.5637 - val\_acc: 0.9190  
Epoch 110/250  
- 25s - loss: 0.4535 - acc: 0.9443 - val\_loss: 0.5521 - val\_acc: 0.9215  
Epoch 111/250  
- 25s - loss: 0.4519 - acc: 0.9442 - val\_loss: 0.5480 - val\_acc: 0.9193  
Epoch 112/250  
- 26s - loss: 0.4474 - acc: 0.9447 - val\_loss: 0.5420 - val\_acc: 0.9216  
Epoch 113/250  
- 26s - loss: 0.4472 - acc: 0.9432 - val\_loss: 0.5465 - val\_acc: 0.9199  
Epoch 114/250  
- 25s - loss: 0.4410 - acc: 0.9448 - val\_loss: 0.5360 - val\_acc: 0.9240  
Epoch 115/250  
- 25s - loss: 0.4375 - acc: 0.9455 - val\_loss: 0.5677 - val\_acc: 0.9116

# Machine Learning Project Report

Epoch 116/250  
- 26s - loss: 0.4354 - acc: 0.9444 - val\_loss: 0.5392 - val\_acc: 0.9197  
Epoch 117/250  
- 25s - loss: 0.4373 - acc: 0.9456 - val\_loss: 0.5536 - val\_acc: 0.9142  
Epoch 118/250  
- 25s - loss: 0.4344 - acc: 0.9452 - val\_loss: 0.5414 - val\_acc: 0.9177  
Epoch 119/250  
- 26s - loss: 0.4311 - acc: 0.9433 - val\_loss: 0.5358 - val\_acc: 0.9196  
Epoch 120/250  
- 25s - loss: 0.4285 - acc: 0.9462 - val\_loss: 0.5566 - val\_acc: 0.9124  
Epoch 121/250  
- 26s - loss: 0.4106 - acc: 0.9495 - val\_loss: 0.5175 - val\_acc: 0.9231  
Epoch 122/250  
- 26s - loss: 0.3942 - acc: 0.9546 - val\_loss: 0.5628 - val\_acc: 0.9144  
Epoch 123/250  
- 25s - loss: 0.3891 - acc: 0.9579 - val\_loss: 0.5226 - val\_acc: 0.9243  
Epoch 124/250  
- 25s - loss: 0.3846 - acc: 0.9583 - val\_loss: 0.5164 - val\_acc: 0.9254  
Epoch 125/250  
- 26s - loss: 0.3813 - acc: 0.9594 - val\_loss: 0.5022 - val\_acc: 0.9263  
Epoch 126/250  
- 25s - loss: 0.3807 - acc: 0.9580 - val\_loss: 0.5138 - val\_acc: 0.9242  
Epoch 127/250  
- 26s - loss: 0.3722 - acc: 0.9586 - val\_loss: 0.5083 - val\_acc: 0.9257  
Epoch 128/250  
- 26s - loss: 0.3705 - acc: 0.9610 - val\_loss: 0.5180 - val\_acc: 0.9240  
Epoch 129/250  
- 25s - loss: 0.3689 - acc: 0.9596 - val\_loss: 0.5063 - val\_acc: 0.9273  
Epoch 130/250  
- 26s - loss: 0.3678 - acc: 0.9596 - val\_loss: 0.5078 - val\_acc: 0.9262  
Epoch 131/250  
- 26s - loss: 0.3609 - acc: 0.9618 - val\_loss: 0.5176 - val\_acc: 0.9221  
Epoch 132/250  
- 25s - loss: 0.3613 - acc: 0.9607 - val\_loss: 0.4954 - val\_acc: 0.9284  
Epoch 133/250  
- 26s - loss: 0.3558 - acc: 0.9619 - val\_loss: 0.4997 - val\_acc: 0.9249  
Epoch 134/250  
- 25s - loss: 0.3506 - acc: 0.9634 - val\_loss: 0.4969 - val\_acc: 0.9275  
Epoch 135/250  
- 26s - loss: 0.3528 - acc: 0.9622 - val\_loss: 0.5027 - val\_acc: 0.9254  
Epoch 136/250  
- 25s - loss: 0.3465 - acc: 0.9633 - val\_loss: 0.5257 - val\_acc: 0.9215  
Epoch 137/250  
- 26s - loss: 0.3488 - acc: 0.9624 - val\_loss: 0.5019 - val\_acc: 0.9255  
Epoch 138/250  
- 26s - loss: 0.3495 - acc: 0.9619 - val\_loss: 0.4962 - val\_acc: 0.9277  
Epoch 139/250  
- 26s - loss: 0.3449 - acc: 0.9626 - val\_loss: 0.4905 - val\_acc: 0.9289  
Epoch 140/250  
- 26s - loss: 0.3405 - acc: 0.9633 - val\_loss: 0.4916 - val\_acc: 0.9269  
Epoch 141/250  
- 26s - loss: 0.3401 - acc: 0.9635 - val\_loss: 0.4860 - val\_acc: 0.9296

# Machine Learning Project Report

Epoch 142/250  
- 26s - loss: 0.3299 - acc: 0.9664 - val\_loss: 0.4716 - val\_acc: 0.9295  
Epoch 143/250  
- 25s - loss: 0.3223 - acc: 0.9686 - val\_loss: 0.4749 - val\_acc: 0.9307  
Epoch 144/250  
- 26s - loss: 0.3192 - acc: 0.9697 - val\_loss: 0.4854 - val\_acc: 0.9269  
Epoch 145/250  
- 25s - loss: 0.3202 - acc: 0.9696 - val\_loss: 0.4780 - val\_acc: 0.9286  
Epoch 146/250  
- 25s - loss: 0.3164 - acc: 0.9695 - val\_loss: 0.4860 - val\_acc: 0.9276  
Epoch 147/250  
- 26s - loss: 0.3138 - acc: 0.9706 - val\_loss: 0.4753 - val\_acc: 0.9313  
Epoch 148/250  
- 25s - loss: 0.3116 - acc: 0.9703 - val\_loss: 0.4784 - val\_acc: 0.9302  
Epoch 149/250  
- 26s - loss: 0.3108 - acc: 0.9707 - val\_loss: 0.4736 - val\_acc: 0.9322  
Epoch 150/250  
- 26s - loss: 0.3097 - acc: 0.9708 - val\_loss: 0.4760 - val\_acc: 0.9308  
Epoch 151/250  
- 26s - loss: 0.3093 - acc: 0.9708 - val\_loss: 0.4734 - val\_acc: 0.9287  
Epoch 152/250  
- 25s - loss: 0.3049 - acc: 0.9719 - val\_loss: 0.4751 - val\_acc: 0.9293  
Epoch 153/250  
- 25s - loss: 0.3047 - acc: 0.9714 - val\_loss: 0.4758 - val\_acc: 0.9300  
Epoch 154/250  
- 25s - loss: 0.3032 - acc: 0.9713 - val\_loss: 0.4820 - val\_acc: 0.9280  
Epoch 155/250  
- 25s - loss: 0.3028 - acc: 0.9716 - val\_loss: 0.4759 - val\_acc: 0.9276  
Epoch 156/250  
- 26s - loss: 0.2977 - acc: 0.9727 - val\_loss: 0.4707 - val\_acc: 0.9311  
Epoch 157/250  
- 26s - loss: 0.3000 - acc: 0.9720 - val\_loss: 0.4855 - val\_acc: 0.9276  
Epoch 158/250  
- 25s - loss: 0.2943 - acc: 0.9736 - val\_loss: 0.4693 - val\_acc: 0.9323  
Epoch 159/250  
- 25s - loss: 0.2961 - acc: 0.9722 - val\_loss: 0.4684 - val\_acc: 0.9310  
Epoch 160/250  
- 26s - loss: 0.2929 - acc: 0.9735 - val\_loss: 0.4822 - val\_acc: 0.9275  
Epoch 161/250  
- 26s - loss: 0.2876 - acc: 0.9747 - val\_loss: 0.4760 - val\_acc: 0.9286  
Epoch 162/250  
- 26s - loss: 0.2847 - acc: 0.9751 - val\_loss: 0.4732 - val\_acc: 0.9299  
Epoch 163/250  
- 26s - loss: 0.2805 - acc: 0.9770 - val\_loss: 0.4789 - val\_acc: 0.9277  
Epoch 164/250  
- 26s - loss: 0.2812 - acc: 0.9764 - val\_loss: 0.4702 - val\_acc: 0.9294  
Epoch 165/250  
- 26s - loss: 0.2819 - acc: 0.9756 - val\_loss: 0.4788 - val\_acc: 0.9289  
Epoch 166/250  
- 26s - loss: 0.2791 - acc: 0.9769 - val\_loss: 0.4671 - val\_acc: 0.9320  
Epoch 167/250  
- 25s - loss: 0.2774 - acc: 0.9780 - val\_loss: 0.4773 - val\_acc: 0.9292

# Machine Learning Project Report

Epoch 168/250  
- 26s - loss: 0.2786 - acc: 0.9769 - val\_loss: 0.4694 - val\_acc: 0.9306  
Epoch 169/250  
- 26s - loss: 0.2748 - acc: 0.9776 - val\_loss: 0.4665 - val\_acc: 0.9319  
Epoch 170/250  
- 25s - loss: 0.2770 - acc: 0.9768 - val\_loss: 0.4680 - val\_acc: 0.9309  
Epoch 171/250  
- 25s - loss: 0.2762 - acc: 0.9765 - val\_loss: 0.4721 - val\_acc: 0.9302  
Epoch 172/250  
- 26s - loss: 0.2764 - acc: 0.9771 - val\_loss: 0.4824 - val\_acc: 0.9282  
Epoch 173/250  
- 25s - loss: 0.2727 - acc: 0.9783 - val\_loss: 0.4733 - val\_acc: 0.9295  
Epoch 174/250  
- 25s - loss: 0.2730 - acc: 0.9778 - val\_loss: 0.4712 - val\_acc: 0.9308  
Epoch 175/250  
- 26s - loss: 0.2687 - acc: 0.9784 - val\_loss: 0.4782 - val\_acc: 0.9280  
Epoch 176/250  
- 25s - loss: 0.2717 - acc: 0.9773 - val\_loss: 0.4701 - val\_acc: 0.9305  
Epoch 177/250  
- 25s - loss: 0.2719 - acc: 0.9772 - val\_loss: 0.4765 - val\_acc: 0.9291  
Epoch 178/250  
- 25s - loss: 0.2687 - acc: 0.9779 - val\_loss: 0.4700 - val\_acc: 0.9318  
Epoch 179/250  
- 25s - loss: 0.2666 - acc: 0.9799 - val\_loss: 0.4632 - val\_acc: 0.9329  
Epoch 180/250  
- 25s - loss: 0.2678 - acc: 0.9790 - val\_loss: 0.4699 - val\_acc: 0.9315  
Epoch 181/250  
- 25s - loss: 0.2635 - acc: 0.9795 - val\_loss: 0.4698 - val\_acc: 0.9310  
Epoch 182/250  
- 26s - loss: 0.2633 - acc: 0.9798 - val\_loss: 0.4681 - val\_acc: 0.9322  
Epoch 183/250  
- 25s - loss: 0.2619 - acc: 0.9800 - val\_loss: 0.4704 - val\_acc: 0.9303  
Epoch 184/250  
- 25s - loss: 0.2606 - acc: 0.9810 - val\_loss: 0.4741 - val\_acc: 0.9304  
Epoch 185/250  
- 25s - loss: 0.2591 - acc: 0.9814 - val\_loss: 0.4680 - val\_acc: 0.9315  
Epoch 186/250  
- 26s - loss: 0.2606 - acc: 0.9805 - val\_loss: 0.4640 - val\_acc: 0.9341  
Epoch 187/250  
- 25s - loss: 0.2586 - acc: 0.9806 - val\_loss: 0.4677 - val\_acc: 0.9318  
Epoch 188/250  
- 26s - loss: 0.2569 - acc: 0.9814 - val\_loss: 0.4731 - val\_acc: 0.9304  
Epoch 189/250  
- 26s - loss: 0.2611 - acc: 0.9798 - val\_loss: 0.4702 - val\_acc: 0.9321  
Epoch 190/250  
- 25s - loss: 0.2612 - acc: 0.9804 - val\_loss: 0.4757 - val\_acc: 0.9316  
Epoch 191/250  
- 26s - loss: 0.2553 - acc: 0.9818 - val\_loss: 0.4702 - val\_acc: 0.9326  
Epoch 192/250  
- 25s - loss: 0.2568 - acc: 0.9809 - val\_loss: 0.4679 - val\_acc: 0.9324  
Epoch 193/250  
- 25s - loss: 0.2589 - acc: 0.9801 - val\_loss: 0.4671 - val\_acc: 0.9320

# Machine Learning Project Report

Epoch 194/250  
- 26s - loss: 0.2569 - acc: 0.9812 - val\_loss: 0.4679 - val\_acc: 0.9314  
Epoch 195/250  
- 25s - loss: 0.2532 - acc: 0.9820 - val\_loss: 0.4621 - val\_acc: 0.9336  
Epoch 196/250  
- 25s - loss: 0.2567 - acc: 0.9807 - val\_loss: 0.4604 - val\_acc: 0.9336  
Epoch 197/250  
- 25s - loss: 0.2551 - acc: 0.9815 - val\_loss: 0.4676 - val\_acc: 0.9327  
Epoch 198/250  
- 26s - loss: 0.2542 - acc: 0.9808 - val\_loss: 0.4622 - val\_acc: 0.9341  
Epoch 199/250  
- 25s - loss: 0.2527 - acc: 0.9818 - val\_loss: 0.4676 - val\_acc: 0.9329  
Epoch 200/250  
- 26s - loss: 0.2529 - acc: 0.9817 - val\_loss: 0.4721 - val\_acc: 0.9315  
Epoch 201/250  
- 26s - loss: 0.2533 - acc: 0.9810 - val\_loss: 0.4696 - val\_acc: 0.9314  
Epoch 202/250  
- 25s - loss: 0.2530 - acc: 0.9814 - val\_loss: 0.4658 - val\_acc: 0.9326  
Epoch 203/250  
- 25s - loss: 0.2499 - acc: 0.9824 - val\_loss: 0.4663 - val\_acc: 0.9328  
Epoch 204/250  
- 26s - loss: 0.2515 - acc: 0.9828 - val\_loss: 0.4670 - val\_acc: 0.9326  
Epoch 205/250  
- 25s - loss: 0.2498 - acc: 0.9820 - val\_loss: 0.4657 - val\_acc: 0.9330  
Epoch 206/250  
- 25s - loss: 0.2504 - acc: 0.9822 - val\_loss: 0.4650 - val\_acc: 0.9323  
Epoch 207/250  
- 26s - loss: 0.2490 - acc: 0.9823 - val\_loss: 0.4611 - val\_acc: 0.9322  
Epoch 208/250  
- 25s - loss: 0.2466 - acc: 0.9831 - val\_loss: 0.4653 - val\_acc: 0.9317  
Epoch 209/250  
- 25s - loss: 0.2486 - acc: 0.9825 - val\_loss: 0.4643 - val\_acc: 0.9318  
Epoch 210/250  
- 26s - loss: 0.2480 - acc: 0.9829 - val\_loss: 0.4664 - val\_acc: 0.9314  
Epoch 211/250  
- 26s - loss: 0.2472 - acc: 0.9833 - val\_loss: 0.4683 - val\_acc: 0.9319  
Epoch 212/250  
- 25s - loss: 0.2507 - acc: 0.9812 - val\_loss: 0.4648 - val\_acc: 0.9340  
Epoch 213/250  
- 26s - loss: 0.2495 - acc: 0.9818 - val\_loss: 0.4658 - val\_acc: 0.9316  
Epoch 214/250  
- 25s - loss: 0.2469 - acc: 0.9833 - val\_loss: 0.4672 - val\_acc: 0.9321  
Epoch 215/250  
- 25s - loss: 0.2495 - acc: 0.9821 - val\_loss: 0.4644 - val\_acc: 0.9323  
Epoch 216/250  
- 25s - loss: 0.2486 - acc: 0.9828 - val\_loss: 0.4632 - val\_acc: 0.9328  
Epoch 217/250  
- 25s - loss: 0.2487 - acc: 0.9824 - val\_loss: 0.4655 - val\_acc: 0.9322  
Epoch 218/250  
- 25s - loss: 0.2466 - acc: 0.9828 - val\_loss: 0.4671 - val\_acc: 0.9317  
Epoch 219/250  
- 25s - loss: 0.2480 - acc: 0.9826 - val\_loss: 0.4621 - val\_acc: 0.9332

# Machine Learning Project Report

Epoch 220/250  
- 25s - loss: 0.2467 - acc: 0.9836 - val\_loss: 0.4614 - val\_acc: 0.9323  
Epoch 221/250  
- 25s - loss: 0.2451 - acc: 0.9831 - val\_loss: 0.4609 - val\_acc: 0.9326  
Epoch 222/250  
- 26s - loss: 0.2470 - acc: 0.9828 - val\_loss: 0.4605 - val\_acc: 0.9332  
Epoch 223/250  
- 26s - loss: 0.2449 - acc: 0.9841 - val\_loss: 0.4606 - val\_acc: 0.9338  
Epoch 224/250  
- 25s - loss: 0.2479 - acc: 0.9825 - val\_loss: 0.4619 - val\_acc: 0.9332  
Epoch 225/250  
- 25s - loss: 0.2441 - acc: 0.9829 - val\_loss: 0.4597 - val\_acc: 0.9340  
Epoch 226/250  
- 25s - loss: 0.2442 - acc: 0.9833 - val\_loss: 0.4601 - val\_acc: 0.9330  
Epoch 227/250  
- 25s - loss: 0.2465 - acc: 0.9819 - val\_loss: 0.4608 - val\_acc: 0.9331  
Epoch 228/250  
- 25s - loss: 0.2441 - acc: 0.9835 - val\_loss: 0.4598 - val\_acc: 0.9330  
Epoch 229/250  
- 26s - loss: 0.2415 - acc: 0.9840 - val\_loss: 0.4625 - val\_acc: 0.9326  
Epoch 230/250  
- 25s - loss: 0.2446 - acc: 0.9832 - val\_loss: 0.4588 - val\_acc: 0.9336  
Epoch 231/250  
- 25s - loss: 0.2457 - acc: 0.9829 - val\_loss: 0.4593 - val\_acc: 0.9333  
Epoch 232/250  
- 25s - loss: 0.2440 - acc: 0.9835 - val\_loss: 0.4620 - val\_acc: 0.9337  
Epoch 233/250  
- 25s - loss: 0.2468 - acc: 0.9823 - val\_loss: 0.4601 - val\_acc: 0.9331  
Epoch 234/250  
- 26s - loss: 0.2456 - acc: 0.9829 - val\_loss: 0.4606 - val\_acc: 0.9338  
Epoch 235/250  
- 26s - loss: 0.2440 - acc: 0.9835 - val\_loss: 0.4611 - val\_acc: 0.9342  
Epoch 236/250  
- 25s - loss: 0.2428 - acc: 0.9843 - val\_loss: 0.4615 - val\_acc: 0.9339  
Epoch 237/250  
- 25s - loss: 0.2434 - acc: 0.9830 - val\_loss: 0.4606 - val\_acc: 0.9338  
Epoch 238/250  
- 25s - loss: 0.2452 - acc: 0.9835 - val\_loss: 0.4595 - val\_acc: 0.9336  
Epoch 239/250  
- 26s - loss: 0.2420 - acc: 0.9832 - val\_loss: 0.4589 - val\_acc: 0.9332  
Epoch 240/250  
- 26s - loss: 0.2429 - acc: 0.9839 - val\_loss: 0.4609 - val\_acc: 0.9332  
Epoch 241/250  
- 26s - loss: 0.2434 - acc: 0.9834 - val\_loss: 0.4595 - val\_acc: 0.9337  
Epoch 242/250  
- 26s - loss: 0.2427 - acc: 0.9832 - val\_loss: 0.4596 - val\_acc: 0.9337  
Epoch 243/250  
- 26s - loss: 0.2432 - acc: 0.9832 - val\_loss: 0.4602 - val\_acc: 0.9337  
Epoch 244/250  
- 25s - loss: 0.2409 - acc: 0.9847 - val\_loss: 0.4612 - val\_acc: 0.9336  
Epoch 245/250  
- 25s - loss: 0.2430 - acc: 0.9835 - val\_loss: 0.4595 - val\_acc: 0.9339



# Machine Learning Project Report

Epoch 246/250  
- 25s - loss: 0.2448 - acc: 0.9827 - val\_loss: 0.4597 - val\_acc: 0.9340  
Epoch 247/250  
- 26s - loss: 0.2445 - acc: 0.9832 - val\_loss: 0.4589 - val\_acc: 0.9344  
Epoch 248/250  
- 26s - loss: 0.2424 - acc: 0.9835 - val\_loss: 0.4581 - val\_acc: 0.9344  
Epoch 249/250  
- 25s - loss: 0.2408 - acc: 0.9848 - val\_loss: 0.4596 - val\_acc: 0.9338  
Epoch 250/250  
- 25s - loss: 0.2418 - acc: 0.9836 - val\_loss: 0.4583 - val\_acc: 0.9345  
Training Time: 6452.848551273346  
Accuracy: 93.45 %

## Second Run:

Epoch 1/250  
- 33s - loss: 22.0208 - acc: 0.1882 - val\_loss: 17.7246 - val\_acc: 0.1387  
Epoch 2/250  
- 27s - loss: 11.6296 - acc: 0.2733 - val\_loss: 9.3304 - val\_acc: 0.1512  
Epoch 3/250  
- 27s - loss: 7.3973 - acc: 0.2360 - val\_loss: 7.8169 - val\_acc: 0.1271  
Epoch 4/250  
- 27s - loss: 6.6897 - acc: 0.2450 - val\_loss: 5.8401 - val\_acc: 0.2071  
Epoch 5/250  
- 27s - loss: 4.8519 - acc: 0.3264 - val\_loss: 4.3972 - val\_acc: 0.2241  
Epoch 6/250  
- 27s - loss: 3.2290 - acc: 0.4056 - val\_loss: 3.0917 - val\_acc: 0.2920  
Epoch 7/250  
- 27s - loss: 2.3470 - acc: 0.4721 - val\_loss: 2.2098 - val\_acc: 0.4187  
Epoch 8/250  
- 27s - loss: 1.9149 - acc: 0.5345 - val\_loss: 2.1548 - val\_acc: 0.4675  
Epoch 9/250  
- 27s - loss: 1.7388 - acc: 0.5777 - val\_loss: 1.7121 - val\_acc: 0.5983  
Epoch 10/250  
- 27s - loss: 1.6145 - acc: 0.6198 - val\_loss: 1.5373 - val\_acc: 0.6542  
Epoch 11/250  
- 27s - loss: 1.5361 - acc: 0.6504 - val\_loss: 1.5947 - val\_acc: 0.6437  
Epoch 12/250  
- 27s - loss: 1.5094 - acc: 0.6720 - val\_loss: 1.3756 - val\_acc: 0.7259  
Epoch 13/250  
- 27s - loss: 1.4905 - acc: 0.6846 - val\_loss: 1.4603 - val\_acc: 0.6914  
Epoch 14/250  
- 26s - loss: 1.4783 - acc: 0.6960 - val\_loss: 1.9989 - val\_acc: 0.5897  
Epoch 15/250  
- 27s - loss: 1.4834 - acc: 0.7035 - val\_loss: 1.4743 - val\_acc: 0.7181  
Epoch 16/250  
- 26s - loss: 1.4790 - acc: 0.7106 - val\_loss: 1.5071 - val\_acc: 0.6970  
Epoch 17/250  
- 27s - loss: 1.4862 - acc: 0.7169 - val\_loss: 1.6156 - val\_acc: 0.6750  
Epoch 18/250

# Machine Learning Project Report

```
- 27s - loss: 1.4895 - acc: 0.7224 - val_loss: 1.5565 - val_acc: 0.7065
Epoch 19/250
- 27s - loss: 1.4953 - acc: 0.7246 - val_loss: 1.4234 - val_acc: 0.7500
Epoch 20/250
- 27s - loss: 1.5010 - acc: 0.7275 - val_loss: 1.4648 - val_acc: 0.7468
Epoch 21/250
- 27s - loss: 1.3387 - acc: 0.7669 - val_loss: 1.1798 - val_acc: 0.8017
Epoch 22/250
- 26s - loss: 1.2430 - acc: 0.7759 - val_loss: 1.1332 - val_acc: 0.8053
Epoch 23/250
- 26s - loss: 1.2199 - acc: 0.7757 - val_loss: 1.1164 - val_acc: 0.8076
Epoch 24/250
- 26s - loss: 1.2170 - acc: 0.7747 - val_loss: 1.2335 - val_acc: 0.7720
Epoch 25/250
- 26s - loss: 1.2014 - acc: 0.7827 - val_loss: 1.1221 - val_acc: 0.8073
Epoch 26/250
- 26s - loss: 1.2125 - acc: 0.7801 - val_loss: 1.1337 - val_acc: 0.8107
Epoch 27/250
- 26s - loss: 1.2075 - acc: 0.7859 - val_loss: 1.3345 - val_acc: 0.7494
Epoch 28/250
- 26s - loss: 1.2113 - acc: 0.7860 - val_loss: 1.0902 - val_acc: 0.8245
Epoch 29/250
- 26s - loss: 1.2193 - acc: 0.7895 - val_loss: 1.2040 - val_acc: 0.7898
Epoch 30/250
- 26s - loss: 1.2219 - acc: 0.7897 - val_loss: 1.1523 - val_acc: 0.8075
Epoch 31/250
- 26s - loss: 1.2264 - acc: 0.7921 - val_loss: 1.1140 - val_acc: 0.8289
Epoch 32/250
- 26s - loss: 1.2302 - acc: 0.7915 - val_loss: 1.1519 - val_acc: 0.8213
Epoch 33/250
- 26s - loss: 1.2333 - acc: 0.7936 - val_loss: 1.2256 - val_acc: 0.7942
Epoch 34/250
- 26s - loss: 1.2299 - acc: 0.7962 - val_loss: 1.2050 - val_acc: 0.8071
Epoch 35/250
- 26s - loss: 1.2319 - acc: 0.7961 - val_loss: 1.1674 - val_acc: 0.8164
Epoch 36/250
- 26s - loss: 1.2352 - acc: 0.7979 - val_loss: 1.2697 - val_acc: 0.7917
Epoch 37/250
- 26s - loss: 1.2325 - acc: 0.7996 - val_loss: 1.1997 - val_acc: 0.8044
Epoch 38/250
- 26s - loss: 1.2316 - acc: 0.8006 - val_loss: 1.3241 - val_acc: 0.7739
Epoch 39/250
- 26s - loss: 1.2484 - acc: 0.7980 - val_loss: 1.2515 - val_acc: 0.7996
Epoch 40/250
- 26s - loss: 1.2393 - acc: 0.7986 - val_loss: 1.2109 - val_acc: 0.8139
Epoch 41/250
- 26s - loss: 1.1157 - acc: 0.8336 - val_loss: 0.9886 - val_acc: 0.8619
Epoch 42/250
- 26s - loss: 1.0401 - acc: 0.8434 - val_loss: 0.9573 - val_acc: 0.8616
Epoch 43/250
- 26s - loss: 1.0015 - acc: 0.8440 - val_loss: 0.9486 - val_acc: 0.8586
Epoch 44/250
```

# Machine Learning Project Report

```
- 26s - loss: 0.9900 - acc: 0.8423 - val_loss: 0.9344 - val_acc: 0.8599
Epoch 45/250
- 26s - loss: 0.9716 - acc: 0.8453 - val_loss: 0.9060 - val_acc: 0.8627
Epoch 46/250
- 26s - loss: 0.9686 - acc: 0.8424 - val_loss: 0.9383 - val_acc: 0.8555
Epoch 47/250
- 26s - loss: 0.9704 - acc: 0.8420 - val_loss: 1.0192 - val_acc: 0.8267
Epoch 48/250
- 27s - loss: 0.9670 - acc: 0.8406 - val_loss: 0.9301 - val_acc: 0.8570
Epoch 49/250
- 26s - loss: 0.9607 - acc: 0.8439 - val_loss: 0.9553 - val_acc: 0.8505
Epoch 50/250
- 26s - loss: 0.9667 - acc: 0.8451 - val_loss: 0.9154 - val_acc: 0.8599
Epoch 51/250
- 26s - loss: 0.9719 - acc: 0.8429 - val_loss: 0.9046 - val_acc: 0.8654
Epoch 52/250
- 26s - loss: 0.9627 - acc: 0.8474 - val_loss: 0.9515 - val_acc: 0.8478
Epoch 53/250
- 26s - loss: 0.9765 - acc: 0.8415 - val_loss: 0.9314 - val_acc: 0.8606
Epoch 54/250
- 26s - loss: 0.9711 - acc: 0.8463 - val_loss: 0.9683 - val_acc: 0.8473
Epoch 55/250
- 26s - loss: 0.9656 - acc: 0.8481 - val_loss: 0.9260 - val_acc: 0.8636
Epoch 56/250
- 26s - loss: 0.9782 - acc: 0.8457 - val_loss: 1.0713 - val_acc: 0.8203
Epoch 57/250
- 26s - loss: 0.9768 - acc: 0.8472 - val_loss: 0.9433 - val_acc: 0.8618
Epoch 58/250
- 26s - loss: 0.9708 - acc: 0.8502 - val_loss: 0.9161 - val_acc: 0.8663
Epoch 59/250
- 26s - loss: 0.9766 - acc: 0.8516 - val_loss: 0.9366 - val_acc: 0.8618
Epoch 60/250
- 26s - loss: 0.9737 - acc: 0.8495 - val_loss: 0.9541 - val_acc: 0.8545
Epoch 61/250
- 26s - loss: 0.8840 - acc: 0.8747 - val_loss: 0.8103 - val_acc: 0.8949
Epoch 62/250
- 26s - loss: 0.8353 - acc: 0.8841 - val_loss: 0.8033 - val_acc: 0.8909
Epoch 63/250
- 26s - loss: 0.8157 - acc: 0.8828 - val_loss: 0.7751 - val_acc: 0.8915
Epoch 64/250
- 26s - loss: 0.7914 - acc: 0.8866 - val_loss: 0.8417 - val_acc: 0.8669
Epoch 65/250
- 26s - loss: 0.7836 - acc: 0.8851 - val_loss: 0.7633 - val_acc: 0.8927
Epoch 66/250
- 26s - loss: 0.7726 - acc: 0.8850 - val_loss: 0.7650 - val_acc: 0.8898
Epoch 67/250
- 26s - loss: 0.7679 - acc: 0.8854 - val_loss: 0.7404 - val_acc: 0.8919
Epoch 68/250
- 26s - loss: 0.7622 - acc: 0.8839 - val_loss: 0.7063 - val_acc: 0.9002
Epoch 69/250
- 26s - loss: 0.7564 - acc: 0.8853 - val_loss: 0.7528 - val_acc: 0.8868
Epoch 70/250
```

# Machine Learning Project Report

```
- 26s - loss: 0.7533 - acc: 0.8842 - val_loss: 0.7876 - val_acc: 0.8793
Epoch 71/250
- 26s - loss: 0.7518 - acc: 0.8841 - val_loss: 0.7208 - val_acc: 0.8962
Epoch 72/250
- 26s - loss: 0.7549 - acc: 0.8834 - val_loss: 0.7602 - val_acc: 0.8821
Epoch 73/250
- 26s - loss: 0.7470 - acc: 0.8849 - val_loss: 0.7281 - val_acc: 0.8913
Epoch 74/250
- 26s - loss: 0.7501 - acc: 0.8847 - val_loss: 0.7463 - val_acc: 0.8854
Epoch 75/250
- 26s - loss: 0.7521 - acc: 0.8853 - val_loss: 0.7446 - val_acc: 0.8905
Epoch 76/250
- 26s - loss: 0.7499 - acc: 0.8836 - val_loss: 0.7261 - val_acc: 0.8946
Epoch 77/250
- 26s - loss: 0.7513 - acc: 0.8842 - val_loss: 0.7320 - val_acc: 0.8907
Epoch 78/250
- 26s - loss: 0.7417 - acc: 0.8888 - val_loss: 0.7554 - val_acc: 0.8899
Epoch 79/250
- 26s - loss: 0.7454 - acc: 0.8872 - val_loss: 0.7454 - val_acc: 0.8896
Epoch 80/250
- 26s - loss: 0.7472 - acc: 0.8857 - val_loss: 0.7414 - val_acc: 0.8911
Epoch 81/250
- 26s - loss: 0.6916 - acc: 0.9028 - val_loss: 0.6760 - val_acc: 0.9094
Epoch 82/250
- 26s - loss: 0.6534 - acc: 0.9129 - val_loss: 0.6549 - val_acc: 0.9113
Epoch 83/250
- 26s - loss: 0.6351 - acc: 0.9159 - val_loss: 0.6543 - val_acc: 0.9103
Epoch 84/250
- 26s - loss: 0.6276 - acc: 0.9162 - val_loss: 0.6619 - val_acc: 0.9070
Epoch 85/250
- 26s - loss: 0.6126 - acc: 0.9171 - val_loss: 0.6453 - val_acc: 0.9106
Epoch 86/250
- 26s - loss: 0.6166 - acc: 0.9150 - val_loss: 0.6259 - val_acc: 0.9149
Epoch 87/250
- 26s - loss: 0.6057 - acc: 0.9169 - val_loss: 0.6698 - val_acc: 0.8995
Epoch 88/250
- 25s - loss: 0.5998 - acc: 0.9164 - val_loss: 0.6395 - val_acc: 0.9073
Epoch 89/250
- 26s - loss: 0.5921 - acc: 0.9175 - val_loss: 0.6370 - val_acc: 0.9021
Epoch 90/250
- 26s - loss: 0.5936 - acc: 0.9165 - val_loss: 0.6307 - val_acc: 0.9075
Epoch 91/250
- 26s - loss: 0.5875 - acc: 0.9186 - val_loss: 0.6376 - val_acc: 0.9033
Epoch 92/250
- 26s - loss: 0.5807 - acc: 0.9177 - val_loss: 0.6141 - val_acc: 0.9131
Epoch 93/250
- 26s - loss: 0.5762 - acc: 0.9191 - val_loss: 0.6566 - val_acc: 0.9023
Epoch 94/250
- 26s - loss: 0.5692 - acc: 0.9214 - val_loss: 0.6255 - val_acc: 0.9070
Epoch 95/250
- 26s - loss: 0.5778 - acc: 0.9176 - val_loss: 0.5972 - val_acc: 0.9147
Epoch 96/250
```

# Machine Learning Project Report

```
- 26s - loss: 0.5755 - acc: 0.9164 - val_loss: 0.6188 - val_acc: 0.9069
Epoch 97/250
- 26s - loss: 0.5769 - acc: 0.9150 - val_loss: 0.6653 - val_acc: 0.8957
Epoch 98/250
- 26s - loss: 0.5716 - acc: 0.9177 - val_loss: 0.6448 - val_acc: 0.9003
Epoch 99/250
- 26s - loss: 0.5717 - acc: 0.9178 - val_loss: 0.6104 - val_acc: 0.9102
Epoch 100/250
- 26s - loss: 0.5681 - acc: 0.9169 - val_loss: 0.6218 - val_acc: 0.9035
Epoch 101/250
- 26s - loss: 0.5340 - acc: 0.9287 - val_loss: 0.6074 - val_acc: 0.9098
Epoch 102/250
- 26s - loss: 0.5060 - acc: 0.9377 - val_loss: 0.5766 - val_acc: 0.9188
Epoch 103/250
- 26s - loss: 0.4995 - acc: 0.9377 - val_loss: 0.5781 - val_acc: 0.9185
Epoch 104/250
- 26s - loss: 0.4929 - acc: 0.9390 - val_loss: 0.5655 - val_acc: 0.9209
Epoch 105/250
- 26s - loss: 0.4867 - acc: 0.9402 - val_loss: 0.5773 - val_acc: 0.9163
Epoch 106/250
- 26s - loss: 0.4805 - acc: 0.9387 - val_loss: 0.5697 - val_acc: 0.9177
Epoch 107/250
- 26s - loss: 0.4763 - acc: 0.9409 - val_loss: 0.5657 - val_acc: 0.9176
Epoch 108/250
- 26s - loss: 0.4736 - acc: 0.9396 - val_loss: 0.5695 - val_acc: 0.9180
Epoch 109/250
- 26s - loss: 0.4666 - acc: 0.9428 - val_loss: 0.5516 - val_acc: 0.9192
Epoch 110/250
- 26s - loss: 0.4623 - acc: 0.9428 - val_loss: 0.5606 - val_acc: 0.9170
Epoch 111/250
- 26s - loss: 0.4584 - acc: 0.9435 - val_loss: 0.5482 - val_acc: 0.9231
Epoch 112/250
- 26s - loss: 0.4586 - acc: 0.9423 - val_loss: 0.5682 - val_acc: 0.9157
Epoch 113/250
- 26s - loss: 0.4494 - acc: 0.9452 - val_loss: 0.5403 - val_acc: 0.9234
Epoch 114/250
- 26s - loss: 0.4478 - acc: 0.9439 - val_loss: 0.5379 - val_acc: 0.9236
Epoch 115/250
- 26s - loss: 0.4459 - acc: 0.9440 - val_loss: 0.5756 - val_acc: 0.9123
Epoch 116/250
- 26s - loss: 0.4459 - acc: 0.9435 - val_loss: 0.5547 - val_acc: 0.9160
Epoch 117/250
- 26s - loss: 0.4441 - acc: 0.9428 - val_loss: 0.5339 - val_acc: 0.9215
Epoch 118/250
- 26s - loss: 0.4402 - acc: 0.9452 - val_loss: 0.5268 - val_acc: 0.9238
Epoch 119/250
- 26s - loss: 0.4357 - acc: 0.9444 - val_loss: 0.5484 - val_acc: 0.9199
Epoch 120/250
- 26s - loss: 0.4397 - acc: 0.9430 - val_loss: 0.5544 - val_acc: 0.9166
Epoch 121/250
- 26s - loss: 0.4164 - acc: 0.9501 - val_loss: 0.5400 - val_acc: 0.9182
Epoch 122/250
```

# Machine Learning Project Report

- 26s - loss: 0.4009 - acc: 0.9545 - val\_loss: 0.5308 - val\_acc: 0.9226  
Epoch 123/250  
- 26s - loss: 0.3954 - acc: 0.9562 - val\_loss: 0.5214 - val\_acc: 0.9256  
Epoch 124/250  
- 26s - loss: 0.3903 - acc: 0.9556 - val\_loss: 0.5198 - val\_acc: 0.9278  
Epoch 125/250  
- 25s - loss: 0.3876 - acc: 0.9572 - val\_loss: 0.5141 - val\_acc: 0.9265  
Epoch 126/250  
- 26s - loss: 0.3848 - acc: 0.9570 - val\_loss: 0.5209 - val\_acc: 0.9272  
Epoch 127/250  
- 26s - loss: 0.3806 - acc: 0.9598 - val\_loss: 0.5275 - val\_acc: 0.9230  
Epoch 128/250  
- 26s - loss: 0.3768 - acc: 0.9579 - val\_loss: 0.5334 - val\_acc: 0.9206  
Epoch 129/250  
- 26s - loss: 0.3748 - acc: 0.9590 - val\_loss: 0.5072 - val\_acc: 0.9279  
Epoch 130/250  
- 26s - loss: 0.3701 - acc: 0.9596 - val\_loss: 0.5136 - val\_acc: 0.9254  
Epoch 131/250  
- 26s - loss: 0.3667 - acc: 0.9609 - val\_loss: 0.5071 - val\_acc: 0.9281  
Epoch 132/250  
- 26s - loss: 0.3690 - acc: 0.9591 - val\_loss: 0.5158 - val\_acc: 0.9250  
Epoch 133/250  
- 26s - loss: 0.3628 - acc: 0.9612 - val\_loss: 0.5124 - val\_acc: 0.9287  
Epoch 134/250  
- 26s - loss: 0.3612 - acc: 0.9607 - val\_loss: 0.5144 - val\_acc: 0.9254  
Epoch 135/250  
- 26s - loss: 0.3584 - acc: 0.9613 - val\_loss: 0.5078 - val\_acc: 0.9273  
Epoch 136/250  
- 26s - loss: 0.3586 - acc: 0.9601 - val\_loss: 0.5084 - val\_acc: 0.9264  
Epoch 137/250  
- 25s - loss: 0.3571 - acc: 0.9611 - val\_loss: 0.4916 - val\_acc: 0.9285  
Epoch 138/250  
- 25s - loss: 0.3558 - acc: 0.9611 - val\_loss: 0.5088 - val\_acc: 0.9271  
Epoch 139/250  
- 26s - loss: 0.3503 - acc: 0.9624 - val\_loss: 0.5015 - val\_acc: 0.9267  
Epoch 140/250  
- 25s - loss: 0.3511 - acc: 0.9602 - val\_loss: 0.4979 - val\_acc: 0.9292  
Epoch 141/250  
- 26s - loss: 0.3373 - acc: 0.9660 - val\_loss: 0.5021 - val\_acc: 0.9281  
Epoch 142/250  
- 26s - loss: 0.3326 - acc: 0.9668 - val\_loss: 0.4942 - val\_acc: 0.9310  
Epoch 143/250  
- 25s - loss: 0.3267 - acc: 0.9683 - val\_loss: 0.4920 - val\_acc: 0.9330  
Epoch 144/250  
- 26s - loss: 0.3257 - acc: 0.9691 - val\_loss: 0.4956 - val\_acc: 0.9304  
Epoch 145/250  
- 25s - loss: 0.3228 - acc: 0.9695 - val\_loss: 0.5083 - val\_acc: 0.9276  
Epoch 146/250  
- 26s - loss: 0.3220 - acc: 0.9699 - val\_loss: 0.4970 - val\_acc: 0.9301  
Epoch 147/250  
- 25s - loss: 0.3195 - acc: 0.9699 - val\_loss: 0.5012 - val\_acc: 0.9296  
Epoch 148/250

# Machine Learning Project Report

- 25s - loss: 0.3184 - acc: 0.9698 - val\_loss: 0.4989 - val\_acc: 0.9283  
Epoch 149/250  
- 26s - loss: 0.3179 - acc: 0.9691 - val\_loss: 0.4889 - val\_acc: 0.9291  
Epoch 150/250  
- 25s - loss: 0.3139 - acc: 0.9706 - val\_loss: 0.4979 - val\_acc: 0.9274  
Epoch 151/250  
- 26s - loss: 0.3108 - acc: 0.9708 - val\_loss: 0.5048 - val\_acc: 0.9269  
Epoch 152/250  
- 26s - loss: 0.3095 - acc: 0.9715 - val\_loss: 0.4918 - val\_acc: 0.9310  
Epoch 153/250  
- 26s - loss: 0.3094 - acc: 0.9709 - val\_loss: 0.5001 - val\_acc: 0.9284  
Epoch 154/250  
- 25s - loss: 0.3042 - acc: 0.9724 - val\_loss: 0.4878 - val\_acc: 0.9293  
Epoch 155/250  
- 26s - loss: 0.3078 - acc: 0.9706 - val\_loss: 0.4882 - val\_acc: 0.9317  
Epoch 156/250  
- 26s - loss: 0.3078 - acc: 0.9712 - val\_loss: 0.4879 - val\_acc: 0.9307  
Epoch 157/250  
- 26s - loss: 0.3078 - acc: 0.9710 - val\_loss: 0.4811 - val\_acc: 0.9322  
Epoch 158/250  
- 26s - loss: 0.3005 - acc: 0.9720 - val\_loss: 0.4889 - val\_acc: 0.9299  
Epoch 159/250  
- 26s - loss: 0.3000 - acc: 0.9719 - val\_loss: 0.4740 - val\_acc: 0.9344  
Epoch 160/250  
- 26s - loss: 0.2971 - acc: 0.9725 - val\_loss: 0.4790 - val\_acc: 0.9317  
Epoch 161/250  
- 26s - loss: 0.2930 - acc: 0.9739 - val\_loss: 0.4819 - val\_acc: 0.9330  
Epoch 162/250  
- 26s - loss: 0.2915 - acc: 0.9739 - val\_loss: 0.4751 - val\_acc: 0.9341  
Epoch 163/250  
- 26s - loss: 0.2895 - acc: 0.9752 - val\_loss: 0.4809 - val\_acc: 0.9325  
Epoch 164/250  
- 26s - loss: 0.2848 - acc: 0.9768 - val\_loss: 0.4778 - val\_acc: 0.9335  
Epoch 165/250  
- 26s - loss: 0.2857 - acc: 0.9758 - val\_loss: 0.4805 - val\_acc: 0.9331  
Epoch 166/250  
- 26s - loss: 0.2845 - acc: 0.9766 - val\_loss: 0.4818 - val\_acc: 0.9335  
Epoch 167/250  
- 26s - loss: 0.2812 - acc: 0.9779 - val\_loss: 0.4763 - val\_acc: 0.9350  
Epoch 168/250  
- 25s - loss: 0.2799 - acc: 0.9777 - val\_loss: 0.4867 - val\_acc: 0.9312  
Epoch 169/250  
- 26s - loss: 0.2845 - acc: 0.9754 - val\_loss: 0.4861 - val\_acc: 0.9318  
Epoch 170/250  
- 26s - loss: 0.2824 - acc: 0.9757 - val\_loss: 0.4772 - val\_acc: 0.9332  
Epoch 171/250  
- 25s - loss: 0.2783 - acc: 0.9775 - val\_loss: 0.4743 - val\_acc: 0.9324  
Epoch 172/250  
- 26s - loss: 0.2775 - acc: 0.9767 - val\_loss: 0.4778 - val\_acc: 0.9323  
Epoch 173/250  
- 25s - loss: 0.2787 - acc: 0.9767 - val\_loss: 0.4816 - val\_acc: 0.9309  
Epoch 174/250

# Machine Learning Project Report

```
- 26s - loss: 0.2760 - acc: 0.9779 - val_loss: 0.4824 - val_acc: 0.9312
Epoch 175/250
- 25s - loss: 0.2757 - acc: 0.9773 - val_loss: 0.4823 - val_acc: 0.9319
Epoch 176/250
- 25s - loss: 0.2757 - acc: 0.9773 - val_loss: 0.4814 - val_acc: 0.9318
Epoch 177/250
- 26s - loss: 0.2773 - acc: 0.9765 - val_loss: 0.4815 - val_acc: 0.9330
Epoch 178/250
- 26s - loss: 0.2731 - acc: 0.9779 - val_loss: 0.4794 - val_acc: 0.9321
Epoch 179/250
- 26s - loss: 0.2706 - acc: 0.9786 - val_loss: 0.4787 - val_acc: 0.9322
Epoch 180/250
- 26s - loss: 0.2712 - acc: 0.9782 - val_loss: 0.4835 - val_acc: 0.9309
Epoch 181/250
- 26s - loss: 0.2692 - acc: 0.9795 - val_loss: 0.4797 - val_acc: 0.9322
Epoch 182/250
- 26s - loss: 0.2715 - acc: 0.9786 - val_loss: 0.4800 - val_acc: 0.9318
Epoch 183/250
- 26s - loss: 0.2666 - acc: 0.9795 - val_loss: 0.4783 - val_acc: 0.9333
Epoch 184/250
- 26s - loss: 0.2673 - acc: 0.9794 - val_loss: 0.4770 - val_acc: 0.9325
Epoch 185/250
- 26s - loss: 0.2629 - acc: 0.9810 - val_loss: 0.4771 - val_acc: 0.9334
Epoch 186/250
- 26s - loss: 0.2693 - acc: 0.9784 - val_loss: 0.4766 - val_acc: 0.9335
Epoch 187/250
- 26s - loss: 0.2629 - acc: 0.9807 - val_loss: 0.4762 - val_acc: 0.9334
Epoch 188/250
- 26s - loss: 0.2667 - acc: 0.9799 - val_loss: 0.4785 - val_acc: 0.9329
Epoch 189/250
- 26s - loss: 0.2617 - acc: 0.9808 - val_loss: 0.4751 - val_acc: 0.9321
Epoch 190/250
- 26s - loss: 0.2655 - acc: 0.9789 - val_loss: 0.4757 - val_acc: 0.9320
Epoch 191/250
- 26s - loss: 0.2616 - acc: 0.9803 - val_loss: 0.4793 - val_acc: 0.9324
Epoch 192/250
- 25s - loss: 0.2632 - acc: 0.9802 - val_loss: 0.4782 - val_acc: 0.9333
Epoch 193/250
- 26s - loss: 0.2624 - acc: 0.9805 - val_loss: 0.4781 - val_acc: 0.9320
Epoch 194/250
- 26s - loss: 0.2588 - acc: 0.9811 - val_loss: 0.4754 - val_acc: 0.9337
Epoch 195/250
- 26s - loss: 0.2623 - acc: 0.9799 - val_loss: 0.4716 - val_acc: 0.9336
Epoch 196/250
- 26s - loss: 0.2616 - acc: 0.9798 - val_loss: 0.4738 - val_acc: 0.9334
Epoch 197/250
- 25s - loss: 0.2618 - acc: 0.9802 - val_loss: 0.4746 - val_acc: 0.9334
Epoch 198/250
- 26s - loss: 0.2592 - acc: 0.9810 - val_loss: 0.4781 - val_acc: 0.9319
Epoch 199/250
- 26s - loss: 0.2568 - acc: 0.9817 - val_loss: 0.4780 - val_acc: 0.9318
Epoch 200/250
```



# Machine Learning Project Report

```
- 26s - loss: 0.2614 - acc: 0.9800 - val_loss: 0.4802 - val_acc: 0.9320
Epoch 201/250
- 25s - loss: 0.2575 - acc: 0.9814 - val_loss: 0.4789 - val_acc: 0.9325
Epoch 202/250
- 26s - loss: 0.2560 - acc: 0.9818 - val_loss: 0.4760 - val_acc: 0.9318
Epoch 203/250
- 25s - loss: 0.2551 - acc: 0.9816 - val_loss: 0.4750 - val_acc: 0.9333
Epoch 204/250
- 26s - loss: 0.2553 - acc: 0.9822 - val_loss: 0.4739 - val_acc: 0.9333
Epoch 205/250
- 26s - loss: 0.2577 - acc: 0.9809 - val_loss: 0.4752 - val_acc: 0.9326
Epoch 206/250
- 25s - loss: 0.2551 - acc: 0.9812 - val_loss: 0.4726 - val_acc: 0.9340
Epoch 207/250
- 26s - loss: 0.2559 - acc: 0.9818 - val_loss: 0.4736 - val_acc: 0.9327
Epoch 208/250
- 26s - loss: 0.2550 - acc: 0.9809 - val_loss: 0.4754 - val_acc: 0.9318
Epoch 209/250
- 26s - loss: 0.2528 - acc: 0.9825 - val_loss: 0.4747 - val_acc: 0.9330
Epoch 210/250
- 26s - loss: 0.2555 - acc: 0.9814 - val_loss: 0.4726 - val_acc: 0.9330
Epoch 211/250
- 26s - loss: 0.2537 - acc: 0.9813 - val_loss: 0.4722 - val_acc: 0.9333
Epoch 212/250
- 26s - loss: 0.2539 - acc: 0.9819 - val_loss: 0.4745 - val_acc: 0.9330
Epoch 213/250
- 25s - loss: 0.2525 - acc: 0.9822 - val_loss: 0.4732 - val_acc: 0.9329
Epoch 214/250
- 26s - loss: 0.2531 - acc: 0.9815 - val_loss: 0.4739 - val_acc: 0.9329
Epoch 215/250
- 26s - loss: 0.2503 - acc: 0.9826 - val_loss: 0.4736 - val_acc: 0.9333
Epoch 216/250
- 26s - loss: 0.2517 - acc: 0.9820 - val_loss: 0.4738 - val_acc: 0.9331
Epoch 217/250
- 26s - loss: 0.2521 - acc: 0.9821 - val_loss: 0.4726 - val_acc: 0.9329
Epoch 218/250
- 26s - loss: 0.2535 - acc: 0.9815 - val_loss: 0.4741 - val_acc: 0.9333
Epoch 219/250
- 26s - loss: 0.2519 - acc: 0.9820 - val_loss: 0.4740 - val_acc: 0.9330
Epoch 220/250
- 26s - loss: 0.2521 - acc: 0.9824 - val_loss: 0.4780 - val_acc: 0.9318
Epoch 221/250
- 26s - loss: 0.2512 - acc: 0.9826 - val_loss: 0.4780 - val_acc: 0.9317
Epoch 222/250
- 26s - loss: 0.2541 - acc: 0.9809 - val_loss: 0.4776 - val_acc: 0.9315
Epoch 223/250
- 26s - loss: 0.2491 - acc: 0.9834 - val_loss: 0.4756 - val_acc: 0.9323
Epoch 224/250
- 26s - loss: 0.2488 - acc: 0.9835 - val_loss: 0.4765 - val_acc: 0.9326
Epoch 225/250
- 26s - loss: 0.2538 - acc: 0.9810 - val_loss: 0.4763 - val_acc: 0.9332
Epoch 226/250
```

# Machine Learning Project Report

```
- 26s - loss: 0.2484 - acc: 0.9832 - val_loss: 0.4757 - val_acc: 0.9328
Epoch 227/250
- 26s - loss: 0.2482 - acc: 0.9832 - val_loss: 0.4736 - val_acc: 0.9333
Epoch 228/250
- 26s - loss: 0.2509 - acc: 0.9826 - val_loss: 0.4726 - val_acc: 0.9329
Epoch 229/250
- 26s - loss: 0.2515 - acc: 0.9818 - val_loss: 0.4748 - val_acc: 0.9326
Epoch 230/250
- 26s - loss: 0.2500 - acc: 0.9826 - val_loss: 0.4751 - val_acc: 0.9327
Epoch 231/250
- 25s - loss: 0.2505 - acc: 0.9821 - val_loss: 0.4728 - val_acc: 0.9334
Epoch 232/250
- 26s - loss: 0.2505 - acc: 0.9828 - val_loss: 0.4746 - val_acc: 0.9339
Epoch 233/250
- 26s - loss: 0.2492 - acc: 0.9827 - val_loss: 0.4743 - val_acc: 0.9338
Epoch 234/250
- 26s - loss: 0.2497 - acc: 0.9832 - val_loss: 0.4725 - val_acc: 0.9342
Epoch 235/250
- 25s - loss: 0.2479 - acc: 0.9830 - val_loss: 0.4741 - val_acc: 0.9332
Epoch 236/250
- 26s - loss: 0.2498 - acc: 0.9827 - val_loss: 0.4740 - val_acc: 0.9327
Epoch 237/250
- 26s - loss: 0.2484 - acc: 0.9830 - val_loss: 0.4755 - val_acc: 0.9333
Epoch 238/250
- 26s - loss: 0.2492 - acc: 0.9825 - val_loss: 0.4746 - val_acc: 0.9331
Epoch 239/250
- 26s - loss: 0.2478 - acc: 0.9837 - val_loss: 0.4751 - val_acc: 0.9323
Epoch 240/250
- 26s - loss: 0.2483 - acc: 0.9822 - val_loss: 0.4748 - val_acc: 0.9338
Epoch 241/250
- 25s - loss: 0.2462 - acc: 0.9833 - val_loss: 0.4755 - val_acc: 0.9337
Epoch 242/250
- 26s - loss: 0.2503 - acc: 0.9821 - val_loss: 0.4744 - val_acc: 0.9331
Epoch 243/250
- 26s - loss: 0.2448 - acc: 0.9839 - val_loss: 0.4749 - val_acc: 0.9330
Epoch 244/250
- 25s - loss: 0.2497 - acc: 0.9825 - val_loss: 0.4748 - val_acc: 0.9334
Epoch 245/250
- 26s - loss: 0.2490 - acc: 0.9832 - val_loss: 0.4741 - val_acc: 0.9339
Epoch 246/250
- 26s - loss: 0.2478 - acc: 0.9830 - val_loss: 0.4749 - val_acc: 0.9336
Epoch 247/250
- 26s - loss: 0.2507 - acc: 0.9821 - val_loss: 0.4748 - val_acc: 0.9341
Epoch 248/250
- 26s - loss: 0.2467 - acc: 0.9831 - val_loss: 0.4733 - val_acc: 0.9335
Epoch 249/250
- 26s - loss: 0.2440 - acc: 0.9844 - val_loss: 0.4730 - val_acc: 0.9341
Epoch 250/250
- 26s - loss: 0.2469 - acc: 0.9841 - val_loss: 0.4756 - val_acc: 0.9339
Training Time: 6461.9510769844055
Accuracy: 93.39 %
```

# Machine Learning Project Report

## Third Run

```
Epoch 1/250
- 31s - loss: 21.9502 - acc: 0.2262 - val_loss: 24.6493 - val_acc: 0.0951
Epoch 2/250
- 27s - loss: 11.4498 - acc: 0.3245 - val_loss: 10.7797 - val_acc: 0.1102
Epoch 3/250
- 27s - loss: 6.5736 - acc: 0.3558 - val_loss: 18.1988 - val_acc: 0.1023
Epoch 4/250
- 27s - loss: 5.3329 - acc: 0.2797 - val_loss: 4.5784 - val_acc: 0.2320
Epoch 5/250
- 26s - loss: 4.2790 - acc: 0.2733 - val_loss: 4.6100 - val_acc: 0.1474
Epoch 6/250
- 26s - loss: 3.2948 - acc: 0.3369 - val_loss: 3.1202 - val_acc: 0.2552
Epoch 7/250
- 27s - loss: 2.3483 - acc: 0.4322 - val_loss: 2.1629 - val_acc: 0.4201
Epoch 8/250
- 26s - loss: 1.9024 - acc: 0.5010 - val_loss: 1.8787 - val_acc: 0.4677
Epoch 9/250
- 27s - loss: 1.7048 - acc: 0.5511 - val_loss: 1.7039 - val_acc: 0.5599
Epoch 10/250
- 26s - loss: 1.6062 - acc: 0.5883 - val_loss: 1.4815 - val_acc: 0.6374
Epoch 11/250
- 26s - loss: 1.5467 - acc: 0.6235 - val_loss: 1.4711 - val_acc: 0.6394
Epoch 12/250
- 26s - loss: 1.4935 - acc: 0.6551 - val_loss: 1.6533 - val_acc: 0.6205
Epoch 13/250
- 26s - loss: 1.4751 - acc: 0.6752 - val_loss: 1.4192 - val_acc: 0.6989
Epoch 14/250
- 26s - loss: 1.4693 - acc: 0.6867 - val_loss: 1.6493 - val_acc: 0.6345
Epoch 15/250
- 26s - loss: 1.4591 - acc: 0.6997 - val_loss: 1.4682 - val_acc: 0.6948
Epoch 16/250
- 27s - loss: 1.4573 - acc: 0.7063 - val_loss: 1.4564 - val_acc: 0.7139
Epoch 17/250
- 26s - loss: 1.4622 - acc: 0.7141 - val_loss: 1.4293 - val_acc: 0.7227
Epoch 18/250
- 26s - loss: 1.4641 - acc: 0.7179 - val_loss: 1.4165 - val_acc: 0.7321
Epoch 19/250
- 26s - loss: 1.4757 - acc: 0.7219 - val_loss: 1.6612 - val_acc: 0.6853
Epoch 20/250
- 26s - loss: 1.4755 - acc: 0.7274 - val_loss: 1.4602 - val_acc: 0.7316
Epoch 21/250
- 26s - loss: 1.3157 - acc: 0.7658 - val_loss: 1.1901 - val_acc: 0.7829
Epoch 22/250
- 26s - loss: 1.2159 - acc: 0.7765 - val_loss: 1.1474 - val_acc: 0.7971
Epoch 23/250
- 26s - loss: 1.1991 - acc: 0.7766 - val_loss: 1.1631 - val_acc: 0.7841
Epoch 24/250
- 26s - loss: 1.1872 - acc: 0.7771 - val_loss: 1.0820 - val_acc: 0.8136
Epoch 25/250
```

# Machine Learning Project Report

```
- 26s - loss: 1.1886 - acc: 0.7782 - val_loss: 1.1479 - val_acc: 0.7958
Epoch 26/250
- 26s - loss: 1.1877 - acc: 0.7813 - val_loss: 1.2066 - val_acc: 0.7766
Epoch 27/250
- 26s - loss: 1.1925 - acc: 0.7854 - val_loss: 1.1859 - val_acc: 0.7888
Epoch 28/250
- 26s - loss: 1.1938 - acc: 0.7870 - val_loss: 1.1842 - val_acc: 0.7909
Epoch 29/250
- 26s - loss: 1.1972 - acc: 0.7896 - val_loss: 1.0930 - val_acc: 0.8227
Epoch 30/250
- 26s - loss: 1.2036 - acc: 0.7893 - val_loss: 1.1792 - val_acc: 0.8000
Epoch 31/250
- 26s - loss: 1.2029 - acc: 0.7925 - val_loss: 1.1930 - val_acc: 0.7982
Epoch 32/250
- 26s - loss: 1.2071 - acc: 0.7922 - val_loss: 1.1346 - val_acc: 0.8185
Epoch 33/250
- 26s - loss: 1.2075 - acc: 0.7950 - val_loss: 1.2183 - val_acc: 0.7907
Epoch 34/250
- 26s - loss: 1.2136 - acc: 0.7939 - val_loss: 1.1572 - val_acc: 0.8140
Epoch 35/250
- 26s - loss: 1.2058 - acc: 0.7994 - val_loss: 1.1263 - val_acc: 0.8224
Epoch 36/250
- 26s - loss: 1.2163 - acc: 0.7957 - val_loss: 1.1311 - val_acc: 0.8267
Epoch 37/250
- 26s - loss: 1.2169 - acc: 0.7999 - val_loss: 1.1166 - val_acc: 0.8332
Epoch 38/250
- 26s - loss: 1.2242 - acc: 0.7986 - val_loss: 1.1439 - val_acc: 0.8246
Epoch 39/250
- 26s - loss: 1.2172 - acc: 0.8029 - val_loss: 1.1134 - val_acc: 0.8355
Epoch 40/250
- 26s - loss: 1.2181 - acc: 0.8017 - val_loss: 1.1796 - val_acc: 0.8109
Epoch 41/250
- 26s - loss: 1.0991 - acc: 0.8307 - val_loss: 0.9634 - val_acc: 0.8690
Epoch 42/250
- 26s - loss: 1.0205 - acc: 0.8417 - val_loss: 0.9860 - val_acc: 0.8466
Epoch 43/250
- 26s - loss: 0.9932 - acc: 0.8432 - val_loss: 0.9213 - val_acc: 0.8605
Epoch 44/250
- 26s - loss: 0.9761 - acc: 0.8407 - val_loss: 0.9348 - val_acc: 0.8510
Epoch 45/250
- 26s - loss: 0.9658 - acc: 0.8406 - val_loss: 0.9890 - val_acc: 0.8326
Epoch 46/250
- 26s - loss: 0.9614 - acc: 0.8424 - val_loss: 0.8941 - val_acc: 0.8578
Epoch 47/250
- 26s - loss: 0.9559 - acc: 0.8443 - val_loss: 0.9385 - val_acc: 0.8482
Epoch 48/250
- 26s - loss: 0.9570 - acc: 0.8441 - val_loss: 0.9702 - val_acc: 0.8452
Epoch 49/250
- 26s - loss: 0.9452 - acc: 0.8446 - val_loss: 0.9311 - val_acc: 0.8490
Epoch 50/250
- 26s - loss: 0.9513 - acc: 0.8436 - val_loss: 0.9148 - val_acc: 0.8550
Epoch 51/250
```

# Machine Learning Project Report

```
- 26s - loss: 0.9561 - acc: 0.8442 - val_loss: 0.9155 - val_acc: 0.8574
Epoch 52/250
- 26s - loss: 0.9556 - acc: 0.8446 - val_loss: 0.9564 - val_acc: 0.8438
Epoch 53/250
- 26s - loss: 0.9588 - acc: 0.8438 - val_loss: 0.8678 - val_acc: 0.8723
Epoch 54/250
- 26s - loss: 0.9589 - acc: 0.8454 - val_loss: 0.9817 - val_acc: 0.8374
Epoch 55/250
- 26s - loss: 0.9590 - acc: 0.8466 - val_loss: 0.9094 - val_acc: 0.8619
Epoch 56/250
- 26s - loss: 0.9592 - acc: 0.8472 - val_loss: 0.9251 - val_acc: 0.8592
Epoch 57/250
- 26s - loss: 0.9626 - acc: 0.8481 - val_loss: 0.9012 - val_acc: 0.8669
Epoch 58/250
- 26s - loss: 0.9597 - acc: 0.8478 - val_loss: 0.8887 - val_acc: 0.8699
Epoch 59/250
- 26s - loss: 0.9560 - acc: 0.8523 - val_loss: 0.9056 - val_acc: 0.8677
Epoch 60/250
- 26s - loss: 0.9598 - acc: 0.8492 - val_loss: 0.9048 - val_acc: 0.8687
Epoch 61/250
- 26s - loss: 0.8739 - acc: 0.8751 - val_loss: 0.8284 - val_acc: 0.8850
Epoch 62/250
- 26s - loss: 0.8310 - acc: 0.8822 - val_loss: 0.8066 - val_acc: 0.8829
Epoch 63/250
- 26s - loss: 0.8025 - acc: 0.8849 - val_loss: 0.7702 - val_acc: 0.8929
Epoch 64/250
- 26s - loss: 0.7830 - acc: 0.8845 - val_loss: 0.7729 - val_acc: 0.8890
Epoch 65/250
- 26s - loss: 0.7756 - acc: 0.8849 - val_loss: 0.7411 - val_acc: 0.8935
Epoch 66/250
- 26s - loss: 0.7595 - acc: 0.8860 - val_loss: 0.7638 - val_acc: 0.8868
Epoch 67/250
- 26s - loss: 0.7591 - acc: 0.8839 - val_loss: 0.7345 - val_acc: 0.8933
Epoch 68/250
- 26s - loss: 0.7485 - acc: 0.8879 - val_loss: 0.7308 - val_acc: 0.8892
Epoch 69/250
- 26s - loss: 0.7517 - acc: 0.8845 - val_loss: 0.7407 - val_acc: 0.8885
Epoch 70/250
- 26s - loss: 0.7443 - acc: 0.8846 - val_loss: 0.7480 - val_acc: 0.8860
Epoch 71/250
- 26s - loss: 0.7345 - acc: 0.8873 - val_loss: 0.7673 - val_acc: 0.8779
Epoch 72/250
- 26s - loss: 0.7430 - acc: 0.8860 - val_loss: 0.7799 - val_acc: 0.8739
Epoch 73/250
- 26s - loss: 0.7358 - acc: 0.8873 - val_loss: 0.7470 - val_acc: 0.8825
Epoch 74/250
- 26s - loss: 0.7410 - acc: 0.8857 - val_loss: 0.7092 - val_acc: 0.8940
Epoch 75/250
- 26s - loss: 0.7389 - acc: 0.8866 - val_loss: 0.7369 - val_acc: 0.8905
Epoch 76/250
- 26s - loss: 0.7382 - acc: 0.8867 - val_loss: 0.7498 - val_acc: 0.8844
Epoch 77/250
```

# Machine Learning Project Report

```
- 26s - loss: 0.7419 - acc: 0.8864 - val_loss: 0.7584 - val_acc: 0.8821
Epoch 78/250
- 26s - loss: 0.7332 - acc: 0.8890 - val_loss: 0.7181 - val_acc: 0.8953
Epoch 79/250
- 26s - loss: 0.7356 - acc: 0.8889 - val_loss: 0.7487 - val_acc: 0.8846
Epoch 80/250
- 26s - loss: 0.7384 - acc: 0.8881 - val_loss: 0.7537 - val_acc: 0.8851
Epoch 81/250
- 26s - loss: 0.6792 - acc: 0.9047 - val_loss: 0.6754 - val_acc: 0.9056
Epoch 82/250
- 26s - loss: 0.6455 - acc: 0.9132 - val_loss: 0.6638 - val_acc: 0.9086
Epoch 83/250
- 26s - loss: 0.6316 - acc: 0.9150 - val_loss: 0.6499 - val_acc: 0.9096
Epoch 84/250
- 26s - loss: 0.6243 - acc: 0.9150 - val_loss: 0.6323 - val_acc: 0.9134
Epoch 85/250
- 26s - loss: 0.6093 - acc: 0.9169 - val_loss: 0.6313 - val_acc: 0.9133
Epoch 86/250
- 26s - loss: 0.6055 - acc: 0.9163 - val_loss: 0.6460 - val_acc: 0.9086
Epoch 87/250
- 26s - loss: 0.6030 - acc: 0.9150 - val_loss: 0.6247 - val_acc: 0.9089
Epoch 88/250
- 26s - loss: 0.5855 - acc: 0.9195 - val_loss: 0.6268 - val_acc: 0.9111
Epoch 89/250
- 26s - loss: 0.5852 - acc: 0.9186 - val_loss: 0.6223 - val_acc: 0.9076
Epoch 90/250
- 26s - loss: 0.5804 - acc: 0.9174 - val_loss: 0.6642 - val_acc: 0.8978
Epoch 91/250
- 26s - loss: 0.5832 - acc: 0.9171 - val_loss: 0.6142 - val_acc: 0.9093
Epoch 92/250
- 26s - loss: 0.5738 - acc: 0.9176 - val_loss: 0.6241 - val_acc: 0.9070
Epoch 93/250
- 26s - loss: 0.5723 - acc: 0.9179 - val_loss: 0.6050 - val_acc: 0.9091
Epoch 94/250
- 26s - loss: 0.5742 - acc: 0.9174 - val_loss: 0.6202 - val_acc: 0.9058
Epoch 95/250
- 26s - loss: 0.5702 - acc: 0.9183 - val_loss: 0.6323 - val_acc: 0.9023
Epoch 96/250
- 26s - loss: 0.5688 - acc: 0.9186 - val_loss: 0.6241 - val_acc: 0.9046
Epoch 97/250
- 26s - loss: 0.5666 - acc: 0.9174 - val_loss: 0.5909 - val_acc: 0.9154
Epoch 98/250
- 26s - loss: 0.5681 - acc: 0.9175 - val_loss: 0.6100 - val_acc: 0.9058
Epoch 99/250
- 26s - loss: 0.5606 - acc: 0.9193 - val_loss: 0.6055 - val_acc: 0.9065
Epoch 100/250
- 26s - loss: 0.5577 - acc: 0.9209 - val_loss: 0.6004 - val_acc: 0.9106
Epoch 101/250
- 26s - loss: 0.5217 - acc: 0.9297 - val_loss: 0.5757 - val_acc: 0.9169
Epoch 102/250
- 26s - loss: 0.5017 - acc: 0.9369 - val_loss: 0.5882 - val_acc: 0.9104
Epoch 103/250
```

# Machine Learning Project Report

- 26s - loss: 0.4963 - acc: 0.9373 - val\_loss: 0.5784 - val\_acc: 0.9124  
Epoch 104/250  
- 26s - loss: 0.4829 - acc: 0.9406 - val\_loss: 0.5842 - val\_acc: 0.9151  
Epoch 105/250  
- 26s - loss: 0.4776 - acc: 0.9394 - val\_loss: 0.5693 - val\_acc: 0.9154  
Epoch 106/250  
- 26s - loss: 0.4728 - acc: 0.9421 - val\_loss: 0.5490 - val\_acc: 0.9232  
Epoch 107/250  
- 26s - loss: 0.4682 - acc: 0.9411 - val\_loss: 0.5459 - val\_acc: 0.9204  
Epoch 108/250  
- 26s - loss: 0.4622 - acc: 0.9423 - val\_loss: 0.5458 - val\_acc: 0.9204  
Epoch 109/250  
- 26s - loss: 0.4637 - acc: 0.9417 - val\_loss: 0.5413 - val\_acc: 0.9196  
Epoch 110/250  
- 26s - loss: 0.4598 - acc: 0.9415 - val\_loss: 0.5578 - val\_acc: 0.9139  
Epoch 111/250  
- 27s - loss: 0.4526 - acc: 0.9438 - val\_loss: 0.5497 - val\_acc: 0.9196  
Epoch 112/250  
- 26s - loss: 0.4453 - acc: 0.9452 - val\_loss: 0.5519 - val\_acc: 0.9150  
Epoch 113/250  
- 26s - loss: 0.4492 - acc: 0.9426 - val\_loss: 0.5556 - val\_acc: 0.9141  
Epoch 114/250  
- 26s - loss: 0.4477 - acc: 0.9405 - val\_loss: 0.5426 - val\_acc: 0.9184  
Epoch 115/250  
- 26s - loss: 0.4418 - acc: 0.9424 - val\_loss: 0.5350 - val\_acc: 0.9190  
Epoch 116/250  
- 26s - loss: 0.4414 - acc: 0.9432 - val\_loss: 0.5468 - val\_acc: 0.9162  
Epoch 117/250  
- 26s - loss: 0.4427 - acc: 0.9414 - val\_loss: 0.5426 - val\_acc: 0.9155  
Epoch 118/250  
- 26s - loss: 0.4344 - acc: 0.9438 - val\_loss: 0.5377 - val\_acc: 0.9180  
Epoch 119/250  
- 26s - loss: 0.4280 - acc: 0.9441 - val\_loss: 0.5402 - val\_acc: 0.9172  
Epoch 120/250  
- 26s - loss: 0.4319 - acc: 0.9447 - val\_loss: 0.5374 - val\_acc: 0.9182  
Epoch 121/250  
- 26s - loss: 0.4080 - acc: 0.9504 - val\_loss: 0.5124 - val\_acc: 0.9256  
Epoch 122/250  
- 26s - loss: 0.3958 - acc: 0.9541 - val\_loss: 0.5183 - val\_acc: 0.9246  
Epoch 123/250  
- 26s - loss: 0.3875 - acc: 0.9558 - val\_loss: 0.5254 - val\_acc: 0.9235  
Epoch 124/250  
- 26s - loss: 0.3874 - acc: 0.9562 - val\_loss: 0.5024 - val\_acc: 0.9261  
Epoch 125/250  
- 26s - loss: 0.3815 - acc: 0.9574 - val\_loss: 0.5141 - val\_acc: 0.9252  
Epoch 126/250  
- 26s - loss: 0.3778 - acc: 0.9581 - val\_loss: 0.5122 - val\_acc: 0.9245  
Epoch 127/250  
- 26s - loss: 0.3724 - acc: 0.9603 - val\_loss: 0.5085 - val\_acc: 0.9256  
Epoch 128/250  
- 26s - loss: 0.3711 - acc: 0.9582 - val\_loss: 0.5044 - val\_acc: 0.9245  
Epoch 129/250

# Machine Learning Project Report

```
- 26s - loss: 0.3666 - acc: 0.9607 - val_loss: 0.4982 - val_acc: 0.9260
Epoch 130/250
- 26s - loss: 0.3669 - acc: 0.9590 - val_loss: 0.5059 - val_acc: 0.9239
Epoch 131/250
- 26s - loss: 0.3661 - acc: 0.9588 - val_loss: 0.5074 - val_acc: 0.9240
Epoch 132/250
- 26s - loss: 0.3631 - acc: 0.9595 - val_loss: 0.5060 - val_acc: 0.9249
Epoch 133/250
- 26s - loss: 0.3630 - acc: 0.9588 - val_loss: 0.5017 - val_acc: 0.9250
Epoch 134/250
- 26s - loss: 0.3572 - acc: 0.9604 - val_loss: 0.5019 - val_acc: 0.9237
Epoch 135/250
- 26s - loss: 0.3532 - acc: 0.9616 - val_loss: 0.5029 - val_acc: 0.9251
Epoch 136/250
- 26s - loss: 0.3536 - acc: 0.9612 - val_loss: 0.4877 - val_acc: 0.9264
Epoch 137/250
- 26s - loss: 0.3538 - acc: 0.9597 - val_loss: 0.4914 - val_acc: 0.9260
Epoch 138/250
- 26s - loss: 0.3498 - acc: 0.9608 - val_loss: 0.4873 - val_acc: 0.9269
Epoch 139/250
- 26s - loss: 0.3477 - acc: 0.9618 - val_loss: 0.4887 - val_acc: 0.9253
Epoch 140/250
- 26s - loss: 0.3470 - acc: 0.9609 - val_loss: 0.5015 - val_acc: 0.9217
Epoch 141/250
- 26s - loss: 0.3347 - acc: 0.9653 - val_loss: 0.4817 - val_acc: 0.9299
Epoch 142/250
- 26s - loss: 0.3253 - acc: 0.9684 - val_loss: 0.4792 - val_acc: 0.9301
Epoch 143/250
- 26s - loss: 0.3270 - acc: 0.9668 - val_loss: 0.4919 - val_acc: 0.9240
Epoch 144/250
- 26s - loss: 0.3213 - acc: 0.9686 - val_loss: 0.4777 - val_acc: 0.9294
Epoch 145/250
- 26s - loss: 0.3184 - acc: 0.9694 - val_loss: 0.4803 - val_acc: 0.9300
Epoch 146/250
- 26s - loss: 0.3170 - acc: 0.9694 - val_loss: 0.4840 - val_acc: 0.9291
Epoch 147/250
- 26s - loss: 0.3135 - acc: 0.9707 - val_loss: 0.4815 - val_acc: 0.9293
Epoch 148/250
- 26s - loss: 0.3116 - acc: 0.9702 - val_loss: 0.4794 - val_acc: 0.9302
Epoch 149/250
- 26s - loss: 0.3105 - acc: 0.9706 - val_loss: 0.4845 - val_acc: 0.9293
Epoch 150/250
- 26s - loss: 0.3089 - acc: 0.9702 - val_loss: 0.4786 - val_acc: 0.9308
Epoch 151/250
- 26s - loss: 0.3080 - acc: 0.9710 - val_loss: 0.4831 - val_acc: 0.9295
Epoch 152/250
- 26s - loss: 0.3030 - acc: 0.9719 - val_loss: 0.4794 - val_acc: 0.9286
Epoch 153/250
- 26s - loss: 0.3060 - acc: 0.9716 - val_loss: 0.4799 - val_acc: 0.9315
Epoch 154/250
- 26s - loss: 0.3003 - acc: 0.9725 - val_loss: 0.4857 - val_acc: 0.9288
Epoch 155/250
```



# Machine Learning Project Report

```
- 26s - loss: 0.3030 - acc: 0.9700 - val_loss: 0.4744 - val_acc: 0.9304
Epoch 156/250
- 26s - loss: 0.2965 - acc: 0.9724 - val_loss: 0.4893 - val_acc: 0.9262
Epoch 157/250
- 26s - loss: 0.2953 - acc: 0.9726 - val_loss: 0.4737 - val_acc: 0.9326
Epoch 158/250
- 26s - loss: 0.2992 - acc: 0.9712 - val_loss: 0.4759 - val_acc: 0.9304
Epoch 159/250
- 26s - loss: 0.2960 - acc: 0.9719 - val_loss: 0.4885 - val_acc: 0.9277
Epoch 160/250
- 26s - loss: 0.2938 - acc: 0.9720 - val_loss: 0.4736 - val_acc: 0.9331
Epoch 161/250
- 26s - loss: 0.2884 - acc: 0.9741 - val_loss: 0.4747 - val_acc: 0.9299
Epoch 162/250
- 26s - loss: 0.2848 - acc: 0.9752 - val_loss: 0.4747 - val_acc: 0.9311
Epoch 163/250
- 26s - loss: 0.2886 - acc: 0.9740 - val_loss: 0.4661 - val_acc: 0.9320
Epoch 164/250
- 26s - loss: 0.2837 - acc: 0.9751 - val_loss: 0.4655 - val_acc: 0.9317
Epoch 165/250
- 26s - loss: 0.2818 - acc: 0.9759 - val_loss: 0.4660 - val_acc: 0.9333
Epoch 166/250
- 26s - loss: 0.2837 - acc: 0.9750 - val_loss: 0.4678 - val_acc: 0.9321
Epoch 167/250
- 26s - loss: 0.2765 - acc: 0.9769 - val_loss: 0.4735 - val_acc: 0.9282
Epoch 168/250
- 26s - loss: 0.2828 - acc: 0.9754 - val_loss: 0.4643 - val_acc: 0.9311
Epoch 169/250
- 26s - loss: 0.2765 - acc: 0.9770 - val_loss: 0.4705 - val_acc: 0.9306
Epoch 170/250
- 26s - loss: 0.2757 - acc: 0.9768 - val_loss: 0.4621 - val_acc: 0.9345
Epoch 171/250
- 26s - loss: 0.2757 - acc: 0.9771 - val_loss: 0.4673 - val_acc: 0.9327
Epoch 172/250
- 26s - loss: 0.2730 - acc: 0.9778 - val_loss: 0.4667 - val_acc: 0.9328
Epoch 173/250
- 26s - loss: 0.2721 - acc: 0.9782 - val_loss: 0.4651 - val_acc: 0.9333
Epoch 174/250
- 26s - loss: 0.2726 - acc: 0.9779 - val_loss: 0.4714 - val_acc: 0.9315
Epoch 175/250
- 26s - loss: 0.2730 - acc: 0.9772 - val_loss: 0.4692 - val_acc: 0.9321
Epoch 176/250
- 26s - loss: 0.2711 - acc: 0.9776 - val_loss: 0.4684 - val_acc: 0.9316
Epoch 177/250
- 26s - loss: 0.2723 - acc: 0.9777 - val_loss: 0.4680 - val_acc: 0.9326
Epoch 178/250
- 26s - loss: 0.2675 - acc: 0.9784 - val_loss: 0.4645 - val_acc: 0.9335
Epoch 179/250
- 26s - loss: 0.2676 - acc: 0.9787 - val_loss: 0.4612 - val_acc: 0.9345
Epoch 180/250
- 26s - loss: 0.2682 - acc: 0.9778 - val_loss: 0.4720 - val_acc: 0.9307
Epoch 181/250
```

# Machine Learning Project Report

```
- 26s - loss: 0.2658 - acc: 0.9790 - val_loss: 0.4620 - val_acc: 0.9325
Epoch 182/250
- 26s - loss: 0.2669 - acc: 0.9784 - val_loss: 0.4617 - val_acc: 0.9333
Epoch 183/250
- 26s - loss: 0.2640 - acc: 0.9793 - val_loss: 0.4618 - val_acc: 0.9322
Epoch 184/250
- 26s - loss: 0.2593 - acc: 0.9809 - val_loss: 0.4593 - val_acc: 0.9333
Epoch 185/250
- 26s - loss: 0.2582 - acc: 0.9810 - val_loss: 0.4601 - val_acc: 0.9338
Epoch 186/250
- 26s - loss: 0.2598 - acc: 0.9802 - val_loss: 0.4632 - val_acc: 0.9335
Epoch 187/250
- 26s - loss: 0.2582 - acc: 0.9806 - val_loss: 0.4628 - val_acc: 0.9340
Epoch 188/250
- 26s - loss: 0.2614 - acc: 0.9801 - val_loss: 0.4650 - val_acc: 0.9330
Epoch 189/250
- 26s - loss: 0.2591 - acc: 0.9805 - val_loss: 0.4688 - val_acc: 0.9327
Epoch 190/250
- 26s - loss: 0.2577 - acc: 0.9804 - val_loss: 0.4659 - val_acc: 0.9319
Epoch 191/250
- 26s - loss: 0.2585 - acc: 0.9803 - val_loss: 0.4653 - val_acc: 0.9335
Epoch 192/250
- 26s - loss: 0.2598 - acc: 0.9805 - val_loss: 0.4704 - val_acc: 0.9316
Epoch 193/250
- 26s - loss: 0.2583 - acc: 0.9803 - val_loss: 0.4676 - val_acc: 0.9335
Epoch 194/250
- 26s - loss: 0.2558 - acc: 0.9815 - val_loss: 0.4631 - val_acc: 0.9334
Epoch 195/250
- 26s - loss: 0.2577 - acc: 0.9797 - val_loss: 0.4654 - val_acc: 0.9332
Epoch 196/250
- 26s - loss: 0.2568 - acc: 0.9800 - val_loss: 0.4697 - val_acc: 0.9316
Epoch 197/250
- 26s - loss: 0.2558 - acc: 0.9805 - val_loss: 0.4634 - val_acc: 0.9335
Epoch 198/250
- 26s - loss: 0.2552 - acc: 0.9806 - val_loss: 0.4664 - val_acc: 0.9333
Epoch 199/250
- 26s - loss: 0.2539 - acc: 0.9807 - val_loss: 0.4616 - val_acc: 0.9341
Epoch 200/250
- 26s - loss: 0.2559 - acc: 0.9802 - val_loss: 0.4631 - val_acc: 0.9330
Epoch 201/250
- 26s - loss: 0.2544 - acc: 0.9810 - val_loss: 0.4630 - val_acc: 0.9328
Epoch 202/250
- 26s - loss: 0.2506 - acc: 0.9819 - val_loss: 0.4606 - val_acc: 0.9342
Epoch 203/250
- 26s - loss: 0.2531 - acc: 0.9810 - val_loss: 0.4617 - val_acc: 0.9333
Epoch 204/250
- 26s - loss: 0.2505 - acc: 0.9813 - val_loss: 0.4619 - val_acc: 0.9341
Epoch 205/250
- 26s - loss: 0.2512 - acc: 0.9824 - val_loss: 0.4631 - val_acc: 0.9342
Epoch 206/250
- 26s - loss: 0.2513 - acc: 0.9816 - val_loss: 0.4615 - val_acc: 0.9341
Epoch 207/250
```

# Machine Learning Project Report

```
- 26s - loss: 0.2509 - acc: 0.9819 - val_loss: 0.4622 - val_acc: 0.9338
Epoch 208/250
- 26s - loss: 0.2500 - acc: 0.9818 - val_loss: 0.4621 - val_acc: 0.9340
Epoch 209/250
- 26s - loss: 0.2519 - acc: 0.9809 - val_loss: 0.4637 - val_acc: 0.9336
Epoch 210/250
- 26s - loss: 0.2500 - acc: 0.9819 - val_loss: 0.4621 - val_acc: 0.9337
Epoch 211/250
- 26s - loss: 0.2504 - acc: 0.9821 - val_loss: 0.4597 - val_acc: 0.9338
Epoch 212/250
- 26s - loss: 0.2506 - acc: 0.9819 - val_loss: 0.4622 - val_acc: 0.9339
Epoch 213/250
- 26s - loss: 0.2491 - acc: 0.9820 - val_loss: 0.4631 - val_acc: 0.9338
Epoch 214/250
- 26s - loss: 0.2488 - acc: 0.9824 - val_loss: 0.4632 - val_acc: 0.9330
Epoch 215/250
- 26s - loss: 0.2477 - acc: 0.9822 - val_loss: 0.4616 - val_acc: 0.9342
Epoch 216/250
- 26s - loss: 0.2451 - acc: 0.9831 - val_loss: 0.4620 - val_acc: 0.9334
Epoch 217/250
- 26s - loss: 0.2473 - acc: 0.9829 - val_loss: 0.4661 - val_acc: 0.9329
Epoch 218/250
- 26s - loss: 0.2461 - acc: 0.9828 - val_loss: 0.4623 - val_acc: 0.9338
Epoch 219/250
- 26s - loss: 0.2481 - acc: 0.9823 - val_loss: 0.4605 - val_acc: 0.9339
Epoch 220/250
- 26s - loss: 0.2480 - acc: 0.9821 - val_loss: 0.4618 - val_acc: 0.9331
Epoch 221/250
- 26s - loss: 0.2485 - acc: 0.9825 - val_loss: 0.4626 - val_acc: 0.9338
Epoch 222/250
- 26s - loss: 0.2459 - acc: 0.9830 - val_loss: 0.4603 - val_acc: 0.9348
Epoch 223/250
- 26s - loss: 0.2478 - acc: 0.9822 - val_loss: 0.4609 - val_acc: 0.9343
Epoch 224/250
- 26s - loss: 0.2465 - acc: 0.9817 - val_loss: 0.4611 - val_acc: 0.9341
Epoch 225/250
- 26s - loss: 0.2444 - acc: 0.9825 - val_loss: 0.4623 - val_acc: 0.9338
Epoch 226/250
- 26s - loss: 0.2441 - acc: 0.9837 - val_loss: 0.4605 - val_acc: 0.9345
Epoch 227/250
- 26s - loss: 0.2457 - acc: 0.9827 - val_loss: 0.4618 - val_acc: 0.9352
Epoch 228/250
- 26s - loss: 0.2457 - acc: 0.9826 - val_loss: 0.4592 - val_acc: 0.9344
Epoch 229/250
- 26s - loss: 0.2434 - acc: 0.9837 - val_loss: 0.4601 - val_acc: 0.9342
Epoch 230/250
- 26s - loss: 0.2483 - acc: 0.9817 - val_loss: 0.4604 - val_acc: 0.9341
Epoch 231/250
- 26s - loss: 0.2464 - acc: 0.9822 - val_loss: 0.4594 - val_acc: 0.9343
Epoch 232/250
- 26s - loss: 0.2453 - acc: 0.9827 - val_loss: 0.4586 - val_acc: 0.9338
Epoch 233/250
```

# Machine Learning Project Report

```
- 26s - loss: 0.2487 - acc: 0.9823 - val_loss: 0.4605 - val_acc: 0.9341
Epoch 234/250
- 26s - loss: 0.2446 - acc: 0.9832 - val_loss: 0.4601 - val_acc: 0.9342
Epoch 235/250
- 26s - loss: 0.2450 - acc: 0.9828 - val_loss: 0.4607 - val_acc: 0.9334
Epoch 236/250
- 26s - loss: 0.2441 - acc: 0.9835 - val_loss: 0.4593 - val_acc: 0.9342
Epoch 237/250
- 26s - loss: 0.2468 - acc: 0.9823 - val_loss: 0.4593 - val_acc: 0.9336
Epoch 238/250
- 26s - loss: 0.2454 - acc: 0.9829 - val_loss: 0.4611 - val_acc: 0.9338
Epoch 239/250
- 26s - loss: 0.2455 - acc: 0.9829 - val_loss: 0.4597 - val_acc: 0.9337
Epoch 240/250
- 26s - loss: 0.2452 - acc: 0.9827 - val_loss: 0.4609 - val_acc: 0.9330
Epoch 241/250
- 26s - loss: 0.2422 - acc: 0.9833 - val_loss: 0.4603 - val_acc: 0.9334
Epoch 242/250
- 26s - loss: 0.2459 - acc: 0.9830 - val_loss: 0.4613 - val_acc: 0.9334
Epoch 243/250
- 26s - loss: 0.2441 - acc: 0.9829 - val_loss: 0.4624 - val_acc: 0.9333
Epoch 244/250
- 26s - loss: 0.2446 - acc: 0.9828 - val_loss: 0.4618 - val_acc: 0.9333
Epoch 245/250
- 26s - loss: 0.2430 - acc: 0.9835 - val_loss: 0.4608 - val_acc: 0.9330
Epoch 246/250
- 26s - loss: 0.2420 - acc: 0.9840 - val_loss: 0.4614 - val_acc: 0.9332
Epoch 247/250
- 26s - loss: 0.2395 - acc: 0.9843 - val_loss: 0.4616 - val_acc: 0.9335
Epoch 248/250
- 26s - loss: 0.2445 - acc: 0.9828 - val_loss: 0.4626 - val_acc: 0.9326
Epoch 249/250
- 26s - loss: 0.2454 - acc: 0.9820 - val_loss: 0.4628 - val_acc: 0.9331
Epoch 250/250
- 26s - loss: 0.2426 - acc: 0.9836 - val_loss: 0.4616 - val_acc: 0.9331
Training Time: 6523.386714935303
Accuracy: 93.31 %
```

## Result:

	First Run	Second Run	Third Run	Average
Accuracy	93.45%	93.39%	93.31%	93.38%
Training Time	6452.848s	6461.951s	6523.386s	6479.939s

# Machine Learning Project Report

## CIFAR-100

### Introduction:

This dataset is just like the CIFAR-10, except it has 100 classes containing 600 images each. There are 500 training images and 100 testing images per class. The 100 classes in the CIFAR-100 are grouped into 20 superclasses. Each image comes with a "fine" label (the class to which it belongs) and a "coarse" label (the superclass to which it belongs).

Dataset Link: <https://www.cs.toronto.edu/~kriz/cifar-100-python.tar.gz>

### Model:

We have used Keras model based on VGG16 architecture for CIFAR-100 dataset.

### VGG16 Architecture:

This network is characterized by its simplicity, using only  $3 \times 3$  convolutional layers stacked on top of each other in increasing depth. Reducing volume size is handled by max pooling. Two fully-connected layers, each with 4,096 nodes are then followed by a softmax classifier

### Source Code:

```
from __future__ import print_function
import keras
from keras.datasets import cifar100
from keras.preprocessing.image import ImageDataGenerator
from keras.models import Sequential
from keras.layers import Dense, Dropout, Activation, Flatten
from keras.layers import Conv2D, MaxPooling2D, BatchNormalization
from keras import optimizers
import numpy as np
from keras.layers.core import Lambda
from keras import backend as K
from keras import regularizers
import time
from sklearn.metrics import accuracy_score

class cifar100vgg:
    def __init__(self, train=True):
        self.num_classes = 100
        self.weight_decay = 0.0005
        self.x_shape = [32, 32, 3]

        self.model = self.build_model()
        if train:
            self.model = self.train(self.model)
        else:
```

# Machine Learning Project Report

```
self.model.load_weights('cifar100vgg.h5')

def build_model(self):
    # Build the network of vgg for 10 classes with massive dropout and weight decay as described in the paper.

    model = Sequential()
    weight_decay = self.weight_decay

    model.add(Conv2D(64, (3, 3), padding='same',
                     input_shape=self.x_shape, kernel_regularizer=regularizers.l2(weight_decay)))
    model.add(Activation('relu'))
    model.add(BatchNormalization())
    model.add(Dropout(0.3))

    model.add(Conv2D(64, (3, 3), padding='same', kernel_regularizer=regularizers.l2(weight_decay)))
    model.add(Activation('relu'))
    model.add(BatchNormalization())

    model.add(MaxPooling2D(pool_size=(2, 2)))

    model.add(Conv2D(128, (3, 3), padding='same', kernel_regularizer=regularizers.l2(weight_decay)))
    model.add(Activation('relu'))
    model.add(BatchNormalization())
    model.add(Dropout(0.4))

    model.add(Conv2D(128, (3, 3), padding='same', kernel_regularizer=regularizers.l2(weight_decay)))
    model.add(Activation('relu'))
    model.add(BatchNormalization())

    model.add(MaxPooling2D(pool_size=(2, 2)))

    model.add(Conv2D(256, (3, 3), padding='same', kernel_regularizer=regularizers.l2(weight_decay)))
    model.add(Activation('relu'))
    model.add(BatchNormalization())
    model.add(Dropout(0.4))

    model.add(Conv2D(256, (3, 3), padding='same', kernel_regularizer=regularizers.l2(weight_decay)))
    model.add(Activation('relu'))
    model.add(BatchNormalization())
    model.add(Dropout(0.4))

    model.add(Conv2D(256, (3, 3), padding='same', kernel_regularizer=regularizers.l2(weight_decay)))
    model.add(Activation('relu'))
    model.add(BatchNormalization())

    model.add(MaxPooling2D(pool_size=(2, 2)))
```

# Machine Learning Project Report

```
model.add(Conv2D(512, (3, 3), padding='same',kernel_regularizer=regularizers.l2(weight_decay)))
model.add(Activation('relu'))
model.add(BatchNormalization())
model.add(Dropout(0.4))
```

```
model.add(Conv2D(512, (3, 3), padding='same',kernel_regularizer=regularizers.l2(weight_decay)))
model.add(Activation('relu'))
model.add(BatchNormalization())
model.add(Dropout(0.4))
```

```
model.add(Conv2D(512, (3, 3), padding='same',kernel_regularizer=regularizers.l2(weight_decay)))
model.add(Activation('relu'))
model.add(BatchNormalization())
```

```
model.add(MaxPooling2D(pool_size=(2, 2)))
```

```
model.add(Conv2D(512, (3, 3), padding='same',kernel_regularizer=regularizers.l2(weight_decay)))
model.add(Activation('relu'))
model.add(BatchNormalization())
model.add(Dropout(0.4))
```

```
model.add(Conv2D(512, (3, 3), padding='same',kernel_regularizer=regularizers.l2(weight_decay)))
model.add(Activation('relu'))
model.add(BatchNormalization())
model.add(Dropout(0.4))
```

```
model.add(Conv2D(512, (3, 3), padding='same',kernel_regularizer=regularizers.l2(weight_decay)))
model.add(Activation('relu'))
model.add(BatchNormalization())
```

```
model.add(MaxPooling2D(pool_size=(2, 2)))
model.add(Dropout(0.5))
```

```
model.add(Flatten())
model.add(Dense(512,kernel_regularizer=regularizers.l2(weight_decay)))
model.add(Activation('relu'))
model.add(BatchNormalization())
```

```
model.add(Dropout(0.5))
model.add(Dense(self.num_classes))
model.add(Activation('softmax'))
return model
```

```
def normalize(self,X_train,X_test):
    #this function normalize inputs for zero mean and unit variance
    # it is used when training a model.
```

# Machine Learning Project Report

```
# Input: training set and test set
# Output: normalized training set and test set according to the training set statistics.
mean = np.mean(X_train,axis=(0,1,2,3))
std = np.std(X_train, axis=(0, 1, 2, 3))
print(mean)
print(std)
X_train = (X_train-mean)/(std+1e-7)
X_test = (X_test-mean)/(std+1e-7)
return X_train, X_test

def normalize_production(self,x):
    #this function is used to normalize instances in production according to saved training set statistics
    # Input: X - a training set
    # Output X - a normalized training set according to normalization constants.

    #these values produced during first training and are general for the standard cifar10 training set normalization
    mean = 121.936
    std = 68.389
    return (x-mean)/(std+1e-7)

def predict(self,x,normalize=True,batch_size=50):
    if normalize:
        x = self.normalize_production(x)
    return self.model.predict(x,batch_size)

def train(self,model):

    #training parameters
    batch_size = 128
    maxepoches = 250
    learning_rate = 0.1
    lr_decay = 1e-6
    lr_drop = 20

    # The data, shuffled and split between train and test sets:
    (x_train, y_train), (x_test, y_test) = cifar100.load_data()
    x_train = x_train.astype('float32')
    x_test = x_test.astype('float32')
    x_train, x_test = self.normalize(x_train, x_test)

    y_train = keras.utils.to_categorical(y_train, self.num_classes)
    y_test = keras.utils.to_categorical(y_test, self.num_classes)

    def lr_scheduler(epoch):
        return learning_rate * (0.5 ** (epoch // lr_drop))
    reduce_lr = keras.callbacks.LearningRateScheduler(lr_scheduler)
```



# Machine Learning Project Report

```
#data augmentation
datagen = ImageDataGenerator(
    featurewise_center=False, # set input mean to 0 over the dataset
    samplewise_center=False, # set each sample mean to 0
    featurewise_std_normalization=False, # divide inputs by std of the dataset
    samplewise_std_normalization=False, # divide each input by its std
    zca_whitening=False, # apply ZCA whitening
    rotation_range=15, # randomly rotate images in the range (degrees, 0 to 180)
    width_shift_range=0.1, # randomly shift images horizontally (fraction of total width)
    height_shift_range=0.1, # randomly shift images vertically (fraction of total height)
    horizontal_flip=True, # randomly flip images
    vertical_flip=False) # randomly flip images
# (std, mean, and principal components if ZCA whitening is applied).
start = time.time()
datagen.fit(x_train)

#optimization details
sgd = optimizers.SGD(lr=learning_rate, decay=lr_decay, momentum=0.9, nesterov=True)
model.compile(loss='categorical_crossentropy', optimizer=sgd, metrics=['accuracy'])

# training process in a for loop with learning rate drop every 25 epoches.

historytemp = model.fit_generator(datagen.flow(x_train, y_train,
                                                batch_size=batch_size),
                                steps_per_epoch=x_train.shape[0] // batch_size,
                                epochs=maxepoches,
                                validation_data=(x_test, y_test), callbacks=[reduce_lr], verbose=2)
model.save_weights('cifar100vgg.h5')
end = time.time()
print("Training Time :", (end-start))
return model

if __name__ == '__main__':
    (x_train, y_train), (x_test, y_test) = cifar100.load_data()
    x_train = x_train.astype('float32')
    x_test = x_test.astype('float32')

    y_train = keras.utils.to_categorical(y_train, 100)
    y_test = keras.utils.to_categorical(y_test, 100)

    model = cifar100vgg()

    predicted_x = model.predict(x_test)
    accuracy = accuracy_score(np.argmax(predicted_x, 1), np.argmax(y_test, 1))
```

# Machine Learning Project Report

```
print("Accuracy: ",accuracy*100,"%" )
```

Output:

Run 1:

Using TensorFlow backend.

The default version of TensorFlow in Colab will soon switch to TensorFlow 2.x.

We recommend you [upgrade](#) now or ensure your notebook will continue to use TensorFlow 1.x via the `%tensorflow_version 1.x` magic: [more info](#).

Downloading data from <https://www.cs.toronto.edu/~kriz/cifar-100-python.tar.gz>

169009152/169001437 [=====] - 6s 0us/step

WARNING:tensorflow:From /usr/local/lib/python3.6/dist-packages/keras/backend/tensorflow\_backend.py:66: The name tf.get\_default\_graph is deprecated. Please use tf.compat.v1.get\_default\_graph instead.

WARNING:tensorflow:From /usr/local/lib/python3.6/dist-packages/keras/backend/tensorflow\_backend.py:541: The name tf.placeholder is deprecated. Please use tf.compat.v1.placeholder instead.

WARNING:tensorflow:From /usr/local/lib/python3.6/dist-packages/keras/backend/tensorflow\_backend.py:4432: The name tf.random\_uniform is deprecated. Please use tf.random.uniform instead.

WARNING:tensorflow:From /usr/local/lib/python3.6/dist-packages/keras/backend/tensorflow\_backend.py:190: The name tf.get\_default\_session is deprecated. Please use tf.compat.v1.get\_default\_session instead.

WARNING:tensorflow:From /usr/local/lib/python3.6/dist-packages/keras/backend/tensorflow\_backend.py:197: The name tf.ConfigProto is deprecated. Please use tf.compat.v1.ConfigProto instead.

WARNING:tensorflow:From /usr/local/lib/python3.6/dist-packages/keras/backend/tensorflow\_backend.py:203: The name tf.Session is deprecated. Please use tf.compat.v1.Session instead.

WARNING:tensorflow:From /usr/local/lib/python3.6/dist-packages/keras/backend/tensorflow\_backend.py:207: The name tf.global\_variables is deprecated. Please use tf.compat.v1.global\_variables instead.

WARNING:tensorflow:From /usr/local/lib/python3.6/dist-packages/keras/backend/tensorflow\_backend.py:216: The name tf.is\_variable\_initialized is deprecated. Please use tf.compat.v1.is\_variable\_initialized instead.

# Machine Learning Project Report

```
WARNING:tensorflow:From /usr/local/lib/python3.6/dist-packages/keras/backend/tensorflow_backend.py:223: The name tf.variables_initializer is deprecated. Please use tf.compat.v1.variables_initializer instead.
```

```
WARNING:tensorflow:From /usr/local/lib/python3.6/dist-packages/keras/backend/tensorflow_backend.py:2041: The name tf.nn.fused_batch_norm is deprecated. Please use tf.compat.v1.nn.fused_batch_norm instead.
```

```
WARNING:tensorflow:From /usr/local/lib/python3.6/dist-packages/keras/backend/tensorflow_backend.py:148: The name tf.placeholder_with_default is deprecated. Please use tf.compat.v1.placeholder_with_default instead.
```

```
WARNING:tensorflow:From /usr/local/lib/python3.6/dist-packages/keras/backend/tensorflow_backend.py:3733: calling dropout (from tensorflow.python.ops.nn_ops) with keep_prob is deprecated and will be removed in a future version.
Instructions for updating:
Please use `rate` instead of `keep_prob`. Rate should be set to `rate = 1 - keep_prob`.
WARNING:tensorflow:From /usr/local/lib/python3.6/dist-packages/keras/backend/tensorflow_backend.py:4267: The name tf.nn.max_pool is deprecated. Please use tf.nn.max_pool2d instead.
```

```
121.93584
68.38902
```

```
WARNING:tensorflow:From /usr/local/lib/python3.6/dist-packages/keras/optimizers.py:793: The name tf.train.Optimizer is deprecated. Please use tf.compat.v1.train.Optimizer instead.
```

```
WARNING:tensorflow:From /usr/local/lib/python3.6/dist-packages/keras/backend/tensorflow_backend.py:3576: The name tf.log is deprecated. Please use tf.math.log instead.
```

```
WARNING:tensorflow:From /usr/local/lib/python3.6/dist-packages/tensorflow_core/python/ops/math_grad.py:1424: where (from tensorflow.python.ops.array_ops) is deprecated and will be removed in a future version.
Instructions for updating:
Use tf.where in 2.0, which has the same broadcast rule as np.where
WARNING:tensorflow:From /usr/local/lib/python3.6/dist-packages/keras/backend/tensorflow_backend.py:1033: The name tf.assign_add is deprecated. Please use tf.compat.v1.assign_add instead.
```

```
WARNING:tensorflow:From /usr/local/lib/python3.6/dist-packages/keras/backend/tensorflow_backend.py:1020: The name tf.assign is deprecated. Please use tf.compat.v1.assign instead.
```

```
Epoch 1/250
- 84s - loss: 20.6756 - acc: 0.0282 - val_loss: 16.4382 - val_acc: 0.0145
Epoch 2/250
```

# Machine Learning Project Report

```
- 73s - loss: 12.2054 - acc: 0.0417 - val_loss: 14.0973 - val_acc: 0.0138
Epoch 3/250
- 73s - loss: 8.2737 - acc: 0.0460 - val_loss: 11.3437 - val_acc: 0.0075
Epoch 4/250
- 73s - loss: 6.9882 - acc: 0.0422 - val_loss: 6.3838 - val_acc: 0.0446
Epoch 5/250
- 73s - loss: 6.1584 - acc: 0.0506 - val_loss: 5.8782 - val_acc: 0.0310
Epoch 6/250
- 73s - loss: 5.3622 - acc: 0.0641 - val_loss: 5.7121 - val_acc: 0.0223
Epoch 7/250
- 73s - loss: 4.6940 - acc: 0.0828 - val_loss: 4.7939 - val_acc: 0.0576
Epoch 8/250
- 73s - loss: 4.3459 - acc: 0.1007 - val_loss: 4.4188 - val_acc: 0.0803
Epoch 9/250
- 73s - loss: 4.1645 - acc: 0.1140 - val_loss: 4.0707 - val_acc: 0.1247
Epoch 10/250
- 73s - loss: 4.0467 - acc: 0.1274 - val_loss: 4.0644 - val_acc: 0.1303
Epoch 11/250
- 73s - loss: 3.9725 - acc: 0.1430 - val_loss: 4.1975 - val_acc: 0.1344
Epoch 12/250
- 73s - loss: 3.9167 - acc: 0.1597 - val_loss: 3.8333 - val_acc: 0.1840
Epoch 13/250
- 73s - loss: 3.8582 - acc: 0.1800 - val_loss: 3.7158 - val_acc: 0.2114
Epoch 14/250
- 73s - loss: 3.8033 - acc: 0.1985 - val_loss: 3.6846 - val_acc: 0.2296
Epoch 15/250
- 73s - loss: 3.7770 - acc: 0.2207 - val_loss: 3.7565 - val_acc: 0.2406
Epoch 16/250
- 73s - loss: 3.7441 - acc: 0.2416 - val_loss: 3.6451 - val_acc: 0.2672
Epoch 17/250
- 74s - loss: 3.7394 - acc: 0.2618 - val_loss: 3.8330 - val_acc: 0.2626
Epoch 18/250
- 73s - loss: 3.7516 - acc: 0.2750 - val_loss: 4.0233 - val_acc: 0.2538
Epoch 19/250
- 74s - loss: 3.7650 - acc: 0.2882 - val_loss: 3.7117 - val_acc: 0.3134
Epoch 20/250
- 73s - loss: 3.7587 - acc: 0.3033 - val_loss: 3.5769 - val_acc: 0.3408
Epoch 21/250
- 73s - loss: 3.4489 - acc: 0.3568 - val_loss: 3.1665 - val_acc: 0.4040
Epoch 22/250
- 73s - loss: 3.3070 - acc: 0.3729 - val_loss: 3.2326 - val_acc: 0.3806
Epoch 23/250
- 73s - loss: 3.2784 - acc: 0.3772 - val_loss: 3.2470 - val_acc: 0.3956
Epoch 24/250
- 73s - loss: 3.2619 - acc: 0.3846 - val_loss: 3.2542 - val_acc: 0.3960
Epoch 25/250
- 73s - loss: 3.2498 - acc: 0.3921 - val_loss: 3.1304 - val_acc: 0.4146
Epoch 26/250
- 73s - loss: 3.2601 - acc: 0.3981 - val_loss: 3.2662 - val_acc: 0.4084
Epoch 27/250
- 73s - loss: 3.2591 - acc: 0.4040 - val_loss: 3.2157 - val_acc: 0.4259
Epoch 28/250
- 73s - loss: 3.2684 - acc: 0.4112 - val_loss: 3.2219 - val_acc: 0.4337
Epoch 29/250
- 73s - loss: 3.2704 - acc: 0.4167 - val_loss: 3.0671 - val_acc: 0.4620
```

# Machine Learning Project Report

```
Epoch 30/250
- 74s - loss: 3.2665 - acc: 0.4218 - val_loss: 3.2299 - val_acc: 0.4368
Epoch 31/250
- 73s - loss: 3.2959 - acc: 0.4193 - val_loss: 3.4639 - val_acc: 0.4091
Epoch 32/250
- 73s - loss: 3.2896 - acc: 0.4276 - val_loss: 3.1511 - val_acc: 0.4592
Epoch 33/250
- 73s - loss: 3.3003 - acc: 0.4325 - val_loss: 3.2044 - val_acc: 0.4548
Epoch 34/250
- 74s - loss: 3.3017 - acc: 0.4349 - val_loss: 3.3354 - val_acc: 0.4356
Epoch 35/250
- 74s - loss: 3.2970 - acc: 0.4390 - val_loss: 3.3408 - val_acc: 0.4384
Epoch 36/250
- 73s - loss: 3.3114 - acc: 0.4416 - val_loss: 3.0754 - val_acc: 0.4999
Epoch 37/250
- 73s - loss: 3.3192 - acc: 0.4452 - val_loss: 3.1470 - val_acc: 0.4805
Epoch 38/250
- 73s - loss: 3.3101 - acc: 0.4482 - val_loss: 3.3261 - val_acc: 0.4564
Epoch 39/250
- 73s - loss: 3.3244 - acc: 0.4500 - val_loss: 3.4501 - val_acc: 0.4480
Epoch 40/250
- 73s - loss: 3.3403 - acc: 0.4505 - val_loss: 3.2798 - val_acc: 0.4652
Epoch 41/250
- 73s - loss: 3.0713 - acc: 0.5010 - val_loss: 2.9705 - val_acc: 0.5149
Epoch 42/250
- 73s - loss: 2.9162 - acc: 0.5179 - val_loss: 2.8526 - val_acc: 0.5327
Epoch 43/250
- 73s - loss: 2.8495 - acc: 0.5211 - val_loss: 2.9520 - val_acc: 0.5133
Epoch 44/250
- 73s - loss: 2.8191 - acc: 0.5236 - val_loss: 2.8670 - val_acc: 0.5208
Epoch 45/250
- 73s - loss: 2.7979 - acc: 0.5268 - val_loss: 2.8340 - val_acc: 0.5234
Epoch 46/250
- 73s - loss: 2.7886 - acc: 0.5246 - val_loss: 2.8433 - val_acc: 0.5189
Epoch 47/250
- 73s - loss: 2.7797 - acc: 0.5278 - val_loss: 2.8080 - val_acc: 0.5303
Epoch 48/250
- 73s - loss: 2.7754 - acc: 0.5323 - val_loss: 2.9752 - val_acc: 0.5147
Epoch 49/250
- 73s - loss: 2.7869 - acc: 0.5325 - val_loss: 3.0404 - val_acc: 0.5049
Epoch 50/250
- 73s - loss: 2.7783 - acc: 0.5335 - val_loss: 2.7255 - val_acc: 0.5532
Epoch 51/250
- 73s - loss: 2.7931 - acc: 0.5367 - val_loss: 2.9087 - val_acc: 0.5200
Epoch 52/250
- 73s - loss: 2.7913 - acc: 0.5372 - val_loss: 2.7988 - val_acc: 0.5467
Epoch 53/250
- 73s - loss: 2.8003 - acc: 0.5372 - val_loss: 2.7715 - val_acc: 0.5470
Epoch 54/250
- 73s - loss: 2.7951 - acc: 0.5438 - val_loss: 2.7716 - val_acc: 0.5554
Epoch 55/250
- 73s - loss: 2.7870 - acc: 0.5477 - val_loss: 2.7653 - val_acc: 0.5580
Epoch 56/250
- 73s - loss: 2.7956 - acc: 0.5457 - val_loss: 2.8150 - val_acc: 0.5511
Epoch 57/250
```

# Machine Learning Project Report

```
- 73s - loss: 2.7982 - acc: 0.5450 - val_loss: 2.7758 - val_acc: 0.5561
Epoch 58/250
- 73s - loss: 2.8016 - acc: 0.5493 - val_loss: 2.8002 - val_acc: 0.5613
Epoch 59/250
- 73s - loss: 2.8088 - acc: 0.5479 - val_loss: 2.8056 - val_acc: 0.5558
Epoch 60/250
- 73s - loss: 2.7991 - acc: 0.5548 - val_loss: 2.8858 - val_acc: 0.5451
Epoch 61/250
- 73s - loss: 2.5989 - acc: 0.5985 - val_loss: 2.5901 - val_acc: 0.5985
Epoch 62/250
- 73s - loss: 2.4865 - acc: 0.6109 - val_loss: 2.4905 - val_acc: 0.6169
Epoch 63/250
- 73s - loss: 2.4425 - acc: 0.6173 - val_loss: 2.4787 - val_acc: 0.6115
Epoch 64/250
- 73s - loss: 2.3959 - acc: 0.6200 - val_loss: 2.5190 - val_acc: 0.5951
Epoch 65/250
- 73s - loss: 2.3715 - acc: 0.6222 - val_loss: 2.4943 - val_acc: 0.5986
Epoch 66/250
- 73s - loss: 2.3595 - acc: 0.6208 - val_loss: 2.4443 - val_acc: 0.6034
Epoch 67/250
- 73s - loss: 2.3313 - acc: 0.6250 - val_loss: 2.5064 - val_acc: 0.5917
Epoch 68/250
- 73s - loss: 2.3354 - acc: 0.6216 - val_loss: 2.4027 - val_acc: 0.6106
Epoch 69/250
- 73s - loss: 2.3271 - acc: 0.6238 - val_loss: 2.4977 - val_acc: 0.5942
Epoch 70/250
- 73s - loss: 2.3222 - acc: 0.6241 - val_loss: 2.6475 - val_acc: 0.5763
Epoch 71/250
- 73s - loss: 2.3147 - acc: 0.6260 - val_loss: 2.5133 - val_acc: 0.5994
Epoch 72/250
- 74s - loss: 2.3097 - acc: 0.6267 - val_loss: 2.4364 - val_acc: 0.6057
Epoch 73/250
- 73s - loss: 2.3119 - acc: 0.6225 - val_loss: 2.4700 - val_acc: 0.6061
Epoch 74/250
- 73s - loss: 2.3128 - acc: 0.6267 - val_loss: 2.5312 - val_acc: 0.5994
Epoch 75/250
- 73s - loss: 2.3132 - acc: 0.6291 - val_loss: 2.6506 - val_acc: 0.5652
Epoch 76/250
- 73s - loss: 2.3169 - acc: 0.6287 - val_loss: 2.5014 - val_acc: 0.5954
Epoch 77/250
- 73s - loss: 2.3172 - acc: 0.6297 - val_loss: 2.5622 - val_acc: 0.5902
Epoch 78/250
- 73s - loss: 2.3083 - acc: 0.6334 - val_loss: 2.5386 - val_acc: 0.5934
Epoch 79/250
- 73s - loss: 2.3084 - acc: 0.6328 - val_loss: 2.5255 - val_acc: 0.6009
Epoch 80/250
- 73s - loss: 2.3118 - acc: 0.6344 - val_loss: 2.5547 - val_acc: 0.5937
Epoch 81/250
- 73s - loss: 2.1604 - acc: 0.6699 - val_loss: 2.4328 - val_acc: 0.6218
Epoch 82/250
- 73s - loss: 2.0735 - acc: 0.6849 - val_loss: 2.3229 - val_acc: 0.6355
Epoch 83/250
- 73s - loss: 2.0313 - acc: 0.6904 - val_loss: 2.3480 - val_acc: 0.6333
Epoch 84/250
- 73s - loss: 1.9983 - acc: 0.6979 - val_loss: 2.2584 - val_acc: 0.6497
```

# Machine Learning Project Report

```
Epoch 85/250
- 73s - loss: 1.9711 - acc: 0.6975 - val_loss: 2.3124 - val_acc: 0.6397
Epoch 86/250
- 73s - loss: 1.9632 - acc: 0.6962 - val_loss: 2.2306 - val_acc: 0.6476
Epoch 87/250
- 73s - loss: 1.9546 - acc: 0.6973 - val_loss: 2.3628 - val_acc: 0.6294
Epoch 88/250
- 73s - loss: 1.9313 - acc: 0.7009 - val_loss: 2.2666 - val_acc: 0.6441
Epoch 89/250
- 73s - loss: 1.9256 - acc: 0.6995 - val_loss: 2.2452 - val_acc: 0.6386
Epoch 90/250
- 73s - loss: 1.9063 - acc: 0.7011 - val_loss: 2.2394 - val_acc: 0.6416
Epoch 91/250
- 73s - loss: 1.9031 - acc: 0.7024 - val_loss: 2.3344 - val_acc: 0.6254
Epoch 92/250
- 73s - loss: 1.9011 - acc: 0.7022 - val_loss: 2.2714 - val_acc: 0.6357
Epoch 93/250
- 73s - loss: 1.8995 - acc: 0.7032 - val_loss: 2.1884 - val_acc: 0.6479
Epoch 94/250
- 73s - loss: 1.8854 - acc: 0.7027 - val_loss: 2.2678 - val_acc: 0.6275
Epoch 95/250
- 73s - loss: 1.8807 - acc: 0.7071 - val_loss: 2.2454 - val_acc: 0.6316
Epoch 96/250
- 73s - loss: 1.8867 - acc: 0.7032 - val_loss: 2.2106 - val_acc: 0.6476
Epoch 97/250
- 73s - loss: 1.8607 - acc: 0.7105 - val_loss: 2.2614 - val_acc: 0.6425
Epoch 98/250
- 73s - loss: 1.8807 - acc: 0.7037 - val_loss: 2.3134 - val_acc: 0.6326
Epoch 99/250
- 73s - loss: 1.8754 - acc: 0.7070 - val_loss: 2.1912 - val_acc: 0.6494
Epoch 100/250
- 73s - loss: 1.8787 - acc: 0.7048 - val_loss: 2.2177 - val_acc: 0.6452
Epoch 101/250
- 73s - loss: 1.7557 - acc: 0.7363 - val_loss: 2.1566 - val_acc: 0.6594
Epoch 102/250
- 73s - loss: 1.6833 - acc: 0.7535 - val_loss: 2.1369 - val_acc: 0.6666
Epoch 103/250
- 73s - loss: 1.6808 - acc: 0.7510 - val_loss: 2.0935 - val_acc: 0.6723
Epoch 104/250
- 73s - loss: 1.6459 - acc: 0.7574 - val_loss: 2.1751 - val_acc: 0.6586
Epoch 105/250
- 73s - loss: 1.6317 - acc: 0.7587 - val_loss: 2.1415 - val_acc: 0.6631
Epoch 106/250
- 73s - loss: 1.6106 - acc: 0.7617 - val_loss: 2.0827 - val_acc: 0.6772
Epoch 107/250
- 73s - loss: 1.5922 - acc: 0.7649 - val_loss: 2.1169 - val_acc: 0.6669
Epoch 108/250
- 73s - loss: 1.5760 - acc: 0.7653 - val_loss: 2.1186 - val_acc: 0.6724
Epoch 109/250
- 73s - loss: 1.5822 - acc: 0.7660 - val_loss: 2.0705 - val_acc: 0.6755
Epoch 110/250
- 73s - loss: 1.5715 - acc: 0.7637 - val_loss: 2.1859 - val_acc: 0.6564
Epoch 111/250
- 73s - loss: 1.5594 - acc: 0.7686 - val_loss: 2.1140 - val_acc: 0.6670
Epoch 112/250
```

# Machine Learning Project Report

```
- 73s - loss: 1.5567 - acc: 0.7684 - val_loss: 2.0638 - val_acc: 0.6737
Epoch 113/250
- 73s - loss: 1.5445 - acc: 0.7692 - val_loss: 2.1003 - val_acc: 0.6718
Epoch 114/250
- 73s - loss: 1.5482 - acc: 0.7685 - val_loss: 2.0253 - val_acc: 0.6802
Epoch 115/250
- 73s - loss: 1.5348 - acc: 0.7716 - val_loss: 2.0756 - val_acc: 0.6733
Epoch 116/250
- 73s - loss: 1.5350 - acc: 0.7684 - val_loss: 2.0520 - val_acc: 0.6721
Epoch 117/250
- 73s - loss: 1.5145 - acc: 0.7742 - val_loss: 2.1025 - val_acc: 0.6659
Epoch 118/250
- 73s - loss: 1.5180 - acc: 0.7712 - val_loss: 2.0321 - val_acc: 0.6779
Epoch 119/250
- 73s - loss: 1.5132 - acc: 0.7734 - val_loss: 2.0939 - val_acc: 0.6685
Epoch 120/250
- 73s - loss: 1.5153 - acc: 0.7724 - val_loss: 2.0729 - val_acc: 0.6669
Epoch 121/250
- 73s - loss: 1.4355 - acc: 0.7935 - val_loss: 2.0412 - val_acc: 0.6778
Epoch 122/250
- 73s - loss: 1.3951 - acc: 0.8011 - val_loss: 2.0376 - val_acc: 0.6803
Epoch 123/250
- 73s - loss: 1.3620 - acc: 0.8098 - val_loss: 2.0401 - val_acc: 0.6809
Epoch 124/250
- 73s - loss: 1.3548 - acc: 0.8108 - val_loss: 2.0352 - val_acc: 0.6819
Epoch 125/250
- 73s - loss: 1.3386 - acc: 0.8119 - val_loss: 2.0281 - val_acc: 0.6830
Epoch 126/250
- 73s - loss: 1.3282 - acc: 0.8145 - val_loss: 2.0856 - val_acc: 0.6759
Epoch 127/250
- 73s - loss: 1.3227 - acc: 0.8154 - val_loss: 2.0566 - val_acc: 0.6814
Epoch 128/250
- 73s - loss: 1.3068 - acc: 0.8177 - val_loss: 2.0759 - val_acc: 0.6772
Epoch 129/250
- 73s - loss: 1.3063 - acc: 0.8181 - val_loss: 2.0966 - val_acc: 0.6778
Epoch 130/250
- 73s - loss: 1.2926 - acc: 0.8203 - val_loss: 2.0357 - val_acc: 0.6882
Epoch 131/250
- 73s - loss: 1.2872 - acc: 0.8182 - val_loss: 2.0345 - val_acc: 0.6822
Epoch 132/250
- 73s - loss: 1.2785 - acc: 0.8219 - val_loss: 2.0800 - val_acc: 0.6795
Epoch 133/250
- 73s - loss: 1.2691 - acc: 0.8228 - val_loss: 2.0819 - val_acc: 0.6753
Epoch 134/250
- 73s - loss: 1.2635 - acc: 0.8214 - val_loss: 2.0667 - val_acc: 0.6798
Epoch 135/250
- 73s - loss: 1.2646 - acc: 0.8235 - val_loss: 2.0518 - val_acc: 0.6799
Epoch 136/250
- 73s - loss: 1.2592 - acc: 0.8226 - val_loss: 2.0559 - val_acc: 0.6846
Epoch 137/250
- 73s - loss: 1.2488 - acc: 0.8251 - val_loss: 2.1136 - val_acc: 0.6710
Epoch 138/250
- 73s - loss: 1.2429 - acc: 0.8283 - val_loss: 2.1088 - val_acc: 0.6720
Epoch 139/250
- 73s - loss: 1.2355 - acc: 0.8265 - val_loss: 2.0294 - val_acc: 0.6889
```



# Machine Learning Project Report

```
Epoch 140/250
- 73s - loss: 1.2234 - acc: 0.8292 - val_loss: 2.0956 - val_acc: 0.6744
Epoch 141/250
- 73s - loss: 1.1887 - acc: 0.8408 - val_loss: 2.0380 - val_acc: 0.6887
Epoch 142/250
- 73s - loss: 1.1542 - acc: 0.8477 - val_loss: 2.0185 - val_acc: 0.6891
Epoch 143/250
- 73s - loss: 1.1408 - acc: 0.8509 - val_loss: 2.0109 - val_acc: 0.6901
Epoch 144/250
- 73s - loss: 1.1384 - acc: 0.8512 - val_loss: 2.0416 - val_acc: 0.6872
Epoch 145/250
- 73s - loss: 1.1208 - acc: 0.8545 - val_loss: 2.0410 - val_acc: 0.6871
Epoch 146/250
- 73s - loss: 1.1176 - acc: 0.8569 - val_loss: 2.0464 - val_acc: 0.6859
Epoch 147/250
- 73s - loss: 1.1197 - acc: 0.8551 - val_loss: 2.0523 - val_acc: 0.6852
Epoch 148/250
- 73s - loss: 1.1083 - acc: 0.8582 - val_loss: 2.0826 - val_acc: 0.6821
Epoch 149/250
- 73s - loss: 1.0997 - acc: 0.8586 - val_loss: 2.1030 - val_acc: 0.6778
Epoch 150/250
- 73s - loss: 1.0972 - acc: 0.8603 - val_loss: 2.0389 - val_acc: 0.6916
Epoch 151/250
- 73s - loss: 1.0853 - acc: 0.8623 - val_loss: 2.0502 - val_acc: 0.6891
Epoch 152/250
- 73s - loss: 1.0853 - acc: 0.8593 - val_loss: 2.0763 - val_acc: 0.6856
Epoch 153/250
- 73s - loss: 1.0818 - acc: 0.8614 - val_loss: 2.0555 - val_acc: 0.6889
Epoch 154/250
- 73s - loss: 1.0752 - acc: 0.8626 - val_loss: 2.0450 - val_acc: 0.6897
Epoch 155/250
- 73s - loss: 1.0665 - acc: 0.8637 - val_loss: 2.0505 - val_acc: 0.6876
Epoch 156/250
- 73s - loss: 1.0677 - acc: 0.8623 - val_loss: 2.0334 - val_acc: 0.6909
Epoch 157/250
- 73s - loss: 1.0617 - acc: 0.8625 - val_loss: 2.0517 - val_acc: 0.6906
Epoch 158/250
- 73s - loss: 1.0479 - acc: 0.8696 - val_loss: 2.0689 - val_acc: 0.6903
Epoch 159/250
- 73s - loss: 1.0501 - acc: 0.8669 - val_loss: 2.0759 - val_acc: 0.6834
Epoch 160/250
- 73s - loss: 1.0438 - acc: 0.8664 - val_loss: 2.0530 - val_acc: 0.6874
Epoch 161/250
- 73s - loss: 1.0284 - acc: 0.8714 - val_loss: 2.0446 - val_acc: 0.6909
Epoch 162/250
- 73s - loss: 1.0062 - acc: 0.8782 - val_loss: 2.0319 - val_acc: 0.6908
Epoch 163/250
- 74s - loss: 0.9989 - acc: 0.8793 - val_loss: 2.0566 - val_acc: 0.6890
Epoch 164/250
- 73s - loss: 0.9990 - acc: 0.8783 - val_loss: 2.0409 - val_acc: 0.6952
Epoch 165/250
- 73s - loss: 0.9926 - acc: 0.8806 - val_loss: 2.0428 - val_acc: 0.6915
Epoch 166/250
- 73s - loss: 0.9866 - acc: 0.8817 - val_loss: 2.0332 - val_acc: 0.6942
Epoch 167/250
```

# Machine Learning Project Report

```
- 73s - loss: 0.9874 - acc: 0.8813 - val_loss: 2.0519 - val_acc: 0.6927
Epoch 168/250
- 73s - loss: 0.9714 - acc: 0.8846 - val_loss: 2.0595 - val_acc: 0.6881
Epoch 169/250
- 73s - loss: 0.9745 - acc: 0.8823 - val_loss: 2.0393 - val_acc: 0.6941
Epoch 170/250
- 73s - loss: 0.9666 - acc: 0.8872 - val_loss: 2.0530 - val_acc: 0.6920
Epoch 171/250
- 73s - loss: 0.9605 - acc: 0.8877 - val_loss: 2.0461 - val_acc: 0.6921
Epoch 172/250
- 73s - loss: 0.9627 - acc: 0.8845 - val_loss: 2.0504 - val_acc: 0.6938
Epoch 173/250
- 73s - loss: 0.9614 - acc: 0.8862 - val_loss: 2.0626 - val_acc: 0.6900
Epoch 174/250
- 73s - loss: 0.9511 - acc: 0.8900 - val_loss: 2.0691 - val_acc: 0.6925
Epoch 175/250
- 73s - loss: 0.9515 - acc: 0.8878 - val_loss: 2.0751 - val_acc: 0.6906
Epoch 176/250
- 73s - loss: 0.9496 - acc: 0.8900 - val_loss: 2.0552 - val_acc: 0.6939
Epoch 177/250
- 73s - loss: 0.9542 - acc: 0.8865 - val_loss: 2.0670 - val_acc: 0.6943
Epoch 178/250
- 73s - loss: 0.9369 - acc: 0.8921 - val_loss: 2.0627 - val_acc: 0.6932
Epoch 179/250
- 73s - loss: 0.9461 - acc: 0.8881 - val_loss: 2.0726 - val_acc: 0.6866
Epoch 180/250
- 73s - loss: 0.9398 - acc: 0.8903 - val_loss: 2.0425 - val_acc: 0.6942
Epoch 181/250
- 73s - loss: 0.9235 - acc: 0.8941 - val_loss: 2.0549 - val_acc: 0.6923
Epoch 182/250
- 73s - loss: 0.9316 - acc: 0.8914 - val_loss: 2.0492 - val_acc: 0.6921
Epoch 183/250
- 73s - loss: 0.9168 - acc: 0.8949 - val_loss: 2.0620 - val_acc: 0.6923
Epoch 184/250
- 73s - loss: 0.9110 - acc: 0.8979 - val_loss: 2.0660 - val_acc: 0.6901
Epoch 185/250
- 73s - loss: 0.9201 - acc: 0.8960 - val_loss: 2.0621 - val_acc: 0.6942
Epoch 186/250
- 73s - loss: 0.9083 - acc: 0.8978 - val_loss: 2.0575 - val_acc: 0.6922
Epoch 187/250
- 73s - loss: 0.9029 - acc: 0.8997 - val_loss: 2.0565 - val_acc: 0.6929
Epoch 188/250
- 73s - loss: 0.9040 - acc: 0.8993 - val_loss: 2.0582 - val_acc: 0.6923
Epoch 189/250
- 73s - loss: 0.9057 - acc: 0.8977 - val_loss: 2.0653 - val_acc: 0.6896
Epoch 190/250
- 73s - loss: 0.9043 - acc: 0.8978 - val_loss: 2.0591 - val_acc: 0.6922
Epoch 191/250
- 73s - loss: 0.8897 - acc: 0.9026 - val_loss: 2.0576 - val_acc: 0.6940
Epoch 192/250
- 73s - loss: 0.8946 - acc: 0.8996 - val_loss: 2.0679 - val_acc: 0.6921
Epoch 193/250
- 73s - loss: 0.9027 - acc: 0.8987 - val_loss: 2.0553 - val_acc: 0.6944
Epoch 194/250
- 73s - loss: 0.8950 - acc: 0.9011 - val_loss: 2.0702 - val_acc: 0.6924
```

# Machine Learning Project Report

```
Epoch 195/250
- 73s - loss: 0.8905 - acc: 0.9019 - val_loss: 2.0743 - val_acc: 0.6921
Epoch 196/250
- 73s - loss: 0.8875 - acc: 0.9016 - val_loss: 2.0507 - val_acc: 0.6933
Epoch 197/250
- 73s - loss: 0.8897 - acc: 0.9006 - val_loss: 2.0591 - val_acc: 0.6935
Epoch 198/250
- 73s - loss: 0.8825 - acc: 0.9030 - val_loss: 2.0673 - val_acc: 0.6926
Epoch 199/250
- 73s - loss: 0.8834 - acc: 0.9036 - val_loss: 2.0717 - val_acc: 0.6921
Epoch 200/250
- 73s - loss: 0.8822 - acc: 0.9024 - val_loss: 2.0756 - val_acc: 0.6916
Epoch 201/250
- 73s - loss: 0.8742 - acc: 0.9036 - val_loss: 2.0609 - val_acc: 0.6944
Epoch 202/250
- 73s - loss: 0.8804 - acc: 0.9025 - val_loss: 2.0607 - val_acc: 0.6940
Epoch 203/250
- 73s - loss: 0.8745 - acc: 0.9049 - val_loss: 2.0608 - val_acc: 0.6949
Epoch 204/250
- 73s - loss: 0.8723 - acc: 0.9044 - val_loss: 2.0632 - val_acc: 0.6926
Epoch 205/250
- 73s - loss: 0.8644 - acc: 0.9077 - val_loss: 2.0600 - val_acc: 0.6944
Epoch 206/250
- 73s - loss: 0.8625 - acc: 0.9077 - val_loss: 2.0795 - val_acc: 0.6918
Epoch 207/250
- 74s - loss: 0.8642 - acc: 0.9071 - val_loss: 2.0698 - val_acc: 0.6941
Epoch 208/250
- 73s - loss: 0.8649 - acc: 0.9065 - val_loss: 2.0664 - val_acc: 0.6942
Epoch 209/250
- 73s - loss: 0.8636 - acc: 0.9071 - val_loss: 2.0661 - val_acc: 0.6937
Epoch 210/250
- 73s - loss: 0.8586 - acc: 0.9078 - val_loss: 2.0681 - val_acc: 0.6941
Epoch 211/250
- 73s - loss: 0.8586 - acc: 0.9078 - val_loss: 2.0738 - val_acc: 0.6935
Epoch 212/250
- 73s - loss: 0.8573 - acc: 0.9096 - val_loss: 2.0687 - val_acc: 0.6934
Epoch 213/250
- 73s - loss: 0.8570 - acc: 0.9086 - val_loss: 2.0746 - val_acc: 0.6942
Epoch 214/250
- 73s - loss: 0.8608 - acc: 0.9077 - val_loss: 2.0659 - val_acc: 0.6956
Epoch 215/250
- 73s - loss: 0.8508 - acc: 0.9100 - val_loss: 2.0703 - val_acc: 0.6950
Epoch 216/250
- 73s - loss: 0.8513 - acc: 0.9117 - val_loss: 2.0774 - val_acc: 0.6947
Epoch 217/250
- 73s - loss: 0.8537 - acc: 0.9089 - val_loss: 2.0736 - val_acc: 0.6938
Epoch 218/250
- 73s - loss: 0.8532 - acc: 0.9088 - val_loss: 2.0682 - val_acc: 0.6958
Epoch 219/250
- 73s - loss: 0.8481 - acc: 0.9124 - val_loss: 2.0846 - val_acc: 0.6945
Epoch 220/250
- 73s - loss: 0.8549 - acc: 0.9069 - val_loss: 2.0746 - val_acc: 0.6938
Epoch 221/250
- 73s - loss: 0.8531 - acc: 0.9086 - val_loss: 2.0784 - val_acc: 0.6931
Epoch 222/250
```

# Machine Learning Project Report

- 73s - loss: 0.8440 - acc: 0.9113 - val\_loss: 2.0791 - val\_acc: 0.6928  
Epoch 223/250  
- 73s - loss: 0.8486 - acc: 0.9102 - val\_loss: 2.0762 - val\_acc: 0.6929  
Epoch 224/250  
- 73s - loss: 0.8518 - acc: 0.9097 - val\_loss: 2.0737 - val\_acc: 0.6955  
Epoch 225/250  
- 73s - loss: 0.8480 - acc: 0.9109 - val\_loss: 2.0711 - val\_acc: 0.6952  
Epoch 226/250  
- 73s - loss: 0.8508 - acc: 0.9091 - val\_loss: 2.0730 - val\_acc: 0.6951  
Epoch 227/250  
- 73s - loss: 0.8401 - acc: 0.9124 - val\_loss: 2.0784 - val\_acc: 0.6934  
Epoch 228/250  
- 73s - loss: 0.8425 - acc: 0.9111 - val\_loss: 2.0667 - val\_acc: 0.6944  
Epoch 229/250  
- 73s - loss: 0.8495 - acc: 0.9098 - val\_loss: 2.0741 - val\_acc: 0.6941  
Epoch 230/250  
- 73s - loss: 0.8408 - acc: 0.9130 - val\_loss: 2.0675 - val\_acc: 0.6950  
Epoch 231/250  
- 73s - loss: 0.8484 - acc: 0.9099 - val\_loss: 2.0703 - val\_acc: 0.6946  
Epoch 232/250  
- 73s - loss: 0.8401 - acc: 0.9122 - val\_loss: 2.0664 - val\_acc: 0.6942  
Epoch 233/250  
- 73s - loss: 0.8380 - acc: 0.9136 - val\_loss: 2.0768 - val\_acc: 0.6922  
Epoch 234/250  
- 73s - loss: 0.8482 - acc: 0.9086 - val\_loss: 2.0791 - val\_acc: 0.6914  
Epoch 235/250  
- 73s - loss: 0.8383 - acc: 0.9128 - val\_loss: 2.0788 - val\_acc: 0.6931  
Epoch 236/250  
- 73s - loss: 0.8426 - acc: 0.9116 - val\_loss: 2.0740 - val\_acc: 0.6937  
Epoch 237/250  
- 73s - loss: 0.8394 - acc: 0.9108 - val\_loss: 2.0802 - val\_acc: 0.6935  
Epoch 238/250  
- 73s - loss: 0.8446 - acc: 0.9108 - val\_loss: 2.0764 - val\_acc: 0.6938  
Epoch 239/250  
- 73s - loss: 0.8419 - acc: 0.9108 - val\_loss: 2.0730 - val\_acc: 0.6938  
Epoch 240/250  
- 73s - loss: 0.8418 - acc: 0.9104 - val\_loss: 2.0760 - val\_acc: 0.6944  
Epoch 241/250  
- 73s - loss: 0.8368 - acc: 0.9131 - val\_loss: 2.0734 - val\_acc: 0.6937  
Epoch 242/250  
- 73s - loss: 0.8344 - acc: 0.9134 - val\_loss: 2.0655 - val\_acc: 0.6953  
Epoch 243/250  
- 73s - loss: 0.8363 - acc: 0.9139 - val\_loss: 2.0730 - val\_acc: 0.6936  
Epoch 244/250  
- 73s - loss: 0.8360 - acc: 0.9122 - val\_loss: 2.0668 - val\_acc: 0.6948  
Epoch 245/250  
- 73s - loss: 0.8318 - acc: 0.9128 - val\_loss: 2.0664 - val\_acc: 0.6987  
Epoch 246/250  
- 73s - loss: 0.8331 - acc: 0.9136 - val\_loss: 2.0699 - val\_acc: 0.7009  
Epoch 247/250  
- 73s - loss: 0.8328 - acc: 0.9134 - val\_loss: 2.0686 - val\_acc: 0.7021  
Epoch 248/250  
- 73s - loss: 0.8311 - acc: 0.9133 - val\_loss: 2.0716 - val\_acc: 0.7045  
Epoch 249/250  
- 73s - loss: 0.8356 - acc: 0.9131 - val\_loss: 2.0731 - val\_acc: 0.7065

# Machine Learning Project Report

Epoch 250/250  
- 73s - loss: 0.8291 - acc: 0.9137 - val\_loss: 2.0754 - val\_acc: 0.7088  
Training Time : 18284.95653295517  
Accuracy: 70.88 %

Run 2:

Using TensorFlow backend.

The default version of TensorFlow in Colab will soon switch to TensorFlow 2.x.

We recommend you [upgrade](#) now or ensure your notebook will continue to use TensorFlow 1.x via the `%tensorflow_version 1.x` magic: [more info](#).

Downloading data from <https://www.cs.toronto.edu/~kriz/cifar-100-python.tar.gz>

169009152/169001437 [=====] - 11s 0us/step

WARNING:tensorflow:From /usr/local/lib/python3.6/dist-packages/keras/backend/tensorflow\_backend.py:66: The name tf.get\_default\_graph is deprecated. Please use tf.compat.v1.get\_default\_graph instead.

WARNING:tensorflow:From /usr/local/lib/python3.6/dist-packages/keras/backend/tensorflow\_backend.py:541: The name tf.placeholder is deprecated. Please use tf.compat.v1.placeholder instead.

WARNING:tensorflow:From /usr/local/lib/python3.6/dist-packages/keras/backend/tensorflow\_backend.py:4432: The name tf.random\_uniform is deprecated. Please use tf.random.uniform instead.

WARNING:tensorflow:From /usr/local/lib/python3.6/dist-packages/keras/backend/tensorflow\_backend.py:190: The name tf.get\_default\_session is deprecated. Please use tf.compat.v1.get\_default\_session instead.

WARNING:tensorflow:From /usr/local/lib/python3.6/dist-packages/keras/backend/tensorflow\_backend.py:197: The name tf.ConfigProto is deprecated. Please use tf.compat.v1.ConfigProto instead.

WARNING:tensorflow:From /usr/local/lib/python3.6/dist-packages/keras/backend/tensorflow\_backend.py:203: The name tf.Session is deprecated. Please use tf.compat.v1.Session instead.

WARNING:tensorflow:From /usr/local/lib/python3.6/dist-packages/keras/backend/tensorflow\_backend.py:207: The name tf.global\_variables is deprecated. Please use tf.compat.v1.global\_variables instead.

WARNING:tensorflow:From /usr/local/lib/python3.6/dist-packages/keras/backend/tensorflow\_backend.py:216: The name tf.is\_variable\_initialized is deprecated. Please use tf.compat.v1.is\_variable\_initialized instead.

# Machine Learning Project Report

WARNING:tensorflow:From /usr/local/lib/python3.6/dist-packages/keras/backend/tensorflow\_backend.py:223: The name tf.variables\_initializer is deprecated. Please use tf.compat.v1.variables\_initializer instead.

WARNING:tensorflow:From /usr/local/lib/python3.6/dist-packages/keras/backend/tensorflow\_backend.py:2041: The name tf.nn.fused\_batch\_norm is deprecated. Please use tf.compat.v1.nn.fused\_batch\_norm instead.

WARNING:tensorflow:From /usr/local/lib/python3.6/dist-packages/keras/backend/tensorflow\_backend.py:148: The name tf.placeholder\_with\_default is deprecated. Please use tf.compat.v1.placeholder\_with\_default instead.

WARNING:tensorflow:From /usr/local/lib/python3.6/dist-packages/keras/backend/tensorflow\_backend.py:3733: calling dropout (from tensorflow.python.ops.nn\_ops) with keep\_prob is deprecated and will be removed in a future version.  
Instructions for updating:  
Please use `rate` instead of `keep\_prob`. Rate should be set to `rate = 1 - keep\_prob`.  
WARNING:tensorflow:From /usr/local/lib/python3.6/dist-packages/keras/backend/tensorflow\_backend.py:4267: The name tf.nn.max\_pool is deprecated. Please use tf.nn.max\_pool2d instead.

121.93584  
68.38902

WARNING:tensorflow:From /usr/local/lib/python3.6/dist-packages/keras/optimizers.py:793: The name tf.train.Optimizer is deprecated. Please use tf.compat.v1.train.Optimizer instead.

WARNING:tensorflow:From /usr/local/lib/python3.6/dist-packages/keras/backend/tensorflow\_backend.py:3576: The name tf.log is deprecated. Please use tf.math.log instead.

WARNING:tensorflow:From /usr/local/lib/python3.6/dist-packages/tensorflow\_core/python/ops/math\_grad.py:1424: where (from tensorflow.python.ops.array\_ops) is deprecated and will be removed in a future version.  
Instructions for updating:  
Use tf.where in 2.0, which has the same broadcast rule as np.where  
WARNING:tensorflow:From /usr/local/lib/python3.6/dist-packages/keras/backend/tensorflow\_backend.py:1033: The name tf.assign\_add is deprecated. Please use tf.compat.v1.assign\_add instead.

WARNING:tensorflow:From /usr/local/lib/python3.6/dist-packages/keras/backend/tensorflow\_backend.py:1020: The name tf.assign is deprecated. Please use tf.compat.v1.assign instead.

Epoch 1/250  
- 34s - loss: 19.7685 - acc: 0.0266 - val\_loss: 24.3936 - val\_acc: 0.0053

# Machine Learning Project Report

```
Epoch 2/250
- 26s - loss: 11.7540 - acc: 0.0481 - val_loss: 12.9683 - val_acc: 0.0157
Epoch 3/250
- 26s - loss: 8.5206 - acc: 0.0419 - val_loss: 7.5412 - val_acc: 0.0343
Epoch 4/250
- 26s - loss: 6.7142 - acc: 0.0549 - val_loss: 6.5740 - val_acc: 0.0369
Epoch 5/250
- 26s - loss: 5.7953 - acc: 0.0580 - val_loss: 6.0052 - val_acc: 0.0393
Epoch 6/250
- 26s - loss: 5.1539 - acc: 0.0710 - val_loss: 5.3819 - val_acc: 0.0266
Epoch 7/250
- 26s - loss: 4.5888 - acc: 0.0855 - val_loss: 4.9524 - val_acc: 0.0399
Epoch 8/250
- 26s - loss: 4.3123 - acc: 0.0993 - val_loss: 4.8184 - val_acc: 0.0689
Epoch 9/250
- 26s - loss: 4.1317 - acc: 0.1112 - val_loss: 4.2757 - val_acc: 0.0939
Epoch 10/250
- 26s - loss: 4.0376 - acc: 0.1270 - val_loss: 4.1955 - val_acc: 0.1098
Epoch 11/250
- 26s - loss: 3.9728 - acc: 0.1426 - val_loss: 3.9241 - val_acc: 0.1540
Epoch 12/250
- 26s - loss: 3.9039 - acc: 0.1607 - val_loss: 3.8912 - val_acc: 0.1668
Epoch 13/250
- 26s - loss: 3.8229 - acc: 0.1861 - val_loss: 3.7463 - val_acc: 0.2150
Epoch 14/250
- 26s - loss: 3.7817 - acc: 0.2100 - val_loss: 3.6448 - val_acc: 0.2423
Epoch 15/250
- 26s - loss: 3.7583 - acc: 0.2297 - val_loss: 4.1827 - val_acc: 0.1991
Epoch 16/250
- 26s - loss: 3.7360 - acc: 0.2463 - val_loss: 3.6596 - val_acc: 0.2682
Epoch 17/250
- 26s - loss: 3.7420 - acc: 0.2628 - val_loss: 3.8115 - val_acc: 0.2696
Epoch 18/250
- 26s - loss: 3.7506 - acc: 0.2726 - val_loss: 3.8232 - val_acc: 0.2718
Epoch 19/250
- 26s - loss: 3.7495 - acc: 0.2869 - val_loss: 3.5881 - val_acc: 0.3313
Epoch 20/250
- 26s - loss: 3.7644 - acc: 0.2955 - val_loss: 4.1675 - val_acc: 0.2638
Epoch 21/250
- 26s - loss: 3.4662 - acc: 0.3480 - val_loss: 3.3955 - val_acc: 0.3623
Epoch 22/250
- 26s - loss: 3.3145 - acc: 0.3679 - val_loss: 3.0883 - val_acc: 0.4198
Epoch 23/250
- 26s - loss: 3.2904 - acc: 0.3740 - val_loss: 3.3546 - val_acc: 0.3677
Epoch 24/250
- 26s - loss: 3.2593 - acc: 0.3794 - val_loss: 3.2644 - val_acc: 0.3910
Epoch 25/250
- 26s - loss: 3.2598 - acc: 0.3870 - val_loss: 3.2654 - val_acc: 0.3933
Epoch 26/250
- 26s - loss: 3.2654 - acc: 0.3947 - val_loss: 3.2257 - val_acc: 0.4112
Epoch 27/250
- 26s - loss: 3.2796 - acc: 0.3992 - val_loss: 3.3614 - val_acc: 0.3934
Epoch 28/250
- 26s - loss: 3.2714 - acc: 0.4037 - val_loss: 3.3750 - val_acc: 0.3975
Epoch 29/250
```

# Machine Learning Project Report

- 26s - loss: 3.2946 - acc: 0.4073 - val\_loss: 3.1965 - val\_acc: 0.4336  
Epoch 30/250  
- 26s - loss: 3.2907 - acc: 0.4124 - val\_loss: 3.1605 - val\_acc: 0.4409  
Epoch 31/250  
- 26s - loss: 3.3008 - acc: 0.4182 - val\_loss: 3.2512 - val\_acc: 0.4366  
Epoch 32/250  
- 26s - loss: 3.3084 - acc: 0.4189 - val\_loss: 3.4488 - val\_acc: 0.4140  
Epoch 33/250  
- 26s - loss: 3.3174 - acc: 0.4248 - val\_loss: 3.3692 - val\_acc: 0.4319  
Epoch 34/250  
- 26s - loss: 3.3198 - acc: 0.4294 - val\_loss: 3.6615 - val\_acc: 0.4013  
Epoch 35/250  
- 26s - loss: 3.3179 - acc: 0.4349 - val\_loss: 3.4176 - val\_acc: 0.4288  
Epoch 36/250  
- 26s - loss: 3.3196 - acc: 0.4375 - val\_loss: 3.4943 - val\_acc: 0.4339  
Epoch 37/250  
- 26s - loss: 3.3264 - acc: 0.4382 - val\_loss: 3.4441 - val\_acc: 0.4196  
Epoch 38/250  
- 25s - loss: 3.3328 - acc: 0.4439 - val\_loss: 3.5274 - val\_acc: 0.4168  
Epoch 39/250  
- 26s - loss: 3.3257 - acc: 0.4471 - val\_loss: 3.3270 - val\_acc: 0.4529  
Epoch 40/250  
- 26s - loss: 3.3427 - acc: 0.4459 - val\_loss: 3.1870 - val\_acc: 0.4799  
Epoch 41/250  
- 25s - loss: 3.0762 - acc: 0.4987 - val\_loss: 2.9643 - val\_acc: 0.5275  
Epoch 42/250  
- 25s - loss: 2.9258 - acc: 0.5156 - val\_loss: 2.9333 - val\_acc: 0.5139  
Epoch 43/250  
- 26s - loss: 2.8772 - acc: 0.5155 - val\_loss: 2.7859 - val\_acc: 0.5362  
Epoch 44/250  
- 26s - loss: 2.8330 - acc: 0.5196 - val\_loss: 2.8306 - val\_acc: 0.5273  
Epoch 45/250  
- 26s - loss: 2.8133 - acc: 0.5236 - val\_loss: 2.8843 - val\_acc: 0.5184  
Epoch 46/250  
- 26s - loss: 2.8164 - acc: 0.5222 - val\_loss: 2.8089 - val\_acc: 0.5270  
Epoch 47/250  
- 26s - loss: 2.7995 - acc: 0.5238 - val\_loss: 2.8081 - val\_acc: 0.5272  
Epoch 48/250  
- 26s - loss: 2.8107 - acc: 0.5248 - val\_loss: 2.9375 - val\_acc: 0.5152  
Epoch 49/250  
- 26s - loss: 2.8005 - acc: 0.5300 - val\_loss: 2.8647 - val\_acc: 0.5354  
Epoch 50/250  
- 26s - loss: 2.8189 - acc: 0.5260 - val\_loss: 2.9271 - val\_acc: 0.5229  
Epoch 51/250  
- 26s - loss: 2.7965 - acc: 0.5364 - val\_loss: 2.7683 - val\_acc: 0.5571  
Epoch 52/250  
- 26s - loss: 2.8131 - acc: 0.5347 - val\_loss: 2.8927 - val\_acc: 0.5246  
Epoch 53/250  
- 26s - loss: 2.8005 - acc: 0.5389 - val\_loss: 2.9677 - val\_acc: 0.5149  
Epoch 54/250  
- 26s - loss: 2.8205 - acc: 0.5361 - val\_loss: 2.7071 - val\_acc: 0.5595  
Epoch 55/250  
- 26s - loss: 2.8157 - acc: 0.5389 - val\_loss: 2.8835 - val\_acc: 0.5416  
Epoch 56/250  
- 26s - loss: 2.8162 - acc: 0.5390 - val\_loss: 2.9346 - val\_acc: 0.5237



# Machine Learning Project Report

```
Epoch 57/250
- 26s - loss: 2.8272 - acc: 0.5425 - val_loss: 2.8332 - val_acc: 0.5476
Epoch 58/250
- 26s - loss: 2.8194 - acc: 0.5466 - val_loss: 2.9615 - val_acc: 0.5290
Epoch 59/250
- 26s - loss: 2.8302 - acc: 0.5458 - val_loss: 2.8825 - val_acc: 0.5415
Epoch 60/250
- 26s - loss: 2.8270 - acc: 0.5505 - val_loss: 2.7677 - val_acc: 0.5652
Epoch 61/250
- 26s - loss: 2.6185 - acc: 0.5891 - val_loss: 2.6056 - val_acc: 0.5962
Epoch 62/250
- 26s - loss: 2.4968 - acc: 0.6113 - val_loss: 2.5813 - val_acc: 0.5952
Epoch 63/250
- 26s - loss: 2.4497 - acc: 0.6149 - val_loss: 2.7964 - val_acc: 0.5559
Epoch 64/250
- 26s - loss: 2.4200 - acc: 0.6161 - val_loss: 2.5451 - val_acc: 0.5910
Epoch 65/250
- 25s - loss: 2.3868 - acc: 0.6176 - val_loss: 2.5715 - val_acc: 0.5837
Epoch 66/250
- 26s - loss: 2.3881 - acc: 0.6149 - val_loss: 2.4482 - val_acc: 0.6113
Epoch 67/250
- 26s - loss: 2.3585 - acc: 0.6209 - val_loss: 2.4967 - val_acc: 0.5942
Epoch 68/250
- 26s - loss: 2.3543 - acc: 0.6212 - val_loss: 2.5385 - val_acc: 0.5908
Epoch 69/250
- 26s - loss: 2.3480 - acc: 0.6215 - val_loss: 2.4505 - val_acc: 0.6024
Epoch 70/250
- 26s - loss: 2.3402 - acc: 0.6206 - val_loss: 2.4691 - val_acc: 0.6005
Epoch 71/250
- 26s - loss: 2.3278 - acc: 0.6227 - val_loss: 2.4997 - val_acc: 0.5976
Epoch 72/250
- 26s - loss: 2.3301 - acc: 0.6240 - val_loss: 2.4299 - val_acc: 0.6007
Epoch 73/250
- 26s - loss: 2.3259 - acc: 0.6239 - val_loss: 2.4850 - val_acc: 0.6004
Epoch 74/250
- 25s - loss: 2.3172 - acc: 0.6274 - val_loss: 2.4108 - val_acc: 0.6162
Epoch 75/250
- 25s - loss: 2.3247 - acc: 0.6278 - val_loss: 2.5810 - val_acc: 0.5856
Epoch 76/250
- 25s - loss: 2.3332 - acc: 0.6259 - val_loss: 2.4518 - val_acc: 0.6105
Epoch 77/250
- 25s - loss: 2.3224 - acc: 0.6296 - val_loss: 2.5016 - val_acc: 0.5983
Epoch 78/250
- 25s - loss: 2.3216 - acc: 0.6297 - val_loss: 2.5045 - val_acc: 0.6040
Epoch 79/250
- 26s - loss: 2.3214 - acc: 0.6320 - val_loss: 2.4495 - val_acc: 0.6124
Epoch 80/250
- 26s - loss: 2.3335 - acc: 0.6307 - val_loss: 2.5088 - val_acc: 0.6119
Epoch 81/250
- 26s - loss: 2.1692 - acc: 0.6688 - val_loss: 2.3187 - val_acc: 0.6439
Epoch 82/250
- 26s - loss: 2.0821 - acc: 0.6862 - val_loss: 2.3215 - val_acc: 0.6422
Epoch 83/250
- 26s - loss: 2.0432 - acc: 0.6886 - val_loss: 2.2495 - val_acc: 0.6539
Epoch 84/250
```

# Machine Learning Project Report

```
- 26s - loss: 2.0206 - acc: 0.6913 - val_loss: 2.3052 - val_acc: 0.6432
Epoch 85/250
- 26s - loss: 1.9951 - acc: 0.6948 - val_loss: 2.2957 - val_acc: 0.6416
Epoch 86/250
- 26s - loss: 1.9730 - acc: 0.6978 - val_loss: 2.2026 - val_acc: 0.6552
Epoch 87/250
- 25s - loss: 1.9527 - acc: 0.6985 - val_loss: 2.2139 - val_acc: 0.6511
Epoch 88/250
- 26s - loss: 1.9510 - acc: 0.6997 - val_loss: 2.2681 - val_acc: 0.6418
Epoch 89/250
- 26s - loss: 1.9494 - acc: 0.6981 - val_loss: 2.3120 - val_acc: 0.6386
Epoch 90/250
- 26s - loss: 1.9250 - acc: 0.7034 - val_loss: 2.2982 - val_acc: 0.6321
Epoch 91/250
- 26s - loss: 1.9273 - acc: 0.6991 - val_loss: 2.2410 - val_acc: 0.6443
Epoch 92/250
- 26s - loss: 1.9171 - acc: 0.7028 - val_loss: 2.3104 - val_acc: 0.6331
Epoch 93/250
- 26s - loss: 1.9102 - acc: 0.6997 - val_loss: 2.2510 - val_acc: 0.6430
Epoch 94/250
- 26s - loss: 1.9022 - acc: 0.7026 - val_loss: 2.2786 - val_acc: 0.6373
Epoch 95/250
- 26s - loss: 1.9005 - acc: 0.7047 - val_loss: 2.3490 - val_acc: 0.6245
Epoch 96/250
- 26s - loss: 1.9048 - acc: 0.7024 - val_loss: 2.2359 - val_acc: 0.6434
Epoch 97/250
- 26s - loss: 1.8975 - acc: 0.7045 - val_loss: 2.2191 - val_acc: 0.6447
Epoch 98/250
- 26s - loss: 1.8920 - acc: 0.7056 - val_loss: 2.3646 - val_acc: 0.6225
Epoch 99/250
- 26s - loss: 1.8960 - acc: 0.7058 - val_loss: 2.2042 - val_acc: 0.6525
Epoch 100/250
- 26s - loss: 1.8878 - acc: 0.7085 - val_loss: 2.2993 - val_acc: 0.6300
Epoch 101/250
- 26s - loss: 1.7692 - acc: 0.7371 - val_loss: 2.1028 - val_acc: 0.6725
Epoch 102/250
- 26s - loss: 1.7086 - acc: 0.7499 - val_loss: 2.1832 - val_acc: 0.6588
Epoch 103/250
- 26s - loss: 1.6819 - acc: 0.7527 - val_loss: 2.1714 - val_acc: 0.6562
Epoch 104/250
- 26s - loss: 1.6548 - acc: 0.7592 - val_loss: 2.0992 - val_acc: 0.6725
Epoch 105/250
- 26s - loss: 1.6415 - acc: 0.7588 - val_loss: 2.1263 - val_acc: 0.6678
Epoch 106/250
- 26s - loss: 1.6227 - acc: 0.7637 - val_loss: 2.2212 - val_acc: 0.6532
Epoch 107/250
- 26s - loss: 1.6175 - acc: 0.7608 - val_loss: 2.1423 - val_acc: 0.6622
Epoch 108/250
- 26s - loss: 1.5959 - acc: 0.7678 - val_loss: 2.0881 - val_acc: 0.6730
Epoch 109/250
- 26s - loss: 1.6026 - acc: 0.7640 - val_loss: 2.0944 - val_acc: 0.6713
Epoch 110/250
- 26s - loss: 1.5787 - acc: 0.7678 - val_loss: 2.0417 - val_acc: 0.6832
Epoch 111/250
- 26s - loss: 1.5838 - acc: 0.7665 - val_loss: 2.1493 - val_acc: 0.6603
```

# Machine Learning Project Report

```
Epoch 112/250
- 26s - loss: 1.5726 - acc: 0.7692 - val_loss: 2.1042 - val_acc: 0.6695
Epoch 113/250
- 26s - loss: 1.5567 - acc: 0.7716 - val_loss: 2.1922 - val_acc: 0.6540
Epoch 114/250
- 26s - loss: 1.5532 - acc: 0.7704 - val_loss: 2.1597 - val_acc: 0.6565
Epoch 115/250
- 26s - loss: 1.5354 - acc: 0.7753 - val_loss: 2.1220 - val_acc: 0.6668
Epoch 116/250
- 26s - loss: 1.5399 - acc: 0.7708 - val_loss: 2.1141 - val_acc: 0.6687
Epoch 117/250
- 26s - loss: 1.5247 - acc: 0.7740 - val_loss: 2.1849 - val_acc: 0.6578
Epoch 118/250
- 26s - loss: 1.5224 - acc: 0.7759 - val_loss: 2.1530 - val_acc: 0.6622
Epoch 119/250
- 26s - loss: 1.5246 - acc: 0.7741 - val_loss: 2.0905 - val_acc: 0.6726
Epoch 120/250
- 26s - loss: 1.5159 - acc: 0.7736 - val_loss: 2.1328 - val_acc: 0.6599
Epoch 121/250
- 26s - loss: 1.4384 - acc: 0.7954 - val_loss: 2.0659 - val_acc: 0.6777
Epoch 122/250
- 26s - loss: 1.3982 - acc: 0.8056 - val_loss: 2.1013 - val_acc: 0.6732
Epoch 123/250
- 26s - loss: 1.3799 - acc: 0.8071 - val_loss: 2.0523 - val_acc: 0.6806
Epoch 124/250
- 26s - loss: 1.3617 - acc: 0.8120 - val_loss: 2.0836 - val_acc: 0.6779
Epoch 125/250
- 26s - loss: 1.3541 - acc: 0.8147 - val_loss: 2.1260 - val_acc: 0.6739
Epoch 126/250
- 26s - loss: 1.3319 - acc: 0.8165 - val_loss: 2.1047 - val_acc: 0.6755
Epoch 127/250
- 26s - loss: 1.3152 - acc: 0.8201 - val_loss: 2.0933 - val_acc: 0.6798
Epoch 128/250
- 26s - loss: 1.3189 - acc: 0.8200 - val_loss: 2.0582 - val_acc: 0.6851
Epoch 129/250
- 26s - loss: 1.3052 - acc: 0.8225 - val_loss: 2.0383 - val_acc: 0.6863
Epoch 130/250
- 26s - loss: 1.3029 - acc: 0.8208 - val_loss: 2.1430 - val_acc: 0.6745
Epoch 131/250
- 26s - loss: 1.2852 - acc: 0.8258 - val_loss: 2.2152 - val_acc: 0.6627
Epoch 132/250
- 26s - loss: 1.2797 - acc: 0.8259 - val_loss: 2.0992 - val_acc: 0.6740
Epoch 133/250
- 26s - loss: 1.2750 - acc: 0.8253 - val_loss: 2.0821 - val_acc: 0.6768
Epoch 134/250
- 26s - loss: 1.2671 - acc: 0.8261 - val_loss: 2.0845 - val_acc: 0.6749
Epoch 135/250
- 26s - loss: 1.2611 - acc: 0.8275 - val_loss: 2.0680 - val_acc: 0.6794
Epoch 136/250
- 26s - loss: 1.2502 - acc: 0.8298 - val_loss: 2.1274 - val_acc: 0.6677
Epoch 137/250
- 26s - loss: 1.2645 - acc: 0.8263 - val_loss: 2.0627 - val_acc: 0.6811
Epoch 138/250
- 26s - loss: 1.2414 - acc: 0.8296 - val_loss: 2.0861 - val_acc: 0.6740
Epoch 139/250
```

# Machine Learning Project Report

```
- 26s - loss: 1.2390 - acc: 0.8315 - val_loss: 2.0939 - val_acc: 0.6775
Epoch 140/250
- 26s - loss: 1.2405 - acc: 0.8298 - val_loss: 2.1165 - val_acc: 0.6765
Epoch 141/250
- 26s - loss: 1.1938 - acc: 0.8433 - val_loss: 2.0326 - val_acc: 0.6886
Epoch 142/250
- 26s - loss: 1.1699 - acc: 0.8490 - val_loss: 2.0844 - val_acc: 0.6800
Epoch 143/250
- 26s - loss: 1.1508 - acc: 0.8523 - val_loss: 2.1215 - val_acc: 0.6756
Epoch 144/250
- 26s - loss: 1.1405 - acc: 0.8547 - val_loss: 2.0349 - val_acc: 0.6882
Epoch 145/250
- 26s - loss: 1.1256 - acc: 0.8593 - val_loss: 2.0729 - val_acc: 0.6848
Epoch 146/250
- 26s - loss: 1.1101 - acc: 0.8612 - val_loss: 2.0838 - val_acc: 0.6804
Epoch 147/250
- 26s - loss: 1.1211 - acc: 0.8579 - val_loss: 2.0592 - val_acc: 0.6858
Epoch 148/250
- 26s - loss: 1.1110 - acc: 0.8599 - val_loss: 2.1277 - val_acc: 0.6726
Epoch 149/250
- 26s - loss: 1.1057 - acc: 0.8607 - val_loss: 2.0784 - val_acc: 0.6813
Epoch 150/250
- 26s - loss: 1.0929 - acc: 0.8635 - val_loss: 2.1237 - val_acc: 0.6786
Epoch 151/250
- 26s - loss: 1.0930 - acc: 0.8644 - val_loss: 2.0612 - val_acc: 0.6876
Epoch 152/250
- 26s - loss: 1.0757 - acc: 0.8679 - val_loss: 2.0844 - val_acc: 0.6840
Epoch 153/250
- 26s - loss: 1.0878 - acc: 0.8623 - val_loss: 2.0860 - val_acc: 0.6819
Epoch 154/250
- 26s - loss: 1.0880 - acc: 0.8621 - val_loss: 2.1512 - val_acc: 0.6772
Epoch 155/250
- 26s - loss: 1.0690 - acc: 0.8681 - val_loss: 2.1552 - val_acc: 0.6720
Epoch 156/250
- 26s - loss: 1.0627 - acc: 0.8682 - val_loss: 2.1125 - val_acc: 0.6806
Epoch 157/250
- 26s - loss: 1.0570 - acc: 0.8701 - val_loss: 2.0952 - val_acc: 0.6854
Epoch 158/250
- 26s - loss: 1.0545 - acc: 0.8712 - val_loss: 2.0950 - val_acc: 0.6832
Epoch 159/250
- 26s - loss: 1.0550 - acc: 0.8699 - val_loss: 2.0751 - val_acc: 0.6843
Epoch 160/250
- 26s - loss: 1.0451 - acc: 0.8711 - val_loss: 2.0752 - val_acc: 0.6853
Epoch 161/250
- 26s - loss: 1.0288 - acc: 0.8751 - val_loss: 2.0755 - val_acc: 0.6858
Epoch 162/250
- 26s - loss: 1.0048 - acc: 0.8813 - val_loss: 2.0842 - val_acc: 0.6882
Epoch 163/250
- 26s - loss: 1.0056 - acc: 0.8820 - val_loss: 2.0820 - val_acc: 0.6882
Epoch 164/250
- 26s - loss: 0.9977 - acc: 0.8840 - val_loss: 2.1042 - val_acc: 0.6854
Epoch 165/250
- 26s - loss: 0.9901 - acc: 0.8861 - val_loss: 2.0950 - val_acc: 0.6870
Epoch 166/250
- 26s - loss: 0.9775 - acc: 0.8899 - val_loss: 2.0625 - val_acc: 0.6909
```

# Machine Learning Project Report

```
Epoch 167/250
- 26s - loss: 0.9818 - acc: 0.8869 - val_loss: 2.1022 - val_acc: 0.6862
Epoch 168/250
- 26s - loss: 0.9828 - acc: 0.8856 - val_loss: 2.1117 - val_acc: 0.6839
Epoch 169/250
- 26s - loss: 0.9703 - acc: 0.8901 - val_loss: 2.1006 - val_acc: 0.6867
Epoch 170/250
- 26s - loss: 0.9687 - acc: 0.8897 - val_loss: 2.0830 - val_acc: 0.6905
Epoch 171/250
- 26s - loss: 0.9674 - acc: 0.8886 - val_loss: 2.0941 - val_acc: 0.6865
Epoch 172/250
- 26s - loss: 0.9627 - acc: 0.8919 - val_loss: 2.1281 - val_acc: 0.6813
Epoch 173/250
- 26s - loss: 0.9497 - acc: 0.8939 - val_loss: 2.0903 - val_acc: 0.6880
Epoch 174/250
- 26s - loss: 0.9525 - acc: 0.8949 - val_loss: 2.1049 - val_acc: 0.6865
Epoch 175/250
- 26s - loss: 0.9587 - acc: 0.8914 - val_loss: 2.0926 - val_acc: 0.6858
Epoch 176/250
- 26s - loss: 0.9492 - acc: 0.8922 - val_loss: 2.0985 - val_acc: 0.6852
Epoch 177/250
- 26s - loss: 0.9490 - acc: 0.8925 - val_loss: 2.1051 - val_acc: 0.6851
Epoch 178/250
- 26s - loss: 0.9344 - acc: 0.8960 - val_loss: 2.1334 - val_acc: 0.6851
Epoch 179/250
- 26s - loss: 0.9501 - acc: 0.8906 - val_loss: 2.1180 - val_acc: 0.6871
Epoch 180/250
- 26s - loss: 0.9386 - acc: 0.8946 - val_loss: 2.1169 - val_acc: 0.6840
Epoch 181/250
- 26s - loss: 0.9347 - acc: 0.8951 - val_loss: 2.1136 - val_acc: 0.6859
Epoch 182/250
- 26s - loss: 0.9201 - acc: 0.8999 - val_loss: 2.0911 - val_acc: 0.6879
Epoch 183/250
- 26s - loss: 0.9136 - acc: 0.9003 - val_loss: 2.1020 - val_acc: 0.6866
Epoch 184/250
- 26s - loss: 0.9188 - acc: 0.8987 - val_loss: 2.0981 - val_acc: 0.6880
Epoch 185/250
- 26s - loss: 0.9166 - acc: 0.9010 - val_loss: 2.1189 - val_acc: 0.6846
Epoch 186/250
- 26s - loss: 0.9161 - acc: 0.9018 - val_loss: 2.1113 - val_acc: 0.6857
Epoch 187/250
- 26s - loss: 0.9076 - acc: 0.9019 - val_loss: 2.1191 - val_acc: 0.6844
Epoch 188/250
- 26s - loss: 0.9050 - acc: 0.9020 - val_loss: 2.1009 - val_acc: 0.6881
Epoch 189/250
- 26s - loss: 0.9028 - acc: 0.9044 - val_loss: 2.1037 - val_acc: 0.6880
Epoch 190/250
- 26s - loss: 0.8945 - acc: 0.9044 - val_loss: 2.1015 - val_acc: 0.6888
Epoch 191/250
- 26s - loss: 0.8985 - acc: 0.9046 - val_loss: 2.1237 - val_acc: 0.6850
Epoch 192/250
- 26s - loss: 0.8992 - acc: 0.9028 - val_loss: 2.1044 - val_acc: 0.6874
Epoch 193/250
- 26s - loss: 0.8917 - acc: 0.9053 - val_loss: 2.1046 - val_acc: 0.6875
Epoch 194/250
```

# Machine Learning Project Report

```
- 26s - loss: 0.8976 - acc: 0.9047 - val_loss: 2.0956 - val_acc: 0.6921
Epoch 195/250
- 26s - loss: 0.8882 - acc: 0.9058 - val_loss: 2.1317 - val_acc: 0.6874
Epoch 196/250
- 26s - loss: 0.8934 - acc: 0.9050 - val_loss: 2.1004 - val_acc: 0.6914
Epoch 197/250
- 26s - loss: 0.8866 - acc: 0.9066 - val_loss: 2.1143 - val_acc: 0.6890
Epoch 198/250
- 26s - loss: 0.8908 - acc: 0.9039 - val_loss: 2.1179 - val_acc: 0.6896
Epoch 199/250
- 26s - loss: 0.8849 - acc: 0.9065 - val_loss: 2.1083 - val_acc: 0.6913
Epoch 200/250
- 26s - loss: 0.8768 - acc: 0.9092 - val_loss: 2.1075 - val_acc: 0.6875
Epoch 201/250
- 26s - loss: 0.8697 - acc: 0.9108 - val_loss: 2.0998 - val_acc: 0.6898
Epoch 202/250
- 26s - loss: 0.8736 - acc: 0.9091 - val_loss: 2.1054 - val_acc: 0.6875
Epoch 203/250
- 26s - loss: 0.8736 - acc: 0.9092 - val_loss: 2.1071 - val_acc: 0.6880
Epoch 204/250
- 26s - loss: 0.8712 - acc: 0.9101 - val_loss: 2.1118 - val_acc: 0.6887
Epoch 205/250
- 26s - loss: 0.8745 - acc: 0.9093 - val_loss: 2.1012 - val_acc: 0.6896
Epoch 206/250
- 26s - loss: 0.8622 - acc: 0.9120 - val_loss: 2.1030 - val_acc: 0.6907
Epoch 207/250
- 26s - loss: 0.8660 - acc: 0.9112 - val_loss: 2.1078 - val_acc: 0.6876
Epoch 208/250
- 26s - loss: 0.8660 - acc: 0.9114 - val_loss: 2.1000 - val_acc: 0.6891
Epoch 209/250
- 26s - loss: 0.8654 - acc: 0.9102 - val_loss: 2.1071 - val_acc: 0.6874
Epoch 210/250
- 26s - loss: 0.8632 - acc: 0.9112 - val_loss: 2.0927 - val_acc: 0.6915
Epoch 211/250
- 26s - loss: 0.8607 - acc: 0.9131 - val_loss: 2.1102 - val_acc: 0.6885
Epoch 212/250
- 26s - loss: 0.8579 - acc: 0.9130 - val_loss: 2.1067 - val_acc: 0.6904
Epoch 213/250
- 26s - loss: 0.8586 - acc: 0.9139 - val_loss: 2.0984 - val_acc: 0.6913
Epoch 214/250
- 26s - loss: 0.8651 - acc: 0.9112 - val_loss: 2.1113 - val_acc: 0.6898
Epoch 215/250
- 26s - loss: 0.8612 - acc: 0.9121 - val_loss: 2.1016 - val_acc: 0.6888
Epoch 216/250
- 25s - loss: 0.8590 - acc: 0.9141 - val_loss: 2.0915 - val_acc: 0.6923
Epoch 217/250
- 26s - loss: 0.8576 - acc: 0.9111 - val_loss: 2.1094 - val_acc: 0.6877
Epoch 218/250
- 26s - loss: 0.8511 - acc: 0.9144 - val_loss: 2.1187 - val_acc: 0.6878
Epoch 219/250
- 25s - loss: 0.8594 - acc: 0.9108 - val_loss: 2.1124 - val_acc: 0.6880
Epoch 220/250
- 25s - loss: 0.8548 - acc: 0.9127 - val_loss: 2.1165 - val_acc: 0.6888
Epoch 221/250
- 25s - loss: 0.8490 - acc: 0.9142 - val_loss: 2.1217 - val_acc: 0.6890
```

# Machine Learning Project Report

```
Epoch 222/250
- 26s - loss: 0.8496 - acc: 0.9147 - val_loss: 2.1120 - val_acc: 0.6890
Epoch 223/250
- 26s - loss: 0.8483 - acc: 0.9139 - val_loss: 2.1117 - val_acc: 0.6912
Epoch 224/250
- 26s - loss: 0.8452 - acc: 0.9155 - val_loss: 2.1121 - val_acc: 0.6897
Epoch 225/250
- 26s - loss: 0.8501 - acc: 0.9146 - val_loss: 2.1169 - val_acc: 0.6887
Epoch 226/250
- 26s - loss: 0.8510 - acc: 0.9146 - val_loss: 2.1184 - val_acc: 0.6888
Epoch 227/250
- 26s - loss: 0.8443 - acc: 0.9165 - val_loss: 2.1152 - val_acc: 0.6901
Epoch 228/250
- 26s - loss: 0.8450 - acc: 0.9158 - val_loss: 2.1082 - val_acc: 0.6903
Epoch 229/250
- 27s - loss: 0.8445 - acc: 0.9157 - val_loss: 2.1134 - val_acc: 0.6895
Epoch 230/250
- 26s - loss: 0.8441 - acc: 0.9152 - val_loss: 2.1112 - val_acc: 0.6889
Epoch 231/250
- 26s - loss: 0.8444 - acc: 0.9144 - val_loss: 2.1102 - val_acc: 0.6904
Epoch 232/250
- 26s - loss: 0.8491 - acc: 0.9137 - val_loss: 2.1135 - val_acc: 0.6901
Epoch 233/250
- 26s - loss: 0.8458 - acc: 0.9148 - val_loss: 2.1136 - val_acc: 0.6903
Epoch 234/250
- 26s - loss: 0.8380 - acc: 0.9171 - val_loss: 2.1115 - val_acc: 0.6903
Epoch 235/250
- 26s - loss: 0.8396 - acc: 0.9171 - val_loss: 2.1068 - val_acc: 0.6901
Epoch 236/250
- 26s - loss: 0.8457 - acc: 0.9151 - val_loss: 2.1118 - val_acc: 0.6883
Epoch 237/250
- 26s - loss: 0.8442 - acc: 0.9156 - val_loss: 2.1038 - val_acc: 0.6919
Epoch 238/250
- 26s - loss: 0.8452 - acc: 0.9144 - val_loss: 2.1092 - val_acc: 0.6903
Epoch 239/250
- 26s - loss: 0.8399 - acc: 0.9155 - val_loss: 2.1039 - val_acc: 0.6909
Epoch 240/250
- 26s - loss: 0.8401 - acc: 0.9170 - val_loss: 2.1130 - val_acc: 0.6911
Epoch 241/250
- 26s - loss: 0.8355 - acc: 0.9181 - val_loss: 2.1085 - val_acc: 0.6896
Epoch 242/250
- 26s - loss: 0.8355 - acc: 0.9179 - val_loss: 2.1142 - val_acc: 0.6891
Epoch 243/250
- 26s - loss: 0.8374 - acc: 0.9174 - val_loss: 2.1143 - val_acc: 0.6954
Epoch 244/250
- 26s - loss: 0.8294 - acc: 0.9190 - val_loss: 2.1124 - val_acc: 0.6991
Epoch 245/250
- 26s - loss: 0.8310 - acc: 0.9192 - val_loss: 2.0991 - val_acc: 0.7034
Epoch 246/250
- 26s - loss: 0.8367 - acc: 0.9163 - val_loss: 2.1037 - val_acc: 0.7056
Epoch 247/250
- 26s - loss: 0.8393 - acc: 0.9170 - val_loss: 2.1076 - val_acc: 0.7089
Epoch 248/250
- 26s - loss: 0.8365 - acc: 0.9182 - val_loss: 2.1078 - val_acc: 0.7101
Epoch 249/250
```

# Machine Learning Project Report

```
- 26s - loss: 0.8260 - acc: 0.9201 - val_loss: 2.1024 - val_acc: 0.7121
Epoch 250/250
- 26s - loss: 0.8319 - acc: 0.9193 - val_loss: 2.1072 - val_acc: 0.7143
Training Time : 6461.053843021393
Accuracy: 71.43 %
```

Run 3:

```
121.93584
68.38902
Epoch 1/250
- 30s - loss: 18.7502 - acc: 0.0239 - val_loss: 14.6522 - val_acc: 0.0163
Epoch 2/250
- 26s - loss: 11.4647 - acc: 0.0439 - val_loss: 9.4003 - val_acc: 0.0218
Epoch 3/250
- 26s - loss: 8.5017 - acc: 0.0394 - val_loss: 7.8104 - val_acc: 0.0445
Epoch 4/250
- 26s - loss: 7.0439 - acc: 0.0468 - val_loss: 6.3975 - val_acc: 0.0336
Epoch 5/250
- 26s - loss: 5.5937 - acc: 0.0657 - val_loss: 5.6558 - val_acc: 0.0202
Epoch 6/250
- 26s - loss: 4.7609 - acc: 0.0842 - val_loss: 4.9156 - val_acc: 0.0449
Epoch 7/250
- 26s - loss: 4.3819 - acc: 0.0957 - val_loss: 4.5202 - val_acc: 0.0738
Epoch 8/250
- 26s - loss: 4.1964 - acc: 0.1127 - val_loss: 4.2389 - val_acc: 0.1076
Epoch 9/250
- 26s - loss: 4.0565 - acc: 0.1296 - val_loss: 3.9747 - val_acc: 0.1556
Epoch 10/250
- 26s - loss: 3.9777 - acc: 0.1447 - val_loss: 3.8448 - val_acc: 0.1701
Epoch 11/250
- 26s - loss: 3.9027 - acc: 0.1598 - val_loss: 3.9674 - val_acc: 0.1752
Epoch 12/250
- 26s - loss: 3.8708 - acc: 0.1702 - val_loss: 3.7472 - val_acc: 0.1995
Epoch 13/250
- 26s - loss: 3.8097 - acc: 0.1989 - val_loss: 3.6767 - val_acc: 0.2364
Epoch 14/250
- 26s - loss: 3.7564 - acc: 0.2171 - val_loss: 3.6926 - val_acc: 0.2432
Epoch 15/250
- 26s - loss: 3.7363 - acc: 0.2364 - val_loss: 3.6505 - val_acc: 0.2561
Epoch 16/250
- 26s - loss: 3.7241 - acc: 0.2543 - val_loss: 3.6460 - val_acc: 0.2812
Epoch 17/250
- 26s - loss: 3.7397 - acc: 0.2694 - val_loss: 3.8514 - val_acc: 0.2640
Epoch 18/250
- 26s - loss: 3.7257 - acc: 0.2822 - val_loss: 3.6232 - val_acc: 0.3153
Epoch 19/250
- 26s - loss: 3.7412 - acc: 0.2949 - val_loss: 3.5996 - val_acc: 0.3222
Epoch 20/250
```



# Machine Learning Project Report

```
- 26s - loss: 3.7478 - acc: 0.3044 - val_loss: 3.8868 - val_acc: 0.2966
Epoch 21/250
- 26s - loss: 3.4332 - acc: 0.3616 - val_loss: 3.3439 - val_acc: 0.3855
Epoch 22/250
- 26s - loss: 3.2927 - acc: 0.3756 - val_loss: 3.3064 - val_acc: 0.3900
Epoch 23/250
- 26s - loss: 3.2552 - acc: 0.3831 - val_loss: 3.2420 - val_acc: 0.3918
Epoch 24/250
- 26s - loss: 3.2453 - acc: 0.3893 - val_loss: 3.3796 - val_acc: 0.3756
Epoch 25/250
- 26s - loss: 3.2437 - acc: 0.3984 - val_loss: 3.4726 - val_acc: 0.3756
Epoch 26/250
- 26s - loss: 3.2474 - acc: 0.4043 - val_loss: 3.2384 - val_acc: 0.4103
Epoch 27/250
- 26s - loss: 3.2552 - acc: 0.4106 - val_loss: 3.2692 - val_acc: 0.4261
Epoch 28/250
- 26s - loss: 3.2485 - acc: 0.4178 - val_loss: 3.1857 - val_acc: 0.4476
Epoch 29/250
- 26s - loss: 3.2623 - acc: 0.4196 - val_loss: 3.2810 - val_acc: 0.4322
Epoch 30/250
- 26s - loss: 3.2725 - acc: 0.4271 - val_loss: 3.4296 - val_acc: 0.4036
Epoch 31/250
- 26s - loss: 3.2937 - acc: 0.4279 - val_loss: 3.2417 - val_acc: 0.4453
Epoch 32/250
- 26s - loss: 3.2858 - acc: 0.4350 - val_loss: 3.2418 - val_acc: 0.4514
Epoch 33/250
- 26s - loss: 3.2899 - acc: 0.4369 - val_loss: 3.2702 - val_acc: 0.4470
Epoch 34/250
- 26s - loss: 3.2852 - acc: 0.4416 - val_loss: 3.4443 - val_acc: 0.4305
Epoch 35/250
- 26s - loss: 3.3034 - acc: 0.4431 - val_loss: 3.1901 - val_acc: 0.4729
Epoch 36/250
- 26s - loss: 3.3096 - acc: 0.4445 - val_loss: 3.2266 - val_acc: 0.4733
Epoch 37/250
- 26s - loss: 3.3132 - acc: 0.4490 - val_loss: 3.3243 - val_acc: 0.4543
Epoch 38/250
- 26s - loss: 3.2977 - acc: 0.4539 - val_loss: 3.2419 - val_acc: 0.4777
Epoch 39/250
- 25s - loss: 3.3201 - acc: 0.4546 - val_loss: 3.2470 - val_acc: 0.4721
Epoch 40/250
- 26s - loss: 3.3196 - acc: 0.4579 - val_loss: 3.5200 - val_acc: 0.4310
Epoch 41/250
- 25s - loss: 3.0563 - acc: 0.5077 - val_loss: 2.8816 - val_acc: 0.5386
Epoch 42/250
- 26s - loss: 2.9137 - acc: 0.5225 - val_loss: 2.9258 - val_acc: 0.5148
Epoch 43/250
- 26s - loss: 2.8375 - acc: 0.5288 - val_loss: 2.8640 - val_acc: 0.5211
Epoch 44/250
- 26s - loss: 2.8149 - acc: 0.5281 - val_loss: 3.0841 - val_acc: 0.4838
Epoch 45/250
- 26s - loss: 2.7931 - acc: 0.5290 - val_loss: 2.7688 - val_acc: 0.5437
Epoch 46/250
```

# Machine Learning Project Report

- 26s - loss: 2.7877 - acc: 0.5310 - val\_loss: 2.8550 - val\_acc: 0.5248  
Epoch 47/250  
- 25s - loss: 2.7745 - acc: 0.5318 - val\_loss: 2.7962 - val\_acc: 0.5419  
Epoch 48/250  
- 26s - loss: 2.7858 - acc: 0.5319 - val\_loss: 2.9742 - val\_acc: 0.5090  
Epoch 49/250  
- 26s - loss: 2.7763 - acc: 0.5360 - val\_loss: 3.0202 - val\_acc: 0.5059  
Epoch 50/250  
- 26s - loss: 2.7913 - acc: 0.5366 - val\_loss: 2.7573 - val\_acc: 0.5519  
Epoch 51/250  
- 26s - loss: 2.7821 - acc: 0.5400 - val\_loss: 3.0113 - val\_acc: 0.5130  
Epoch 52/250  
- 26s - loss: 2.7840 - acc: 0.5413 - val\_loss: 2.8658 - val\_acc: 0.5355  
Epoch 53/250  
- 26s - loss: 2.7788 - acc: 0.5443 - val\_loss: 2.8707 - val\_acc: 0.5323  
Epoch 54/250  
- 26s - loss: 2.7966 - acc: 0.5437 - val\_loss: 3.0194 - val\_acc: 0.5117  
Epoch 55/250  
- 26s - loss: 2.7916 - acc: 0.5467 - val\_loss: 3.0979 - val\_acc: 0.5106  
Epoch 56/250  
- 26s - loss: 2.7974 - acc: 0.5501 - val\_loss: 2.8678 - val\_acc: 0.5407  
Epoch 57/250  
- 26s - loss: 2.7963 - acc: 0.5490 - val\_loss: 2.9129 - val\_acc: 0.5433  
Epoch 58/250  
- 26s - loss: 2.8011 - acc: 0.5532 - val\_loss: 2.9803 - val\_acc: 0.5264  
Epoch 59/250  
- 26s - loss: 2.7910 - acc: 0.5575 - val\_loss: 2.9021 - val\_acc: 0.5466  
Epoch 60/250  
- 26s - loss: 2.8054 - acc: 0.5528 - val\_loss: 2.8983 - val\_acc: 0.5358  
Epoch 61/250  
- 26s - loss: 2.5998 - acc: 0.5975 - val\_loss: 2.6724 - val\_acc: 0.5901  
Epoch 62/250  
- 26s - loss: 2.4919 - acc: 0.6134 - val\_loss: 2.5270 - val\_acc: 0.6051  
Epoch 63/250  
- 26s - loss: 2.4184 - acc: 0.6236 - val\_loss: 2.4620 - val\_acc: 0.6162  
Epoch 64/250  
- 26s - loss: 2.3877 - acc: 0.6244 - val\_loss: 2.5150 - val\_acc: 0.5992  
Epoch 65/250  
- 26s - loss: 2.3548 - acc: 0.6250 - val\_loss: 2.6662 - val\_acc: 0.5793  
Epoch 66/250  
- 26s - loss: 2.3481 - acc: 0.6250 - val\_loss: 2.4814 - val\_acc: 0.6067  
Epoch 67/250  
- 26s - loss: 2.3337 - acc: 0.6257 - val\_loss: 2.6290 - val\_acc: 0.5786  
Epoch 68/250  
- 26s - loss: 2.3335 - acc: 0.6243 - val\_loss: 2.4187 - val\_acc: 0.6152  
Epoch 69/250  
- 26s - loss: 2.3195 - acc: 0.6275 - val\_loss: 2.4531 - val\_acc: 0.6141  
Epoch 70/250  
- 26s - loss: 2.3072 - acc: 0.6310 - val\_loss: 2.4255 - val\_acc: 0.6136  
Epoch 71/250  
- 26s - loss: 2.3233 - acc: 0.6237 - val\_loss: 2.5032 - val\_acc: 0.5923  
Epoch 72/250

# Machine Learning Project Report

```
- 26s - loss: 2.2982 - acc: 0.6319 - val_loss: 2.3902 - val_acc: 0.6140
Epoch 73/250
- 26s - loss: 2.3140 - acc: 0.6289 - val_loss: 2.5192 - val_acc: 0.5951
Epoch 74/250
- 26s - loss: 2.3099 - acc: 0.6317 - val_loss: 2.3603 - val_acc: 0.6250
Epoch 75/250
- 26s - loss: 2.3061 - acc: 0.6338 - val_loss: 2.4870 - val_acc: 0.6000
Epoch 76/250
- 26s - loss: 2.3095 - acc: 0.6332 - val_loss: 2.4098 - val_acc: 0.6186
Epoch 77/250
- 26s - loss: 2.3074 - acc: 0.6345 - val_loss: 2.5692 - val_acc: 0.5944
Epoch 78/250
- 26s - loss: 2.3051 - acc: 0.6362 - val_loss: 2.5442 - val_acc: 0.5958
Epoch 79/250
- 26s - loss: 2.2992 - acc: 0.6401 - val_loss: 2.4147 - val_acc: 0.6225
Epoch 80/250
- 26s - loss: 2.3042 - acc: 0.6394 - val_loss: 2.4915 - val_acc: 0.6106
Epoch 81/250
- 26s - loss: 2.1470 - acc: 0.6748 - val_loss: 2.3966 - val_acc: 0.6303
Epoch 82/250
- 26s - loss: 2.0641 - acc: 0.6911 - val_loss: 2.2644 - val_acc: 0.6523
Epoch 83/250
- 26s - loss: 2.0245 - acc: 0.6953 - val_loss: 2.3132 - val_acc: 0.6389
Epoch 84/250
- 26s - loss: 1.9981 - acc: 0.6989 - val_loss: 2.3844 - val_acc: 0.6264
Epoch 85/250
- 26s - loss: 1.9769 - acc: 0.6995 - val_loss: 2.2197 - val_acc: 0.6511
Epoch 86/250
- 26s - loss: 1.9509 - acc: 0.7021 - val_loss: 2.2696 - val_acc: 0.6456
Epoch 87/250
- 26s - loss: 1.9473 - acc: 0.7012 - val_loss: 2.2549 - val_acc: 0.6453
Epoch 88/250
- 26s - loss: 1.9255 - acc: 0.7075 - val_loss: 2.2759 - val_acc: 0.6438
Epoch 89/250
- 26s - loss: 1.9087 - acc: 0.7074 - val_loss: 2.2779 - val_acc: 0.6467
Epoch 90/250
- 26s - loss: 1.9196 - acc: 0.7038 - val_loss: 2.2902 - val_acc: 0.6354
Epoch 91/250
- 26s - loss: 1.8944 - acc: 0.7062 - val_loss: 2.2026 - val_acc: 0.6537
Epoch 92/250
- 26s - loss: 1.8900 - acc: 0.7091 - val_loss: 2.1911 - val_acc: 0.6508
Epoch 93/250
- 26s - loss: 1.8871 - acc: 0.7096 - val_loss: 2.2220 - val_acc: 0.6504
Epoch 94/250
- 26s - loss: 1.8785 - acc: 0.7095 - val_loss: 2.2527 - val_acc: 0.6420
Epoch 95/250
- 26s - loss: 1.8774 - acc: 0.7083 - val_loss: 2.2138 - val_acc: 0.6439
Epoch 96/250
- 26s - loss: 1.8871 - acc: 0.7058 - val_loss: 2.1967 - val_acc: 0.6541
Epoch 97/250
- 26s - loss: 1.8739 - acc: 0.7095 - val_loss: 2.2361 - val_acc: 0.6445
Epoch 98/250
```

# Machine Learning Project Report

```
- 26s - loss: 1.8685 - acc: 0.7108 - val_loss: 2.2807 - val_acc: 0.6407
Epoch 99/250
- 26s - loss: 1.8649 - acc: 0.7125 - val_loss: 2.2264 - val_acc: 0.6524
Epoch 100/250
- 26s - loss: 1.8704 - acc: 0.7099 - val_loss: 2.1528 - val_acc: 0.6618
Epoch 101/250
- 26s - loss: 1.7519 - acc: 0.7416 - val_loss: 2.1890 - val_acc: 0.6602
Epoch 102/250
- 26s - loss: 1.6904 - acc: 0.7543 - val_loss: 2.1657 - val_acc: 0.6637
Epoch 103/250
- 26s - loss: 1.6693 - acc: 0.7585 - val_loss: 2.1556 - val_acc: 0.6645
Epoch 104/250
- 26s - loss: 1.6383 - acc: 0.7631 - val_loss: 2.2171 - val_acc: 0.6561
Epoch 105/250
- 26s - loss: 1.6262 - acc: 0.7610 - val_loss: 2.1359 - val_acc: 0.6670
Epoch 106/250
- 26s - loss: 1.6015 - acc: 0.7671 - val_loss: 2.1451 - val_acc: 0.6663
Epoch 107/250
- 26s - loss: 1.5942 - acc: 0.7674 - val_loss: 2.1227 - val_acc: 0.6711
Epoch 108/250
- 26s - loss: 1.5773 - acc: 0.7726 - val_loss: 2.1091 - val_acc: 0.6749
Epoch 109/250
- 26s - loss: 1.5593 - acc: 0.7740 - val_loss: 2.1275 - val_acc: 0.6687
Epoch 110/250
- 26s - loss: 1.5591 - acc: 0.7712 - val_loss: 2.1299 - val_acc: 0.6711
Epoch 111/250
- 26s - loss: 1.5567 - acc: 0.7731 - val_loss: 2.2058 - val_acc: 0.6511
Epoch 112/250
- 26s - loss: 1.5480 - acc: 0.7711 - val_loss: 2.1612 - val_acc: 0.6620
Epoch 113/250
- 26s - loss: 1.5294 - acc: 0.7761 - val_loss: 2.0955 - val_acc: 0.6703
Epoch 114/250
- 26s - loss: 1.5189 - acc: 0.7781 - val_loss: 2.1476 - val_acc: 0.6616
Epoch 115/250
- 26s - loss: 1.5136 - acc: 0.7793 - val_loss: 2.1449 - val_acc: 0.6688
Epoch 116/250
- 26s - loss: 1.5249 - acc: 0.7758 - val_loss: 2.0870 - val_acc: 0.6715
Epoch 117/250
- 26s - loss: 1.5118 - acc: 0.7790 - val_loss: 2.1889 - val_acc: 0.6560
Epoch 118/250
- 26s - loss: 1.5148 - acc: 0.7752 - val_loss: 2.1818 - val_acc: 0.6535
Epoch 119/250
- 26s - loss: 1.4964 - acc: 0.7781 - val_loss: 2.1277 - val_acc: 0.6657
Epoch 120/250
- 26s - loss: 1.5010 - acc: 0.7786 - val_loss: 2.0444 - val_acc: 0.6781
Epoch 121/250
- 26s - loss: 1.4308 - acc: 0.7962 - val_loss: 2.0853 - val_acc: 0.6760
Epoch 122/250
- 26s - loss: 1.3803 - acc: 0.8072 - val_loss: 2.1075 - val_acc: 0.6738
Epoch 123/250
- 26s - loss: 1.3582 - acc: 0.8148 - val_loss: 2.0834 - val_acc: 0.6770
Epoch 124/250
```

# Machine Learning Project Report

```
- 26s - loss: 1.3390 - acc: 0.8175 - val_loss: 2.1116 - val_acc: 0.6742
Epoch 125/250
- 26s - loss: 1.3278 - acc: 0.8192 - val_loss: 2.0959 - val_acc: 0.6784
Epoch 126/250
- 26s - loss: 1.3115 - acc: 0.8228 - val_loss: 2.1019 - val_acc: 0.6803
Epoch 127/250
- 26s - loss: 1.3015 - acc: 0.8250 - val_loss: 2.0544 - val_acc: 0.6851
Epoch 128/250
- 26s - loss: 1.3053 - acc: 0.8216 - val_loss: 2.0908 - val_acc: 0.6826
Epoch 129/250
- 26s - loss: 1.2773 - acc: 0.8269 - val_loss: 2.1007 - val_acc: 0.6787
Epoch 130/250
- 26s - loss: 1.2769 - acc: 0.8246 - val_loss: 2.1037 - val_acc: 0.6786
Epoch 131/250
- 26s - loss: 1.2767 - acc: 0.8266 - val_loss: 2.1089 - val_acc: 0.6765
Epoch 132/250
- 26s - loss: 1.2672 - acc: 0.8292 - val_loss: 2.1097 - val_acc: 0.6786
Epoch 133/250
- 26s - loss: 1.2448 - acc: 0.8326 - val_loss: 2.0886 - val_acc: 0.6840
Epoch 134/250
- 26s - loss: 1.2476 - acc: 0.8315 - val_loss: 2.0529 - val_acc: 0.6880
Epoch 135/250
- 26s - loss: 1.2521 - acc: 0.8304 - val_loss: 2.1184 - val_acc: 0.6771
Epoch 136/250
- 26s - loss: 1.2311 - acc: 0.8345 - val_loss: 2.0452 - val_acc: 0.6875
Epoch 137/250
- 26s - loss: 1.2306 - acc: 0.8349 - val_loss: 2.1017 - val_acc: 0.6781
Epoch 138/250
- 26s - loss: 1.2332 - acc: 0.8308 - val_loss: 2.0805 - val_acc: 0.6830
Epoch 139/250
- 26s - loss: 1.2173 - acc: 0.8344 - val_loss: 2.0872 - val_acc: 0.6841
Epoch 140/250
- 26s - loss: 1.2137 - acc: 0.8354 - val_loss: 2.0960 - val_acc: 0.6789
Epoch 141/250
- 26s - loss: 1.1783 - acc: 0.8448 - val_loss: 2.0420 - val_acc: 0.6890
Epoch 142/250
- 26s - loss: 1.1503 - acc: 0.8503 - val_loss: 2.0141 - val_acc: 0.6966
Epoch 143/250
- 26s - loss: 1.1312 - acc: 0.8539 - val_loss: 2.0733 - val_acc: 0.6858
Epoch 144/250
- 26s - loss: 1.1238 - acc: 0.8581 - val_loss: 2.0564 - val_acc: 0.6891
Epoch 145/250
- 26s - loss: 1.1092 - acc: 0.8612 - val_loss: 2.0760 - val_acc: 0.6832
Epoch 146/250
- 26s - loss: 1.0986 - acc: 0.8650 - val_loss: 2.1183 - val_acc: 0.6788
Epoch 147/250
- 26s - loss: 1.0911 - acc: 0.8656 - val_loss: 2.0543 - val_acc: 0.6931
Epoch 148/250
- 26s - loss: 1.0903 - acc: 0.8647 - val_loss: 2.0472 - val_acc: 0.6928
Epoch 149/250
- 26s - loss: 1.0825 - acc: 0.8650 - val_loss: 2.0744 - val_acc: 0.6872
Epoch 150/250
```

# Machine Learning Project Report

```
- 26s - loss: 1.0815 - acc: 0.8663 - val_loss: 2.0581 - val_acc: 0.6917
Epoch 151/250
- 26s - loss: 1.0744 - acc: 0.8685 - val_loss: 2.0645 - val_acc: 0.6921
Epoch 152/250
- 26s - loss: 1.0713 - acc: 0.8675 - val_loss: 2.0768 - val_acc: 0.6894
Epoch 153/250
- 26s - loss: 1.0619 - acc: 0.8683 - val_loss: 2.0806 - val_acc: 0.6859
Epoch 154/250
- 26s - loss: 1.0588 - acc: 0.8698 - val_loss: 2.0942 - val_acc: 0.6887
Epoch 155/250
- 26s - loss: 1.0472 - acc: 0.8717 - val_loss: 2.0768 - val_acc: 0.6901
Epoch 156/250
- 26s - loss: 1.0546 - acc: 0.8703 - val_loss: 2.0940 - val_acc: 0.6881
Epoch 157/250
- 26s - loss: 1.0412 - acc: 0.8721 - val_loss: 2.1150 - val_acc: 0.6869
Epoch 158/250
- 26s - loss: 1.0449 - acc: 0.8711 - val_loss: 2.0860 - val_acc: 0.6898
Epoch 159/250
- 26s - loss: 1.0383 - acc: 0.8736 - val_loss: 2.0904 - val_acc: 0.6843
Epoch 160/250
- 26s - loss: 1.0298 - acc: 0.8748 - val_loss: 2.0676 - val_acc: 0.6897
Epoch 161/250
- 26s - loss: 1.0160 - acc: 0.8773 - val_loss: 2.0261 - val_acc: 0.6976
Epoch 162/250
- 26s - loss: 0.9921 - acc: 0.8843 - val_loss: 2.0318 - val_acc: 0.6976
Epoch 163/250
- 26s - loss: 0.9768 - acc: 0.8879 - val_loss: 2.0391 - val_acc: 0.6984
Epoch 164/250
- 26s - loss: 0.9841 - acc: 0.8868 - val_loss: 2.0497 - val_acc: 0.6948
Epoch 165/250
- 26s - loss: 0.9745 - acc: 0.8877 - val_loss: 2.0702 - val_acc: 0.6924
Epoch 166/250
- 26s - loss: 0.9735 - acc: 0.8876 - val_loss: 2.0657 - val_acc: 0.6953
Epoch 167/250
- 26s - loss: 0.9670 - acc: 0.8907 - val_loss: 2.0767 - val_acc: 0.6920
Epoch 168/250
- 26s - loss: 0.9650 - acc: 0.8901 - val_loss: 2.0825 - val_acc: 0.6939
Epoch 169/250
- 26s - loss: 0.9658 - acc: 0.8876 - val_loss: 2.0887 - val_acc: 0.6903
Epoch 170/250
- 26s - loss: 0.9514 - acc: 0.8925 - val_loss: 2.0826 - val_acc: 0.6911
Epoch 171/250
- 26s - loss: 0.9536 - acc: 0.8907 - val_loss: 2.0762 - val_acc: 0.6935
Epoch 172/250
- 26s - loss: 0.9394 - acc: 0.8950 - val_loss: 2.0935 - val_acc: 0.6897
Epoch 173/250
- 26s - loss: 0.9469 - acc: 0.8936 - val_loss: 2.0815 - val_acc: 0.6935
Epoch 174/250
- 26s - loss: 0.9382 - acc: 0.8952 - val_loss: 2.0725 - val_acc: 0.6965
Epoch 175/250
- 26s - loss: 0.9427 - acc: 0.8946 - val_loss: 2.0500 - val_acc: 0.6997
Epoch 176/250
```

# Machine Learning Project Report

```
- 26s - loss: 0.9364 - acc: 0.8950 - val_loss: 2.0801 - val_acc: 0.6933
Epoch 177/250
- 26s - loss: 0.9308 - acc: 0.8971 - val_loss: 2.0972 - val_acc: 0.6906
Epoch 178/250
- 26s - loss: 0.9295 - acc: 0.8948 - val_loss: 2.0688 - val_acc: 0.6950
Epoch 179/250
- 26s - loss: 0.9240 - acc: 0.8969 - val_loss: 2.0775 - val_acc: 0.6951
Epoch 180/250
- 26s - loss: 0.9213 - acc: 0.8976 - val_loss: 2.0616 - val_acc: 0.6958
Epoch 181/250
- 26s - loss: 0.9106 - acc: 0.8999 - val_loss: 2.0761 - val_acc: 0.6939
Epoch 182/250
- 26s - loss: 0.9012 - acc: 0.9044 - val_loss: 2.0734 - val_acc: 0.6929
Epoch 183/250
- 26s - loss: 0.9044 - acc: 0.9017 - val_loss: 2.0819 - val_acc: 0.6940
Epoch 184/250
- 26s - loss: 0.9006 - acc: 0.9035 - val_loss: 2.0736 - val_acc: 0.6952
Epoch 185/250
- 26s - loss: 0.8970 - acc: 0.9034 - val_loss: 2.0632 - val_acc: 0.6995
Epoch 186/250
- 26s - loss: 0.8937 - acc: 0.9043 - val_loss: 2.0884 - val_acc: 0.6939
Epoch 187/250
- 26s - loss: 0.8901 - acc: 0.9049 - val_loss: 2.0800 - val_acc: 0.6950
Epoch 188/250
- 26s - loss: 0.8857 - acc: 0.9072 - val_loss: 2.0778 - val_acc: 0.6963
Epoch 189/250
- 26s - loss: 0.8879 - acc: 0.9059 - val_loss: 2.1019 - val_acc: 0.6943
Epoch 190/250
- 26s - loss: 0.8903 - acc: 0.9038 - val_loss: 2.0928 - val_acc: 0.6947
Epoch 191/250
- 26s - loss: 0.8851 - acc: 0.9065 - val_loss: 2.0849 - val_acc: 0.6972
Epoch 192/250
- 26s - loss: 0.8824 - acc: 0.9077 - val_loss: 2.0826 - val_acc: 0.6965
Epoch 193/250
- 26s - loss: 0.8818 - acc: 0.9055 - val_loss: 2.0758 - val_acc: 0.6968
Epoch 194/250
- 26s - loss: 0.8801 - acc: 0.9059 - val_loss: 2.0852 - val_acc: 0.6928
Epoch 195/250
- 26s - loss: 0.8766 - acc: 0.9082 - val_loss: 2.0826 - val_acc: 0.6947
Epoch 196/250
- 26s - loss: 0.8744 - acc: 0.9076 - val_loss: 2.0991 - val_acc: 0.6945
Epoch 197/250
- 26s - loss: 0.8740 - acc: 0.9071 - val_loss: 2.1060 - val_acc: 0.6923
Epoch 198/250
- 26s - loss: 0.8635 - acc: 0.9111 - val_loss: 2.0908 - val_acc: 0.6951
Epoch 199/250
- 26s - loss: 0.8698 - acc: 0.9087 - val_loss: 2.0813 - val_acc: 0.6984
Epoch 200/250
- 26s - loss: 0.8717 - acc: 0.9077 - val_loss: 2.0956 - val_acc: 0.6950
Epoch 201/250
- 26s - loss: 0.8653 - acc: 0.9100 - val_loss: 2.0863 - val_acc: 0.6960
Epoch 202/250
```

# Machine Learning Project Report

- 26s - loss: 0.8574 - acc: 0.9116 - val\_loss: 2.0876 - val\_acc: 0.6980  
Epoch 203/250  
- 26s - loss: 0.8584 - acc: 0.9117 - val\_loss: 2.0981 - val\_acc: 0.6943  
Epoch 204/250  
- 26s - loss: 0.8578 - acc: 0.9110 - val\_loss: 2.0912 - val\_acc: 0.6981  
Epoch 205/250  
- 26s - loss: 0.8576 - acc: 0.9122 - val\_loss: 2.0836 - val\_acc: 0.7020  
Epoch 206/250  
- 26s - loss: 0.8548 - acc: 0.9124 - val\_loss: 2.0883 - val\_acc: 0.7042  
Epoch 207/250  
- 26s - loss: 0.8522 - acc: 0.9142 - val\_loss: 2.1009 - val\_acc: 0.7053  
Epoch 208/250  
- 26s - loss: 0.8519 - acc: 0.9144 - val\_loss: 2.0963 - val\_acc: 0.7062  
Epoch 209/250  
- 26s - loss: 0.8483 - acc: 0.9143 - val\_loss: 2.0961 - val\_acc: 0.7088  
Epoch 210/250  
- 26s - loss: 0.8546 - acc: 0.9124 - val\_loss: 2.0914 - val\_acc: 0.6985  
Epoch 211/250  
- 26s - loss: 0.8470 - acc: 0.9143 - val\_loss: 2.0836 - val\_acc: 0.6972  
Epoch 212/250  
- 26s - loss: 0.8459 - acc: 0.9145 - val\_loss: 2.0977 - val\_acc: 0.6939  
Epoch 213/250  
- 26s - loss: 0.8497 - acc: 0.9142 - val\_loss: 2.0966 - val\_acc: 0.6943  
Epoch 214/250  
- 26s - loss: 0.8448 - acc: 0.9145 - val\_loss: 2.0970 - val\_acc: 0.6958  
Epoch 215/250  
- 26s - loss: 0.8454 - acc: 0.9140 - val\_loss: 2.0969 - val\_acc: 0.6955  
Epoch 216/250  
- 26s - loss: 0.8417 - acc: 0.9154 - val\_loss: 2.0915 - val\_acc: 0.6958  
Epoch 217/250  
- 26s - loss: 0.8453 - acc: 0.9139 - val\_loss: 2.0833 - val\_acc: 0.6952  
Epoch 218/250  
- 26s - loss: 0.8430 - acc: 0.9151 - val\_loss: 2.0799 - val\_acc: 0.6989  
Epoch 219/250  
- 26s - loss: 0.8446 - acc: 0.9143 - val\_loss: 2.0890 - val\_acc: 0.6940  
Epoch 220/250  
- 26s - loss: 0.8344 - acc: 0.9179 - val\_loss: 2.0882 - val\_acc: 0.6967  
Epoch 221/250  
- 26s - loss: 0.8393 - acc: 0.9136 - val\_loss: 2.0911 - val\_acc: 0.6965  
Epoch 222/250  
- 26s - loss: 0.8367 - acc: 0.9164 - val\_loss: 2.0883 - val\_acc: 0.6980  
Epoch 223/250  
- 26s - loss: 0.8339 - acc: 0.9175 - val\_loss: 2.0890 - val\_acc: 0.6978  
Epoch 224/250  
- 26s - loss: 0.8322 - acc: 0.9176 - val\_loss: 2.0954 - val\_acc: 0.6962  
Epoch 225/250  
- 26s - loss: 0.8355 - acc: 0.9169 - val\_loss: 2.0955 - val\_acc: 0.6952  
Epoch 226/250  
- 26s - loss: 0.8334 - acc: 0.9175 - val\_loss: 2.0959 - val\_acc: 0.6949  
Epoch 227/250  
- 26s - loss: 0.8425 - acc: 0.9161 - val\_loss: 2.0939 - val\_acc: 0.6975  
Epoch 228/250



# Machine Learning Project Report

```
- 26s - loss: 0.8352 - acc: 0.9162 - val_loss: 2.0900 - val_acc: 0.6965
Epoch 229/250
- 26s - loss: 0.8336 - acc: 0.9164 - val_loss: 2.0895 - val_acc: 0.6959
Epoch 230/250
- 26s - loss: 0.8310 - acc: 0.9186 - val_loss: 2.1002 - val_acc: 0.6943
Epoch 231/250
- 26s - loss: 0.8289 - acc: 0.9189 - val_loss: 2.0969 - val_acc: 0.6962
Epoch 232/250
- 26s - loss: 0.8248 - acc: 0.9189 - val_loss: 2.0951 - val_acc: 0.6958
Epoch 233/250
- 26s - loss: 0.8314 - acc: 0.9158 - val_loss: 2.0938 - val_acc: 0.6965
Epoch 234/250
- 26s - loss: 0.8287 - acc: 0.9180 - val_loss: 2.0954 - val_acc: 0.6966
Epoch 235/250
- 26s - loss: 0.8250 - acc: 0.9199 - val_loss: 2.0869 - val_acc: 0.6979
Epoch 236/250
- 26s - loss: 0.8240 - acc: 0.9184 - val_loss: 2.0966 - val_acc: 0.6980
Epoch 237/250
- 26s - loss: 0.8264 - acc: 0.9186 - val_loss: 2.0965 - val_acc: 0.6963
Epoch 238/250
- 26s - loss: 0.8299 - acc: 0.9178 - val_loss: 2.0946 - val_acc: 0.6954
Epoch 239/250
- 26s - loss: 0.8217 - acc: 0.9194 - val_loss: 2.0942 - val_acc: 0.6969
Epoch 240/250
- 26s - loss: 0.8197 - acc: 0.9199 - val_loss: 2.0992 - val_acc: 0.6959
Epoch 241/250
- 26s - loss: 0.8179 - acc: 0.9214 - val_loss: 2.0944 - val_acc: 0.6979
Epoch 242/250
- 26s - loss: 0.8209 - acc: 0.9189 - val_loss: 2.1030 - val_acc: 0.6961
Epoch 243/250
- 26s - loss: 0.8276 - acc: 0.9190 - val_loss: 2.0962 - val_acc: 0.6962
Epoch 244/250
- 26s - loss: 0.8238 - acc: 0.9188 - val_loss: 2.1088 - val_acc: 0.6987
Epoch 245/250
- 26s - loss: 0.8202 - acc: 0.9202 - val_loss: 2.0956 - val_acc: 0.7076
Epoch 246/250
- 26s - loss: 0.8241 - acc: 0.9189 - val_loss: 2.0996 - val_acc: 0.7083
Epoch 247/250
- 26s - loss: 0.8253 - acc: 0.9175 - val_loss: 2.1002 - val_acc: 0.7091
Epoch 248/250
- 26s - loss: 0.8240 - acc: 0.9192 - val_loss: 2.0982 - val_acc: 0.7113
Epoch 249/250
- 26s - loss: 0.8261 - acc: 0.9178 - val_loss: 2.0966 - val_acc: 0.7123
Epoch 250/250
- 26s - loss: 0.8217 - acc: 0.9205 - val_loss: 2.1067 - val_acc: 0.7152
Training Time : 6496.094993114471
Accuracy: 71.52000000000001 %
```

## Results:

	First Run	Second Run	Third Run	Average
Accuracy	70.88%	71.43%	71.52%	71.28%
Training Time	8284.956s	6461.054s	6469.095	7080.701s

# Machine Learning Project Report

## CALTECH-101

### Introduction :

Pictures of objects belonging to 101 categories. About 40 to 800 images per category. Most categories have about 50 images. Collected in September 2003 by Fei-Fei Li, Marco Andreetto, and Marc 'Aurelio Ranzato. The size of each image is roughly 300 x 200 pixels.

Dataset Link: [http://www.vision.caltech.edu/Image\\_Datasets/Caltech101/#Download](http://www.vision.caltech.edu/Image_Datasets/Caltech101/#Download)

### Model:

ResNet50 model with weights pre-trained on ImageNet.

### ResNet50 Architecture:

ResNet, short for Residual Networks is a classic neural network used as a backbone for many computer vision tasks. This model was the winner of ImageNet challenge in 2015. The fundamental breakthrough with ResNet was it allowed us to train extremely deep neural networks with 150+layers successfully. Prior to ResNet training very deep neural networks was difficult due to the problem of vanishing gradients.

### Source Code:

```
# -*- coding: utf-8 -*-  
"""
```

Created on Sun Nov 24 20:02:56 2019

```
@author: Rahul  
"""
```

```
# -*- coding: utf-8 -*-  
"""Untitled8.ipynb
```

Automatically generated by Colaboratory.

Original file is located at

```
https://colab.research.google.com/drive/1wj-aNzWsigjcnczEjz1Qae2X1AtbjWW6  
"""
```

```
# Commented out IPython magic to ensure Python compatibility.  
import numpy as np  
import pandas as pd  
import matplotlib.pyplot as plt  
# %matplotlib inline  
#from keras.datasets import mnist  
from keras import *  
from keras.models import Sequential,Model  
from keras.layers import Dense, Dropout, Activation, Flatten  
from keras.optimizers import Adam,SGD  
from keras.layers.normalization import BatchNormalization  
from keras.utils import np_utils
```

# Machine Learning Project Report

```
from keras.layers import Conv2D, MaxPooling2D, ZeroPadding2D, GlobalAveragePooling2D
from keras.layers.advanced_activations import LeakyReLU
from keras.preprocessing.image import ImageDataGenerator
from keras.preprocessing import image
from sklearn.model_selection import train_test_split
import glob
import time
from skimage import io
import os
from imageio import imread
from skimage.transform import resize
from keras import regularizers
import csv
from keras.applications.inception_v3 import InceptionV3, preprocess_input

datasets_path = './101_ObjectCategories/101_ObjectCategories' #Add the path to the unzipped folder

def load_images(path,n=0):
    X_train = []
    Y_train = []
    X_test = []
    Y_test = []
    i=-1
    labels = []
    for label in os.listdir(path):
        imageCount = 0
        back_path = path + '/' + label
        labels.append(label)
        i = i+1
        count = 0
        for filename in os.listdir(back_path):
            imageCount += 1
        for filename in os.listdir(back_path):
            image_path = path + '/' + label + '/' + filename
            print(image_path)
            # os.path.join(back_path,filename)
            img = image.load_img(image_path,target_size=(224,224))
            if count < 30:
                img = image.img_to_array(img)
                #Y.append(image)
                #image = imresize(image,[128,128,3])
                #image = imresize(imread(image_path), [128,128, 3])
                #image = image.astype('float32')
                img[:, :, 0] -= 123.68
                img[:, :, 1] -= 116.78
                img[:, :, 2] -= 103.94
                #image = image/255
                #image = 1-image
                Y_train.append(i)
                X_train.append(img)
                count += 1
            else:
                img = image.img_to_array(img)
                #Y.append(image)
```

# Machine Learning Project Report

```
#image = imresize(image,[128,128,3])
#image = imresize(imread(image_path), [128,128, 3])
#image = image.astype('float32')
img[:, :, 0] -= 123.68
img[:, :, 1] -= 116.78
img[:, :, 2] -= 103.94
#image = image/255
#image = 1-image
Y_test.append(i)
X_test.append(img)
count += 1

print(count)
#X.append(image.img_to_array(img))
#X.append(image)
return X_train,Y_train,X_test,Y_test,labels

X_train,Y_train,X_test,Y_test,labels = load_images(datasets_path)
X_train = np.array(X_train)
Y_train = np.array(Y_train)
X_test = np.array(X_test)
Y_test = np.array(Y_test)
##X = np.array(x_train)
###print(len(x_train))
###print(x_train.shape)
##Y = np.array(y)
##print(X.shape)
##
##X_train, X_test, y_train, y_test = train_test_split(X, Y, test_size = 0.2)
##
##print(y_train.shape)
##print(X_train.shape)
##
number_of_classes = 101
Y_train = np_utils.to_categorical(Y_train-1, number_of_classes)
Y_test = np_utils.to_categorical(Y_test-1, number_of_classes)
gen = ImageDataGenerator(width_shift_range=.2,
                          height_shift_range=.2,
                          zoom_range=0.2)

test_gen = ImageDataGenerator()
train_generator = gen.flow(X_train, Y_train, batch_size=16)
test_generator = test_gen.flow(X_test, Y_test, batch_size=16)

from keras.applications.vgg16 import VGG16
from keras.applications.resnet50 import ResNet50

def add_new_last_layer(base_model, nb_classes):
    x = base_model.output
    x = GlobalAveragePooling2D()(x)
    x = Dense(512, activation='relu')(x)
    predictions = Dense(nb_classes, activation='softmax')(x)
    model = Model(input=base_model.input, output=predictions)
    return model
```

# Machine Learning Project Report

```
base_model = ResNet50(weights='imagenet', include_top=False) #include_top=False excludes final FC
layer
model = add_new_last_layer(base_model, 101)

model.compile(loss='categorical_crossentropy', optimizer=SGD(lr = 1e-3,momentum =
0.9),metrics=['accuracy'])
start = time.time()
x = model.fit_generator(train_generator, epochs=5, shuffle = True,
                        validation_data=test_generator)
end = time.time()
#print("Training time: ",(end-start))

score = model.evaluate(X_test, Y_test)
print()
print('Test loss: ', score[0])
print('Test Accuracy', score[1])
```

## Output:

### Run 1:

```
WARNING:tensorflow:From /usr/local/lib/python3.6/dist-
packages/tensorflow_core/python/ops/math_grad.py:1424: where (from
tensorflow.python.ops.array_ops) is deprecated and will be removed in a
future version.
Instructions for updating:
Use tf.where in 2.0, which has the same broadcast rule as np.where
WARNING:tensorflow:From /usr/local/lib/python3.6/dist-
packages/keras/backend/tensorflow_backend.py:1033: The name tf.assign_add
is deprecated. Please use tf.compat.v1.assign_add instead.
```

```
WARNING:tensorflow:From /usr/local/lib/python3.6/dist-
packages/keras/backend/tensorflow_backend.py:1020: The name tf.assign is
deprecated. Please use tf.compat.v1.assign instead.
```

```
Epoch 1/5
457/458 [=====>.] - ETA: 0s - loss: 1.7060 - acc:
0.6379Epoch 2/5
458/458 [=====] - 81s 176ms/step - loss: 0.3349 -
acc: 0.9142 - val_loss: 0.3554 - val_acc: 0.9032
Epoch 3/5
458/458 [=====] - 81s 176ms/step - loss: 0.1533 -
acc: 0.9611 - val_loss: 0.2457 - val_acc: 0.9295
Epoch 4/5
458/458 [=====] - 80s 175ms/step - loss: 0.0911 -
acc: 0.9775 - val_loss: 0.2680 - val_acc: 0.9338
Epoch 5/5
458/458 [=====] - 82s 178ms/step - loss: 0.0622 -
acc: 0.9850 - val_loss: 0.2467 - val_acc: 0.9338
```

# Machine Learning Project Report

Training time: 412.1241245702147

1829/1829 [=====] - 5s 2ms/step

Test loss: 0.24673965746030707

Test Accuracy 0.9338436304317138

Run 2:

Epoch 1/5

458/458 [=====] - 91s 199ms/step - loss: 1.6579 -  
acc: 0.6431 - val\_loss: 0.5237 - val\_acc: 0.8562

Epoch 2/5

458/458 [=====] - 82s 178ms/step - loss: 0.3224 -  
acc: 0.9168 - val\_loss: 0.2791 - val\_acc: 0.9229

Epoch 3/5

458/458 [=====] - 81s 178ms/step - loss: 0.1511 -  
acc: 0.9623 - val\_loss: 0.2713 - val\_acc: 0.9218

Epoch 4/5

458/458 [=====] - 81s 177ms/step - loss: 0.0898 -  
acc: 0.9772 - val\_loss: 0.2368 - val\_acc: 0.9344

Epoch 5/5

458/458 [=====] - 82s 179ms/step - loss: 0.0696 -  
acc: 0.9828 - val\_loss: 0.2550 - val\_acc: 0.9349

Training time: 458.3651282504564

1829/1829 [=====] - 4s 2ms/step

Test loss: 0.2549842904668613

Test Accuracy 0.934937124144125

Run 3:

Epoch 1/5

458/458 [=====] - 94s 205ms/step - loss: 1.6349 -  
acc: 0.6502 - val\_loss: 0.4725 - val\_acc: 0.8721

Epoch 2/5

458/458 [=====] - 81s 177ms/step - loss: 0.3229 -  
acc: 0.9178 - val\_loss: 0.3017 - val\_acc: 0.9142

Epoch 3/5

458/458 [=====] - 81s 176ms/step - loss: 0.1407 -  
acc: 0.9636 - val\_loss: 0.2450 - val\_acc: 0.9284

Epoch 4/5

458/458 [=====] - 81s 177ms/step - loss: 0.0868 -  
acc: 0.9794 - val\_loss: 0.2322 - val\_acc: 0.9273

Epoch 5/5

458/458 [=====] - 81s 178ms/step - loss: 0.0568 -  
acc: 0.9869 - val\_loss: 0.2173 - val\_acc: 0.9404

Training time: 422.0074882507324

1829/1829 [=====] - 4s 2ms/step

# Machine Learning Project Report

Test loss: 0.2173448165887819  
Test Accuracy 0.9404045927387694

## Result:

	First Run	Second Run	Third Run	Average
Accuracy	93.38%	93.49%	94.04%	93.64%
Training Time	412.129s	458.365s	422.007s	430.83s

## CALTECH-256

### Introduction:

Caltech-256 dataset have 257 categories with 30608 total pictures.

Dataset Link: [http://www.vision.caltech.edu/Image\\_Datasets/Caltech256/256\\_ObjectCategories.tar](http://www.vision.caltech.edu/Image_Datasets/Caltech256/256_ObjectCategories.tar)

### Model:

InceptionResNetV2 model with weights pre-trained on ImageNet.

### InceptionResNetV2:

Inception-v4, evolved from GoogLeNet / Inception-v1, has a more uniform simplified architecture and more inception modules than Inception-v3

### Source Code:

```
import os

import numpy as np

import pandas as pd

import matplotlib.pyplot as plt

import tarfile

#import cv2

import pickle

import time

from keras.preprocessing import image
```

# Machine Learning Project Report

#root path is /content/drive/My Drive/caltech256

datasets\_path = './256\_ObjectCategories' #Add the path to the unzipped folder

```
def load_images(path,n=0):
    X_train = []
    Y_train = []
    X_test = []
    Y_test = []
    i=-1
    labels = []
    for label in os.listdir(path):
        imageCount = 0
        back_path = path + '/' + label
        labels.append(label)
        i = i+1
        count = 0
        for filename in os.listdir(back_path):
            imageCount += 1
        for filename in os.listdir(back_path):
            image_path = path + '/' + label + '/' + filename
            print(image_path)
            # os.path.join(back_path,filename)
            img = image.load_img(image_path,target_size=(224,224))
            if count < 30:
                img = image.img_to_array(img)
                #Y.append(image)
                #image = imresize(image,[128,128,3])
                #image = imresize(imread(image_path), [128,128, 3])
```



# Machine Learning Project Report

```
#image = image.astype('float32')

img[:, :, 0] -= 123.68
img[:, :, 1] -= 116.78
img[:, :, 2] -= 103.94

#image = image/255

#image = 1-image
Y_train.append(i)
X_train.append(img)

count += 1

else:

    img = image.img_to_array(img)
    #Y.append(image)

    #image = imresize(image,[128,128,3])
    #image = imresize(imread(image_path), [128,128, 3])
    #image = image.astype('float32')

    img[:, :, 0] -= 123.68
    img[:, :, 1] -= 116.78
    img[:, :, 2] -= 103.94

    #image = image/255

    #image = 1-image
    Y_test.append(i)
    X_test.append(img)

    count += 1


print(count)

#X.append(image.img_to_array(img))
#X.append(image)

return X_train,Y_train,X_test,Y_test,labels
```

# Machine Learning Project Report

```
X_train,Y_train,X_test,Y_test,labels = load_images(datasets_path)

X_train = np.array(X_train)
y_train = np.array(Y_train)
X_test = np.array(X_test)
y_test = np.array(Y_test)
img_size = 128
X_train = X_train.reshape(-1,img_size,img_size,3)

X_test = X_test.reshape(-1,img_size,img_size,3)

from keras.preprocessing.image import ImageDataGenerator, array_to_img, img_to_array, load_img

generator = ImageDataGenerator(
    rotation_range=40,
    width_shift_range=0.2,
    height_shift_range=0.2,
    shear_range=0.2,
    zoom_range=0.2,
    horizontal_flip=True,
    fill_mode='nearest')

from keras.models import Sequential
from keras.models import Model
from keras.layers import Dense, Activation, Dropout
import keras.backend as K
from keras.optimizers import Adam
from keras.layers.normalization import BatchNormalization
```

# Machine Learning Project Report

```
from keras.layers import Conv2D, MaxPool2D, AvgPool2D, Flatten, GlobalAveragePooling2D
```

```
batch_size = 64
```

```
time.sleep(100)
```

```
from keras.applications.inception_resnet_v2 import InceptionResNetV2
```

```
base_model = InceptionResNetV2(weights='imagenet', include_top=False)
```

```
x = base_model.output
```

```
x = GlobalAveragePooling2D()(x)
```

```
x = Dense(1024, activation='relu')(x)
```

```
predictions = Dense(257, activation='softmax')(x)
```

```
model = Model(inputs=base_model.input, outputs=predictions)
```

```
for layer in base_model.layers:
```

```
    layer.trainable = False
```

```
adam = Adam(lr=0.0001)
```

```
model.compile(optimizer=adam, loss='categorical_crossentropy', metrics=['accuracy'])
```

```
model.fit_generator(generator.flow(X_train, y_train.values, batch_size=batch_size), len(X_train) /  
batch_size, epochs=20, verbose=1, validation_data=(X_test, y_test.values))
```

```
for i, layer in enumerate(base_model.layers):
```

```
    print(i, layer.name)
```

# Machine Learning Project Report

```
for layer in model.layers[:249]:
```

```
    layer.trainable = False
```

```
for layer in model.layers[249:]:
```

```
    layer.trainable = True
```

```
time.sleep(100)
```

```
from keras.optimizers import SGD
```

```
model.compile(optimizer=adam, loss='categorical_crossentropy', metrics=['accuracy'])
```

```
start = time.time()
```

```
model.fit_generator(generator.flow(X_train, y_train.values, batch_size=batch_size), len(X_train) /  
batch_size, epochs=30, verbose=1, validation_data=(X_test, y_test.values))
```

```
end = time.time()
```

```
print("Training time: ", (end-start))
```

```
from keras.models import load_model
```

```
model.save('inceptionresnet_model.h5')
```

```
del model
```

```
model = load_model('inceptionresnet_model.h5')
```

```
score = model.evaluate(x = X_test, y = y_test.values)
```

```
print("Accuracy: ", score[1], "%")
```

Output:

Run 1:

Epoch 1/30

463/462 [=====] - 249s 538ms/step - loss: 1.8827  
- acc: 0.5618 - val\_loss: 1.4449 - val\_acc: 0.6586

Epoch 2/30

463/462 [=====] - 225s 486ms/step - loss: 1.1246  
- acc: 0.7152 - val\_loss: 1.2406 - val\_acc: 0.7043

Epoch 3/30

# Machine Learning Project Report

```
463/462 [=====] - 225s 486ms/step - loss: 0.8536
- acc: 0.7740 - val_loss: 1.2254 - val_acc: 0.7179
Epoch 4/30
463/462 [=====] - 225s 487ms/step - loss: 0.6785
- acc: 0.8189 - val_loss: 1.1927 - val_acc: 0.7335
Epoch 5/30
463/462 [=====] - 225s 486ms/step - loss: 0.5486
- acc: 0.8500 - val_loss: 1.2124 - val_acc: 0.7247
Epoch 6/30
463/462 [=====] - 225s 486ms/step - loss: 0.4598
- acc: 0.8727 - val_loss: 1.2334 - val_acc: 0.7247
Epoch 7/30
463/462 [=====] - 225s 486ms/step - loss: 0.4006
- acc: 0.8846 - val_loss: 1.3307 - val_acc: 0.7033
Epoch 8/30
463/462 [=====] - 225s 486ms/step - loss: 0.3421
- acc: 0.9012 - val_loss: 1.3250 - val_acc: 0.7296
Epoch 9/30
463/462 [=====] - 225s 486ms/step - loss: 0.3047
- acc: 0.9132 - val_loss: 1.3632 - val_acc: 0.7267
Epoch 10/30
463/462 [=====] - 225s 486ms/step - loss: 0.2635
- acc: 0.9206 - val_loss: 1.3884 - val_acc: 0.7208
Epoch 11/30
463/462 [=====] - 225s 486ms/step - loss: 0.2420
- acc: 0.9288 - val_loss: 1.3988 - val_acc: 0.7150
Epoch 12/30
463/462 [=====] - 225s 486ms/step - loss: 0.2324
- acc: 0.9323 - val_loss: 1.4727 - val_acc: 0.7189
Epoch 13/30
463/462 [=====] - 225s 486ms/step - loss: 0.2139
- acc: 0.9374 - val_loss: 1.4974 - val_acc: 0.7140
Epoch 14/30
463/462 [=====] - 225s 486ms/step - loss: 0.1931
- acc: 0.9422 - val_loss: 1.5419 - val_acc: 0.7150
Epoch 15/30
463/462 [=====] - 225s 486ms/step - loss: 0.1882
- acc: 0.9454 - val_loss: 1.4930 - val_acc: 0.7179
Epoch 16/30
463/462 [=====] - 225s 486ms/step - loss: 0.1702
- acc: 0.9497 - val_loss: 1.5631 - val_acc: 0.7179
Epoch 17/30
463/462 [=====] - 225s 485ms/step - loss: 0.1607
- acc: 0.9533 - val_loss: 1.5291 - val_acc: 0.7276
Epoch 18/30
463/462 [=====] - 225s 486ms/step - loss: 0.1508
- acc: 0.9541 - val_loss: 1.6081 - val_acc: 0.7101
Epoch 19/30
463/462 [=====] - 225s 486ms/step - loss: 0.1598
- acc: 0.9520 - val_loss: 1.5886 - val_acc: 0.7160
Epoch 20/30
```

# Machine Learning Project Report

```
463/462 [=====] - 225s 486ms/step - loss: 0.1474
- acc: 0.9558 - val_loss: 1.5343 - val_acc: 0.7208
Epoch 21/30
463/462 [=====] - 225s 486ms/step - loss: 0.1309
- acc: 0.9605 - val_loss: 1.6192 - val_acc: 0.7208
Epoch 22/30
463/462 [=====] - 225s 486ms/step - loss: 0.1307
- acc: 0.9607 - val_loss: 1.6476 - val_acc: 0.7198
Epoch 23/30
463/462 [=====] - 225s 486ms/step - loss: 0.1314
- acc: 0.9606 - val_loss: 1.5573 - val_acc: 0.7276
Epoch 24/30
463/462 [=====] - 225s 486ms/step - loss: 0.1226
- acc: 0.9630 - val_loss: 1.6265 - val_acc: 0.7072
Epoch 25/30
463/462 [=====] - 225s 485ms/step - loss: 0.1185
- acc: 0.9636 - val_loss: 1.7588 - val_acc: 0.7053
Epoch 26/30
463/462 [=====] - 225s 486ms/step - loss: 0.1118
- acc: 0.9657 - val_loss: 1.7262 - val_acc: 0.7121
Epoch 27/30
463/462 [=====] - 225s 486ms/step - loss: 0.1152
- acc: 0.9662 - val_loss: 1.7283 - val_acc: 0.7111
Epoch 28/30
463/462 [=====] - 225s 486ms/step - loss: 0.1121
- acc: 0.9668 - val_loss: 1.6804 - val_acc: 0.7023
Epoch 29/30
463/462 [=====] - 225s 486ms/step - loss: 0.1041
- acc: 0.9684 - val_loss: 1.7268 - val_acc: 0.7160
Epoch 30/30
463/462 [=====] - 225s 486ms/step - loss: 0.1121
- acc: 0.9666 - val_loss: 1.6555 - val_acc: 0.7150
```

Training time: 6750.564

```
1028/1028 [=====] - 5s 5ms/step
```

```
[1.655475185697777, 0.71498054474708173]
```

Accuracy: 71.4980544%

## Run 2:

```
Epoch 1/30
463/462 [=====] - 249s 538ms/step - loss: 1.8765
- acc: 0.5656 - val_loss: 1.453 - val_acc: 0.6542
Epoch 2/30
463/462 [=====] - 225s 486ms/step - loss: 1.1289
- acc: 0.7154 - val_loss: 1.2397 - val_acc: 0.7088
Epoch 3/30
463/462 [=====] - 225s 486ms/step - loss: 0.8578
- acc: 0.7696 - val_loss: 1.2298 - val_acc: 0.7165
Epoch 4/30
```

# Machine Learning Project Report

```
463/462 [=====] - 225s 487ms/step - loss: 0.6745
- acc: 0.8154 - val_loss: 1.1976 - val_acc: 0.7332
Epoch 5/30
463/462 [=====] - 225s 486ms/step - loss: 0.5398
- acc: 0.8505 - val_loss: 1.2134 - val_acc: 0.7256
Epoch 6/30
463/462 [=====] - 225s 486ms/step - loss: 0.4589
- acc: 0.8745 - val_loss: 1.2336 - val_acc: 0.7298
Epoch 7/30
463/462 [=====] - 225s 486ms/step - loss: 0.4006
- acc: 0.8846 - val_loss: 1.3307 - val_acc: 0.7033
Epoch 8/30
463/462 [=====] - 225s 486ms/step - loss: 0.3421
- acc: 0.9012 - val_loss: 1.3250 - val_acc: 0.7296
Epoch 9/30
463/462 [=====] - 225s 486ms/step - loss: 0.3047
- acc: 0.9132 - val_loss: 1.3632 - val_acc: 0.7267
Epoch 10/30
463/462 [=====] - 225s 486ms/step - loss: 0.2635
- acc: 0.9206 - val_loss: 1.3884 - val_acc: 0.7208
Epoch 11/30
463/462 [=====] - 225s 486ms/step - loss: 0.2420
- acc: 0.9288 - val_loss: 1.3988 - val_acc: 0.7150
Epoch 12/30
463/462 [=====] - 225s 486ms/step - loss: 0.2324
- acc: 0.9323 - val_loss: 1.4727 - val_acc: 0.7189
Epoch 13/30
463/462 [=====] - 225s 486ms/step - loss: 0.2139
- acc: 0.9374 - val_loss: 1.4974 - val_acc: 0.7140
Epoch 14/30
463/462 [=====] - 225s 486ms/step - loss: 0.1931
- acc: 0.9422 - val_loss: 1.5419 - val_acc: 0.7150
Epoch 15/30
463/462 [=====] - 225s 486ms/step - loss: 0.1882
- acc: 0.9454 - val_loss: 1.4930 - val_acc: 0.7179
Epoch 16/30
463/462 [=====] - 225s 486ms/step - loss: 0.1702
- acc: 0.9497 - val_loss: 1.5631 - val_acc: 0.7179
Epoch 17/30
463/462 [=====] - 225s 485ms/step - loss: 0.1607
- acc: 0.9533 - val_loss: 1.5291 - val_acc: 0.7276
Epoch 18/30
463/462 [=====] - 225s 486ms/step - loss: 0.1508
- acc: 0.9541 - val_loss: 1.6081 - val_acc: 0.7101
Epoch 19/30
463/462 [=====] - 225s 486ms/step - loss: 0.1598
- acc: 0.9520 - val_loss: 1.5886 - val_acc: 0.7160
Epoch 20/30
463/462 [=====] - 225s 486ms/step - loss: 0.1474
- acc: 0.9558 - val_loss: 1.5343 - val_acc: 0.7208
Epoch 21/30
```

# Machine Learning Project Report

```
463/462 [=====] - 225s 486ms/step - loss: 0.1309
- acc: 0.9605 - val_loss: 1.6192 - val_acc: 0.7208
Epoch 22/30
463/462 [=====] - 225s 486ms/step - loss: 0.1307
- acc: 0.9607 - val_loss: 1.6476 - val_acc: 0.7198
Epoch 23/30
463/462 [=====] - 225s 486ms/step - loss: 0.1314
- acc: 0.9606 - val_loss: 1.5573 - val_acc: 0.7276
Epoch 24/30
463/462 [=====] - 225s 486ms/step - loss: 0.1226
- acc: 0.9630 - val_loss: 1.6265 - val_acc: 0.7072
Epoch 25/30
463/462 [=====] - 225s 485ms/step - loss: 0.1185
- acc: 0.9636 - val_loss: 1.7588 - val_acc: 0.7053
Epoch 26/30
463/462 [=====] - 225s 486ms/step - loss: 0.1118
- acc: 0.9657 - val_loss: 1.7262 - val_acc: 0.7121
Epoch 27/30
463/462 [=====] - 225s 486ms/step - loss: 0.1198
- acc: 0.9687 - val_loss: 1.7284 - val_acc: 0.7116
Epoch 28/30
463/462 [=====] - 225s 486ms/step - loss: 0.1167
- acc: 0.9656 - val_loss: 1.6813 - val_acc: 0.7045
Epoch 29/30
463/462 [=====] - 225s 486ms/step - loss: 0.1141
- acc: 0.9666 - val_loss: 1.7260 - val_acc: 0.7148
Epoch 30/30
463/462 [=====] - 225s 486ms/step - loss: 0.1119
- acc: 0.9667 - val_loss: 1.6556 - val_acc: 0.7150
```

Training time: 6767.678

```
1028/1028 [=====] - 5s 5ms/step
```

```
[1.655475185697777, 0.71501234856475435]
```

Accuracy: 71.5012348%

## Run 3:

```
Epoch 1/30
463/462 [=====] - 248s 537ms/step - loss: 1.8727
- acc: 0.5619 - val_loss: 1.4469 - val_acc: 0.6588
Epoch 2/30
463/462 [=====] - 225s 486ms/step - loss: 1.1226
- acc: 0.7052 - val_loss: 1.2406 - val_acc: 0.7033
Epoch 3/30
463/462 [=====] - 225s 486ms/step - loss: 0.8536
- acc: 0.7540 - val_loss: 1.2354 - val_acc: 0.7169
Epoch 4/30
463/462 [=====] - 225s 487ms/step - loss: 0.6775
- acc: 0.8189 - val_loss: 1.1827 - val_acc: 0.7355
Epoch 5/30
```



# Machine Learning Project Report

```
463/462 [=====] - 225s 486ms/step - loss: 0.5486
- acc: 0.8500 - val_loss: 1.2124 - val_acc: 0.7247
Epoch 6/30
463/462 [=====] - 225s 486ms/step - loss: 0.4598
- acc: 0.8727 - val_loss: 1.2334 - val_acc: 0.7247
Epoch 7/30
463/462 [=====] - 225s 486ms/step - loss: 0.4006
- acc: 0.8846 - val_loss: 1.3307 - val_acc: 0.7033
Epoch 8/30
463/462 [=====] - 225s 486ms/step - loss: 0.3421
- acc: 0.9012 - val_loss: 1.3250 - val_acc: 0.7296
Epoch 9/30
463/462 [=====] - 225s 486ms/step - loss: 0.3047
- acc: 0.9132 - val_loss: 1.3632 - val_acc: 0.7267
Epoch 10/30
463/462 [=====] - 225s 486ms/step - loss: 0.2635
- acc: 0.9206 - val_loss: 1.3884 - val_acc: 0.7208
Epoch 11/30
463/462 [=====] - 225s 486ms/step - loss: 0.2420
- acc: 0.9288 - val_loss: 1.3988 - val_acc: 0.7150
Epoch 12/30
463/462 [=====] - 225s 486ms/step - loss: 0.2324
- acc: 0.9323 - val_loss: 1.4727 - val_acc: 0.7189
Epoch 13/30
463/462 [=====] - 225s 486ms/step - loss: 0.2139
- acc: 0.9374 - val_loss: 1.4974 - val_acc: 0.7140
Epoch 14/30
463/462 [=====] - 225s 486ms/step - loss: 0.1931
- acc: 0.9422 - val_loss: 1.5419 - val_acc: 0.7150
Epoch 15/30
463/462 [=====] - 225s 486ms/step - loss: 0.1882
- acc: 0.9454 - val_loss: 1.4930 - val_acc: 0.7179
Epoch 16/30
463/462 [=====] - 225s 486ms/step - loss: 0.1702
- acc: 0.9497 - val_loss: 1.5631 - val_acc: 0.7179
Epoch 17/30
463/462 [=====] - 225s 485ms/step - loss: 0.1607
- acc: 0.9533 - val_loss: 1.5291 - val_acc: 0.7276
Epoch 18/30
463/462 [=====] - 225s 486ms/step - loss: 0.1508
- acc: 0.9541 - val_loss: 1.6081 - val_acc: 0.7101
Epoch 19/30
463/462 [=====] - 225s 486ms/step - loss: 0.1598
- acc: 0.9520 - val_loss: 1.5886 - val_acc: 0.7160
Epoch 20/30
463/462 [=====] - 225s 486ms/step - loss: 0.1474
- acc: 0.9558 - val_loss: 1.5343 - val_acc: 0.7208
Epoch 21/30
463/462 [=====] - 225s 486ms/step - loss: 0.1309
- acc: 0.9605 - val_loss: 1.6192 - val_acc: 0.7208
Epoch 22/30
```

# Machine Learning Project Report

```
463/462 [=====] - 225s 486ms/step - loss: 0.1307
- acc: 0.9607 - val_loss: 1.6476 - val_acc: 0.7198
Epoch 23/30
463/462 [=====] - 225s 486ms/step - loss: 0.1314
- acc: 0.9606 - val_loss: 1.5573 - val_acc: 0.7276
Epoch 24/30
463/462 [=====] - 225s 486ms/step - loss: 0.1226
- acc: 0.9630 - val_loss: 1.6265 - val_acc: 0.7072
Epoch 25/30
463/462 [=====] - 225s 485ms/step - loss: 0.1185
- acc: 0.9636 - val_loss: 1.7588 - val_acc: 0.7053
Epoch 26/30
463/462 [=====] - 225s 486ms/step - loss: 0.1118
- acc: 0.9657 - val_loss: 1.7262 - val_acc: 0.7121
Epoch 27/30
463/462 [=====] - 225s 486ms/step - loss: 0.1145
- acc: 0.9642 - val_loss: 1.7273 - val_acc: 0.7110
Epoch 28/30
463/462 [=====] - 225s 486ms/step - loss: 0.1121
- acc: 0.9768 - val_loss: 1.6704 - val_acc: 0.7021
Epoch 29/30
463/462 [=====] - 225s 486ms/step - loss: 0.1041
- acc: 0.9674 - val_loss: 1.7258 - val_acc: 0.7150
Epoch 30/30
463/462 [=====] - 225s 486ms/step - loss: 0.1120
- acc: 0.9656 - val_loss: 1.6455 - val_acc: 0.7138
```

Training time: 6843.324

1028/1028 [=====] - 5s 5ms/step

[1.655475185697787, 0.71498054484808173]

Accuracy: 71.3980622%

## Result:

	First Run	Second Run	Third Run	Average
Accuracy	71.49%	71.50%	71.40%	71.46%
Training Time	6750.564s	6767.678s	6843.324s	6787.188s

## Final Accuracy Result:

# Machine Learning Project Report

Accuracy from Cifar-10: 93.38%

Accuracy from Cifar-100: 71.28.%

Accuracy from Caltech-101: 93.64%

Accuracy from Caltech-256: 71.46%

Overall Accuracy = 82.44%